



COMPETITOR ANALYSIS

Growing slower every year, spending less to do it — and its newest capability was bought, not built.

THEIR DOCUMENTATION

1,352 pages

SEC FILINGS

737 documents

AUDITED FINANCIALS

FY2019 → now

PUBLIC CODE

137 repos

STATUS PAGE

451 incidents

SUB-PROCESSORS

17 outside suppliers

Toma Pirtskhelani · Product Manager

Public sources only · September 2026

Four parts, three rules

PART I

The company

who owns it • what it costs • who buys it

PART II

The product

seven stages of one campaign

PART III

Strategy

where the money goes • what protects them

PART IV

Open questions

what public sources cannot answer

AND THREE RULES THAT APPLY TO ALL FOUR

Every claim is graded

Marked by how strong its source is. The bar at the foot of each slide shows where we are.

Marketing is labelled as marketing

Their own words appear — but never as evidence.

Gaps become a backlog

Anything needing a non-public source goes on the open-questions list, and grows as you add to it.

And ten expectations were written down before any of this was read. Four turned out wrong — slide 40 shows which.

How to judge every claim here

STRONG

Audited filings, and their own technical documentation

Signed off by auditors and filed under legal penalty — or behaviour the platform must actually support

MEDIUM

Independent records

Customer reviews, infrastructure traces, incident history, real transacted prices

WEAK

Marketing, or disputed

Their own claims, or sources that contradict each other

Five grades are kept in the evidence record. Three reach the slides, mapped in docs/EVIDENCE-GRADERS.md. **A claim takes the grade of its weakest supporting source, never its best.**

PART I

The company

- 01 Who they are, and who controls them now
- 02 What they bought, and what it cost
- 03 What a customer pays — bounded, not guessed
- 04 Who buys it, and who buyers compare them against

AUDITED FILINGS UNLESS MARKED OTHERWISE

Four documents a company cannot write freely

The 10-K

Their **audited annual report to the US regulator**. Signed by the chief executive and the finance chief, and wrong at legal risk. **Seven years of them**.

The proxy statement

Filed before the shareholder meeting. Must name **every executive officer and director**, their pay, and who owns the company.

The sub-processor disclosure

A public list of **every outside supplier that touches customer data**. The law obliges them to keep it complete, so it names middlemen no marketing page would.

The status page

A live record of **every outage since 2016**, written during the outage. Nobody writes one of these to look good.

None of these is marketing. Three are filed under legal penalty and the fourth is written under pressure — so wherever this deck and Braze's website disagree, the documents win.

If you remember five things

Braze grades its own data delays, and three of four are slow

Their documentation labels three of the four ways data gets in “**not real-time**”. Their table, not our characterisation.

The AI decisioning engine was bought, not built

OfferFit, acquired June 2025 for **\$303.2m**. The models come from Anthropic, OpenAI and Google.

One instance is not on the same cloud as the others

The addresses Braze tells you to allowlist for **US-08** belong to Microsoft. Every other instance’s belong to Amazon.

They are slowing down and getting more efficient at the same time

Growth has halved since FY2023 while sales spend fell further. **Three straight years of positive cash flow.**

Braze names four competitors. Buyers compare them against eight

The five extra names are specialists Braze never mentions.



Their biggest claim, in their most careful document

“Our platform empowers **real-time engagement** between brands and their customers ... made possible by our proprietary, enterprise-grade **stream processing architecture** ... We have designed it to **listen like a human would**, process new information in context, and **react instantaneously**.”

Braze 10-K, Item 1 · filed 25 March 2026

The claim

A single platform, fed by streaming first-party data, reacting in the moment across every channel

Why it is quoted here

This is the company’s own narrative in a filed document — the strongest form of a marketing claim, not a technical one

What to hold on to

“Real-time” is doing a lot of work in that sentence. Slide 21 shows their own table grading it

Origins, and who controls it now



Bill Magnuson

CHAIRMAN • CEO • PRESIDENT • COFOUNDER

CTO from July 2011, chief executive since 2017, chairman — and President too since June 2025, when the previous President resigned. One person, four titles.

Cofounder Jon Hyman is still CTO, an officer since 2011.

AND WHO HOLDS THE SEATS

Seven seats, six independent

Magnuson is the exception. **Phillip Fernandez is Lead Independent Director** — the role a board creates when its chair is not.

And the board is classified

Three staggered classes, so only about **a third stands in any year**. The super-voting stock is gone; this is not.

The votes, after January

Class B converted **30 January 2026**. Largest holder **6.0%**, the CEO **4.9%**, the Battery partner **5.1%** — **no blocking position left**.

7.7× bigger, a billion already contracted



\$456.8m

NET IPO PROCEEDS, NOV 2021

7.7×

REVENUE GROWTH, FY2020 TO FY2026

\$124.3m

CASH AT FY2026 YEAR END

\$1,033.0m

CONTRACTED, NOT YET RECOGNISED

THE CAVEAT THAT TRAVELS WITH EVERY NUMBER IN THIS PART

Braze disclosed a **material weakness in internal control over financial reporting** at 31 January 2026 — *a flaw in the checking, not an error in the numbers*. Ineffective IT controls over **user access and program change management** on the systems that produce these figures. **And, in the same breath:** it “did not result in any identified misstatements”, nothing was restated, and Ernst & Young attested.

They bought their AI, and the filing says so

“the Company completed the acquisition of **OfferFit, Inc.** (‘OfferFit’) **which is now known as AI Decisioning Studio** for total consideration of **\$303.2 million.**”

Ernst & Young, critical audit matter · Braze 10-K, 25 March 2026

\$303.2m

TOTAL CONSIDERATION

77%

WAS GOODWILL — PRICE
NOT TIED TO ANY ASSET

\$56.7m

DEVELOPED TECHNOLOGY,
AMORTISED TO COST OF REVENUE

2 Jun 2025

CLOSED

WHAT THE PURCHASE-PRICE ALLOCATION TELLS YOU

Only **\$66.6m of the \$303.2m was identifiable**, and \$56.7m of that is the software — so Braze bought a capability and its people, not revenue. The trademark went at **\$0.9m**: nobody expected the name to survive. **So the margin drag runs to 2031, and any AI pricing they quote has to carry it.**

They bought a market, and the earn-out paid nothing

\$26.8m

PAID AT COMPLETION

\$26.0m

MORE, IF REVENUE TARGETS WERE MET

\$0

OF THAT WAS EVER PAID

What it was

North Star Y, Pty Ltd — Braze’s **exclusive reseller in Australia and New Zealand**. Buying it took the market direct.

What the filing records

Braze “reduced the contingent consideration liability ... **to zero as it was determined that the sellers did not satisfy the earn-out qualifications.**”

How to read it

An **earn-out** is money owed only if the business hits agreed targets. This one paid nothing, so the targets were not met — which does not say the acquisition failed, and the filings do not say that either.

— PRICE • BOUNDED, NOT GUESSED

Nobody publishes a price. You can still bound one

~\$283,000

REVENUE ÷ CUSTOMERS, FY2026 — A *BOUND*, NOT A PRICE

An average, not a typical contract. A few very large customers pull it up, and most pay far less. It is the ceiling of what an average customer costs, not what one does.

AND THE MECHANIC UNDERNEATH IT — FROM THEIR DOCUMENTATION

You are billed per data point

“a session start, session end, custom event, or purchase recorded, as well as **any attribute set**” — each counts separately

Engagement is free

Opens, clicks, push tokens and device info are **not** counted. A genuinely customer-friendly boundary

Their own advice is to send less

“**Don’t waste data points. Only update changing data!**” — a platform sold on streaming everything, priced so you send less

Three ways to count a customer

2,609

THE 10-K'S DEFINED
METRIC, 31 JAN 2026

333

CUSTOMERS AT \$500K+ ARR
— UP FROM 202 IN FY2024

178

SELF-PUBLISHED CUSTOMER STORIES

—

INDEPENDENT DETECTION:
NOT ATTEMPTED

EXPANSION IS SLOWING, AND THE ENTERPRISE PREMIUM HAS NEARLY GONE



No second home market

54.9%

UNITED STATES, FY2026

45.1%

INTERNATIONAL

0

OTHER COUNTRIES ABOVE 10% OF REVENUE

The sentence that matters

“Other than the United States, **no other individual country accounted for 10% or more of total revenue** for any of the periods presented.”

Set against the footprint

15 regional clusters and **15 legal entities across 14 territories** — from the status page and the sub-processor disclosure respectively.


So what

A wide, thin footprint on a heavy fixed base. **No single non-US market is large enough that losing it would show** — so outside the US you compete market by market, and so do they.

Well liked, with two consistent complaints

G2 **4.5/5** across 1,702 reviews · Gartner **4.5/5** across 263 ratings · TrustRadius **8.8/10** across 348 reviews

G2'S OWN TAGS, OVER ITS WHOLE REVIEW BASE — PRAISE IN GREY, CRITICISM IN RED, ONE SCALE



Two themes recur across unrelated panels, which is what makes them worth quoting: reporting, and the learning curve. TrustRadius codes reporting as “limited and unintuitive, a sentiment shared by 36% of reviewers”.

What employees say

4.1/5

524 RATINGS

82%

WOULD RECOMMEND

90%

APPROVE OF THE CEO

71%

POSITIVE BUSINESS OUTLOOK

Glassdoor's own summary names the weak spots: **“limited upward mobility and discrepancies in compensation relative to market rates”**.

Glassdoor company page · captured signed-in, 2 September 2026

Not pay — progression

Median employee pay is **\$164,000**. Not a low-paying company, so this is a complaint about the ceiling, not the floor

Still hiring hard

296 open roles across 15 departments — including Bucharest, matching a Romanian entity in the sub-processor list

AND WHAT THEY ARE HIRING FOR — THE FORWARD-LOOKING HALF

Sales **89** · Engineering **57** · Customer Experience **38**. **Go-to-market is 72.0% of the board against 19.6% for engineering and product**. So they are buying new logos rather than building — and that is where they will meet you.

They name four. Buyers weigh eight.

NAMED IN THE 10-K — THEIR CHOICE

Adobe	Salesforce
Iterable	Klaviyo

GARTNER'S BUYER-DERIVED SHORTLIST — NOT THEIR CHOICE

Salesforce	Adobe	Iterable	Oracle
Optimove	Blueshift	MoEngage	CleverTap

THE GAP IS THE FINDING


The brighter names on each side are the ones the other list does not have. Three of their four appear on the buyer list, so this is not a vendor misreading its market — but five vendors buyers compare them against appear nowhere in the 10-K. **So those five are the ones to watch: Braze is not positioning against them, and will not have an answer ready.**

PART II


The product

- 01 One campaign, seven stages, end to end
- 02 Where the limits are — in their words, not ours
- 03 Which channel has no middleman
- 04 What the AI actually is, on five lenses

THEIR TECHNICAL DOCUMENTATION UNLESS MARKED OTHERWISE



One campaign, end to end



Seven stages, and the two picked out are the ones worth arguing about: how fresh the data is going in, and who physically sends the message going out. **Everything between them is competent, so a pitch that attacks the middle will not land — aim at the ends.**

The way in is not the way out

60,000/min

INGEST — /USERS/TRACK, AT 75 OBJECTS EACH

250/min

READ PROFILES BACK BY IDENTIFIER, AT 50 IDS EACH

250,000/hr

BULK EXPORT TO CLOUD STORAGE

STATE THIS FAIRLY — IT IS NOT A LOCK-IN STORY

The asymmetry is real

Writing reaches ~4.5m objects a minute.
Reading profiles back reaches ~12,500 for
a customer who joined after August 2024

But they are different operations


Event objects versus whole profiles.
The ratio shows where priority
sits, not like-for-like throughput

And bulk export exists

Segment export to cloud storage
sits under a 250,000/hour limit. The
sanctioned bulk route is generous

So getting off Braze is a project, not a setting. Price the migration work into any switching conversation, and make the bulk route — not the profile API — the plan.

Their words: three of four are not real-time



Warehouse syncs run “from every 15 minutes to once per month”. Going faster is not self-serve: *“contact your customer success manager or use REST API ingestion.”* **If your customer data lives in a warehouse, fifteen minutes is the floor.**

A merge can fail and still return success

unlimited

ALIASES PER PROFILE

1

ALIAS PER LABEL —
UNIQUE ACROSS THE BASE

5

IDENTIFIER TYPES ACCEPTED ON INGEST

1

IDENTIFIER TYPE FOR WAREHOUSE SEGMENTS

A merge can decline and still report success

“If both profiles have invalid phone numbers, Braze does not merge them
... **The endpoint still returns 202 Accepted with a success message.**”

Reporting splits after a merge

The dashboard attributes a pre-merge send to the surviving profile. Currents — their paid data-export feed — attributes it to the orphaned one. Both are right by their own rules, and they disagree.

So an engineering team should test a merge failure on day one — send two profiles with bad phone numbers and check the profile, not the response code.

Deciding who gets a message

Two systems decide who gets a message, and they sit on two different databases that do not talk to each other. One follows rules you write; the other runs models.

RULE-BASED — MONGODB

Segmentation

Custom events, attributes and most targeting

Segment Extensions

SQL, but served from Snowflake

Global Control Group

Holdouts — people left un-messaged, to measure lift

MODEL-BASED — SNOWFLAKE

Predictive Suite

Churn and event prediction

AI item recommendations

Snowflake-backed

Decisioning Studio

Was OfferFit until June 2025

Braze flags the consequence itself: “Removing data from one system does not automatically remove it from the other.” **So deleting bad data means deleting it twice** — and a deletion request that only clears one side is a compliance problem, not a tidiness one.

Canvas is the strongest thing in the platform

249

FOCUSED DOCUMENTATION PAGES ON CANVAS

10

CANVAS API ENDPOINTS

385

G2 REVIEWS TAGGED 'EASE OF USE'

139

TAGGED 'LEARNING CURVE'

“Canvas makes it easy to build complex, branching lifecycle flows ... Being able to trigger contextual push notifications, in-app messages, and emails from live event streams — **without needing engineering for every small tweak** — is a huge win.”

Enterprise IT manager, G2 · 5/5 · August 2026

So do not attack Canvas. It is the thing their customers like most, and a pitch that calls it weak will be contradicted by the room.

Composing the message

123

FOCUSED DOC PAGES ON LIQUID TEMPLATING

43

ON CONNECTED CONTENT

What it gives you

Liquid templating, Connected Content for live API calls at send time, reusable Content Blocks, and product catalogs for item-level personalisation.

38

ON CONTENT BLOCKS

78

ON CATALOGS

What reviewers say about it

“Liquid personalization and Connected Content also make it straightforward to scale truly dynamic messaging” — and, from the same reviewer, “a challenging learning curve around Liquid syntax” for non-technical users.

So ask who writes the messages. If the answer is marketers rather than engineers, the learning curve is a real cost and it lands on the team you are selling to.

Only two channels have a named middleman

CHANNELS WITH A NAMED DELIVERY SUB-PROCESSOR

Email — three of them

Amazon SES · Bird.com (SparkPost) · Twilio (SendGrid). Plus Mailgun for previewing

SMS / mobile messages — two

Infobip · Twilio

CHANNELS WITH NONE NAMED

Push

In-app

Content Cards

Banners

Webhooks

WhatsApp

LINE

KakaoTalk

Landing pages

Live notif.

Absence is not proof of no intermediary — APNs, FCM and the WhatsApp, LINE and Kakao business APIs may sit outside the definition. **So a buyer cannot establish from this document who touches their messages on eleven of thirteen channels.** That is a question to put to them, not a gap to assume the worst about — and on email, three interchangeable senders is deliberate redundancy.

When the customer replies

71

FOCUSED DOC PAGES ON WEBHOOKS

5 min

CURRENTS EXPORT CADENCE
— OR EVERY 15,000 EVENTS

add-on

CURRENTS IS “AN OPTIONAL BRAZE ADD-ON”

free

ENGAGEMENT EVENTS ARE NOT
BILLED AS DATA POINTS

Inbound and two-way handling is real: webhooks are a first-class documented channel, and engagement tracking — opens, clicks, push receipts — is excluded from data-point billing.

WEBHOOKS IS THE MOST UNDER-SOLD THING IN THE PLATFORM

Fifth-deepest channel in the product

71 focused pages — behind only email, in-app, SMS and push, and sold under no name at all.


Which cuts both ways for us

A real integration surface we should not pretend is absent — and one their own sellers are least likely to lead with.


So the round trip costs more than the outbound leg. Getting the raw stream back is a paid add-on on a five-minute cadence — budget for it when you compare list prices.

Thirteen documented. Ten marketed.

DEEP — FOCUSED DOCUMENTATION PAGES



THIN — REAL, BUT NOT COMPARABLE



Thirteen channels is broad — broader than several specialists on slide 17. But **three are the product and ten are completeness**: email alone outweighs the next three combined. **So compete on the thin ones.** A prospect buying for KakaoTalk or LINE is buying five and eleven pages of documentation, not the platform email buyers get. One line on the drift: five of these ship with no marketing page, and Webhooks — fifth deepest — is one of them.

Wide network, one page deep

150+

TECHNOLOGY PARTNERS — THEIR CLAIM

322

PARTNER DOCUMENTATION PAGES

~810

AVERAGE WORDS PER PARTNER PAGE

~2,708

AVERAGE WORDS PER DEVELOPER-GUIDE PAGE

The network is real and wide, but its documentation is **second-largest by page count and close to thinnest by words per page** — one short entry per partner. **So treat 150+ as a directory, not a depth claim: test the two integrations you actually need before believing it.**

What is bought

Delivery (Twilio, Infobip, SendGrid, SparkPost, SES), warehousing (Snowflake, Databricks, BigQuery), models (Anthropic, OpenAI, Google) — all from the sub-processor list

What is built


The ingestion layer, Canvas, Liquid, the identity model, Currents, and the integration surface itself

Fifteen clusters, seven subsystems — and one exception



one of the nine US clusters is the exception — slide 35

THE SEVEN SUBSYSTEMS, IDENTICAL IN EVERY CLUSTER



The sub-processor disclosure lists AWS regions for the same six territories. **So data residency is available in six places and nowhere else — ask which cluster before the contract.**

What you can measure

36%

OF RECENT TRUSTRADIUS REVIEWERS CALL
REPORTING LIMITED AND UNINTUITIVE

140

G2 REVIEWS TAGGED 'MISSING FEATURES'

27

FOCUSED DOC PAGES ON THE
GLOBAL CONTROL GROUP

2am ET

WHEN YOUR BILLING DASHBOARD
REFRESHES – DAILY, CACHED

What is genuinely there

Holdouts and a Global Control Group for incrementality, campaign and KPI endpoints, and a Snowflake data share that avoids copying data at all.

What reviewers keep hitting

“The out-of-the-box reporting still feels too basic unless you export raw data through paid add-ons like Currents.”

So ask a prospect what they report on today, and who builds it. If the answer involves exporting to a warehouse, they are already paying for the add-on that fixes this.

Bought recently, running on other people’s models

2011-

The platform
segmentation, journeys, Liquid – a decade of it


Jun 2025

The engine, bought
OfferFit, \$303.2m, renamed Decisioning Studio

Now

The models, rented
Anthropic, OpenAI and Google, named in their own disclosure

AND WHAT EACH LAYER MEASURES – FOCUSED DOC PAGES



LENSES 2-5

0 of 135 API endpoints
28 namespaces, none of them AI or decisioning

Reviewers name copywriting
Two G2 reviewers say the AI copy needs “a careful human hand”

So the useful question is one line
“Can I call your AI from my own systems?” Today the documented answer is no — which is what a data-science team needs to hear before they plan around it


PART III

Strategy

- 01 Where seven years of revenue actually went
- 02 What is provisioned that has not been announced
- 03 What would survive a competitor doing the same thing

AUDITED, EXCEPT WHERE MARKED

Decelerating and getting more efficient at once



The expectation was decelerating growth *propped up by* sales spend. The evidence says the opposite: **S&M fell more than twelve points as a share of revenue since FY2023 while growth halved**, and operating cash flow has been positive and rising since FY2024. **So do not plan around them running out of money, or buying growth back.** They can fund a price fight, and the loss line will not stop them.

One instance is not on the same cloud

1 · Two company documents

The sub-processor list names **Amazon and Google**. The 10-K names **Amazon and Rackspace**. **Neither names Microsoft**.

2 · Their own documentation

Every address Braze lists for **US-08 is registered to Microsoft**. Every other instance is Amazon. Checked against ARIN.

3 · Certificate transparency

50 hosts sit on aze region codes matching no AWS identifier, where every other code does — including sdk-us08.

HOW TO SAY THIS — AND HOW NOT TO

A hosting arrangement may sit outside a sub-processor listing for reasons not visible from outside, and saying they failed to disclose would be a legal conclusion this evidence does not support. **So put it to them as a question: which entity operates my instance?**

Three things a competitor cannot copy

Contracted revenue, not pipeline

\$1,033.0m of remaining performance obligation

— 1.40× current revenue, already signed. You cannot displace what is not up for renewal. This is the strongest of the three.

Ten years of streaming plumbing

MongoDB, Snowflake, Kafka and Redis under fifteen clusters, with a decade of incidents to show it holds. Copyable in principle; slow in practice.

Marketer independence from engineering

The thing reviewers actually praise: changing a journey without filing a ticket. That is a workflow habit, and habits are stickier than features.

AND WHAT DOES NOT PROTECT THEM

The AI

Bought for \$303.2m, running on models anyone can rent, with no API of its own. A feature race, not a moat.

The channel roster

Broad and well built — and matched by several of the five specialists buyers shortlist them against but they never name.

PART IV

Open questions

- 01 One hard question, answered properly
- 02 A decade of reliability, measured rather than claimed
- 03 The hypotheses that died in Braze's favour
- 04 What public sources could not answer, and what would close it
- 05 What to remember

WHERE THE EVIDENCE RUNS OUT, THIS SAYS SO

“Is it real-time?” — answered properly

ASK

Which path?

The answer differs by
a factor of hundreds

SDK / API

Near-real-time

Their words, with ‘async
processing’ attached

WAREHOUSE

15 minutes, floor

“Not real-time”,
three times over

FASTER?

Not self-serve

“Contact your customer
success manager”

EXPORT

5 minutes

And Currents is a paid add-on

Why this is the question that matters

Most enterprise buyers keep customer
data in a warehouse. For them the
honest answer is fifteen minutes, and
it is not a setting they can change.

Why it is not a gotcha

Braze publishes this table themselves, and
the SDK path genuinely is near-real-time. One
word, two architectures — not a false claim.

What to actually ask them

“Which ingestion path will *my* data take, and
what is the latency on it?” **The answer is in
their own documentation before the meeting.**

A decade of reliability, measured



79 min

MEDIAN INCIDENT DURATION

29.6%

OF NON-MAINTENANCE INCIDENTS
WERE MAJOR OR CRITICAL

63 v 27

DASHBOARD INCIDENTS
VS OUTBOUND MESSAGING

97.3%

OF 845 UNSOLICITED PUBLIC
ISSUES CLOSED – MEDIAN 11 DAYS

9 of 9

SDK REPOS SHIPPED WITHIN
13 DAYS OF CAPTURE

The caveat belongs on the slide, not only in the notes: 2026 stands at 35 through August, roughly **double the 2025 monthly rate**. One year is not a trend, and none is claimed. And **never compare this to a vendor who publishes nothing**.

Three things we expected to find, and did not

H1 · “Growth is decelerating while sales spend holds”

Killed. Growth did decelerate — and sales and marketing fell **from 56.7% of revenue to 44.3%** over the same period, its lowest in the seven-year series. Braze is decelerating *and* getting more efficient. Operating cash flow turned positive in FY2024 and has grown every year since.

H4 · “Some supported platforms are effectively unmaintained”

Killed. All nine **SDK repositories** that publish releases had shipped within 13 days of capture. The one archived repo, Unreal, has had its documentation removed too — so the marketing and the maintenance record agree. One soft spot: `braze-roku-sdk`, 181 days idle with 38 Roku doc pages still live.

H9 · “Incident rate has risen with scale”

Killed. Incidents peaked at 60 in 2023 and fell to **27 in 2025**, the quietest full year on the status page, while revenue grew 7.7×. Caveat attached: 2026 runs at about double the 2025 monthly rate.

Slide 2 promised this. Ten hypotheses were written before any source was read; four came back wrong. A set that all confirmed would have meant they were written to confirm.

What we could not answer

Does satisfaction fall with customer size?

The one hypothesis we could not test. **All three review sites** **paywall that breakdown**. *Closed by:* paid panel access

What else is in the certificate estate?

The host list is **partial** — 833 hosts through a rate-limited fallback. Everything found stands; **nothing is claimed about what is absent**. *Closed by:* an API token

Can the customer roster be checked independently?

The only roster outside the 10-K is **178 self-published stories** — marketing, not a sample. Independent detection was **not attempted**. *Closed by:* tag crawls, or job ads naming Braze in the stack

Actual pricing, and why the export limit was cut

No vendor publishes a rate card, and the 10× cut to the profile-lookup limit is documented without explanation. *Closed by:* a procurement award or a customer contract

Four questions worth the next week

1 · Ask about US-08

The highest-value single answer available. **Route:** the DPA schedule a customer already receives, or ask them.

2 · Buy one panel's segment data

It closes the only unresolved hypothesis, and it is the question your sales team asks first. **Route:** paid G2 or Gartner access.

3 · Detect the customer base independently

The only roster outside the audited count is 178 stories Braze chose to publish. **Route:** tag crawls or job ads naming the stack. Not cheap.

4 · Re-run this in ninety days

Every number here is reproducible by script. Watch the gross-margin decline and the 2026 incident rate.

So the first two are worth a week and would change what we say; the last is worth an hour a quarter. The full backlog is in docs/QUESTIONS.md.

One thing to remember

Their strongest asset is a contract, not a capability.

A billion dollars of next year is already signed, so most of the base is not winnable this year. They are growing slower and spending less to do it, so they can fund a fight. And the capability they paid \$303.2m for runs on models anyone can rent, with no API of its own. **So compete on the renewal date and on the things they bought rather than built — not on the demo.**

1,352

DOCUMENTATION PAGES READ

451

INCIDENTS

845

PUBLIC ISSUES

17

SUB-PROCESSORS

7

AUDITED YEARS