



COMPETITOR ANALYSIS

A customer engagement platform — and a public company, which changes what can be known.

THEIR DOCUMENTATION

1,352 pages

SEC FILINGS

737 documents

AUDITED FINANCIALS

FY2019 → now

PUBLIC CODE

137 repos

STATUS PAGE

451 incidents

SUB-PROCESSORS

17 outside suppliers

Toma Pirtskhelani · Product Manager

Public sources only · September 2026

Four parts, three rules

PART I

The company

who owns it • what it costs • who buys it

PART II

The product

seven stages of one campaign

PART III

Strategy

where the money goes • what protects them

PART IV

Open questions

what public sources cannot answer

AND THREE RULES THAT APPLY TO ALL FOUR

Every claim is graded

Marked by how strong its source is. The bar at the foot of each slide shows where we are.

Marketing is labelled as marketing

Their own words appear
— but never as evidence.

Gaps become a backlog

Anything needing a non-public source goes on the open-questions list, and grows as you add to it.

Every claim takes its weakest source

STRONG

Audited filings, and their own technical documentation

Signed off by auditors and filed under legal penalty — or behaviour the platform must actually support

MEDIUM

Independent records

Customer reviews, infrastructure traces, incident history, real transacted prices

WEAK

Marketing, or disputed

Their own claims, or sources that contradict each other

Five grades are kept in the evidence record. Three reach the slides, mapped in docs/EVIDENCE- GRADES .md. **A claim takes the grade of its weakest supporting source, never its best.**

PART I

The company

- 01 Who they are, and who controls them now
- 02 What they bought, and what it cost
- 03 What a customer pays — bounded, not guessed
- 04 Who buys it, and who buyers compare them against

AUDITED FILINGS UNLESS MARKED OTHERWISE

Four documents a company cannot write freely

The 10-K

Their **audited annual report to the US regulator**, filed once a year. Signed by the chief executive and the finance chief, audited by Ernst & Young, and wrong at legal risk. **Seven years of them.**

The proxy statement

Filed before the annual shareholder meeting. It has to name **every executive officer and director**, what each is paid, and who owns the company. It is the only place people appear.

The sub-processor disclosure

A public list of **every outside supplier that touches customer data**. Data-protection law obliges them to keep it complete, so it names middlemen no marketing page would.

The status page

A live record of **every outage since 2016**, written during the outage. Nobody writes a status page to look good, which is exactly why it is worth reading.

None of these is marketing. Three are filed under legal penalty and the fourth is written under pressure — which is why, wherever this deck and Braze's website disagree, the documents win.

If you remember five things

Braze grades its own data delays, and three of four are slow

Their documentation labels three of the four ways data gets in “**not real-time**”. It is their table, not our characterisation.

The AI decisioning engine was bought, not built

OfferFit, acquired June 2025 for **\$303.2m** and renamed AI Decisioning Studio. The models come from Anthropic, OpenAI and Google.

One instance is not on the same cloud as the others

The addresses Braze tells you to allowlist — to permit through your firewall — for **US-08** belong to Microsoft. Every other instance’s belong to Amazon.

They are slowing down and getting more efficient at the same time

Growth has halved since FY2023, and the share of revenue spent on sales and marketing fell further. **Three straight years of positive cash flow.**

Braze names four competitors. Buyers compare them against eight

And the five extra names on the buyer’s list are specialists Braze never mentions.



Their biggest claim, in their most careful document

“Our platform empowers **real-time engagement** between brands and their customers ... made possible by our proprietary, enterprise-grade **stream processing architecture** ... We have designed it to **listen like a human would**, process new information in context, and **react instantaneously**.”

Braze 10-K, Item 1 · filed 25 March 2026

The claim

A single platform, fed by streaming first-party data, reacting in the moment across every channel

Why it is quoted here

This is the company’s own narrative in a filed document — the strongest form of a marketing claim, not a technical one

What to hold on to

“Real-time” is doing a lot of work in that sentence. Slide 21 shows their own table grading it

Two of three cofounders still run it



Bill Magnuson

CHAIRMAN • CEO • PRESIDENT • COFOUNDER

CTO from July 2011, chief executive since January 2017, chairman of the board — and President too since June 2025, when the previous President resigned and the role was not refilled. One person, four titles.

Cofounder Jon Hyman is still CTO, an officer since 2011.

AND WHO HOLDS THE SEATS

Seven seats, six independent

Magnuson is the exception, as President and CEO. **Phillip Fernandez is Lead Independent Director** — the role a board creates when its chair is not.

The chair is the chief executive

The proxy defends combining them rather than glossing it: separating them “has the potential to give rise to divided leadership”.

And the board is classified

Three staggered classes, so only about **a third stands in any year**. The super-voting stock is gone; this is not.

The votes, after January

Class B converted **30 January 2026**; super-voting ended. Largest holder **6.0%**, the CEO **4.9%**, the Battery partner **5.1%** — **no blocking position left**.

7.7× bigger, a billion already contracted



\$456.8m

NET IPO PROCEEDS, NOV 2021

7.7×

REVENUE GROWTH, FY2020 TO FY2026

\$124.3m

CASH AT FY2026 YEAR END

\$1,033.0m

CONTRACTED, NOT YET RECOGNISED

THE CAVEAT THAT TRAVELS WITH EVERY NUMBER IN THIS PART

Braze disclosed a **material weakness in internal control over financial reporting** at 31 January 2026 — *a flaw in the checking, not an error in the numbers*. Ineffective IT controls over **user access and program change management** on the systems that produce these figures. **And, in the same breath:** it “did not result in any identified misstatements”, nothing was restated, and Ernst & Young attested.



They bought their AI, and the filing says so

“the Company completed the acquisition of **OfferFit, Inc.** (‘OfferFit’) **which is now known as AI Decisioning Studio** for total consideration of **\$303.2 million.**”

Ernst & Young, critical audit matter · Braze 10-K, 25 March 2026

\$303.2m

TOTAL CONSIDERATION

77%

WAS GOODWILL — PRICE
NOT TIED TO ANY ASSET

\$56.7m

DEVELOPED TECHNOLOGY,
AMORTISED TO COST OF REVENUE

2 Jun 2025

CLOSED

WHAT THE PURCHASE-PRICE ALLOCATION TELLS YOU

Only **\$66.6m of the \$303.2m was identifiable**, and \$56.7m of that is the software itself — so Braze bought a capability and the people who built it, not a revenue stream. The trademark was valued at **\$0.9m**: nobody expected the OfferFit name to survive, and it did not. And the technology amortises **to cost of revenue over six years**, which is why an AI acquisition lands in gross margin — and why the drag runs to 2031 rather than washing out next year.

They bought a market, and the earn-out paid nothing

\$26.8m

PAID AT COMPLETION

\$26.0m

MORE, IF REVENUE TARGETS WERE MET

\$0

OF THAT WAS EVER PAID

What it was

North Star Y, Pty Ltd — Braze’s **exclusive reseller in Australia and New Zealand**. Buying it took the market direct.

What the filing records

Braze “reduced the contingent consideration liability ... **to zero as it was determined that the sellers did not satisfy the earn-out qualifications.**”

How to read it

An **earn-out** is money owed only if the business hits agreed targets. This one paid nothing, so the targets were not met — which does not say the acquisition failed, and the filings do not say that either.

— PRICE • BOUNDED, NOT GUESSED

Nobody publishes a price. You can still bound one

~\$283,000

REVENUE ÷ CUSTOMERS, FY2026 — A *BOUND*, NOT A PRICE

An average, not a typical contract. A few very large customers pull it up, and most pay far less. It is the ceiling of what an average customer costs, not what one does.

AND THE MECHANIC UNDERNEATH IT — FROM THEIR DOCUMENTATION

You are billed per data point

“a session start, session end, custom event, or purchase recorded, as well as **any attribute set** on an end user profile” — each one counts separately

Engagement is free

Push tokens, device info, email opens and push clicks are **not** counted. That is a genuinely customer-friendly boundary

Their own advice is to send less

“**Don’t waste data points. Only update changing data!**” — a platform sold on streaming everything, priced so that you send less of it

How many customers depends on who counts

2,609

THE 10-K'S DEFINED METRIC, 31 JAN 2026

333

CUSTOMERS AT \$500K+ ARR
— UP FROM 202 IN FY2024

178

SELF-PUBLISHED CUSTOMER STORIES

—

INDEPENDENT DETECTION:
NOT ATTEMPTED

EXPANSION IS SLOWING, AND THE ENTERPRISE PREMIUM HAS NEARLY GONE



No second home market

54.9%

UNITED STATES, FY2026

45.1%

INTERNATIONAL

0

OTHER COUNTRIES ABOVE 10% OF REVENUE

The sentence that matters

“Other than the United States, **no other individual country accounted for 10% or more of total revenue** for any of the periods presented.”

Set against the footprint

15 regional clusters and **15 legal entities across 14 territories** —
from the status page and the sub-processor disclosure respectively.


So what

Forty-five per cent of revenue is international and none of it concentrates.
That is a wide, thin footprint carrying a heavy infrastructure and legal base.

Well liked, with two consistent complaints

G2 **4.5/5** across 1,702 reviews · Gartner **4.5/5** across 263 ratings · TrustRadius **8.8/10** across 348 reviews

G2'S OWN TAGS, OVER ITS WHOLE REVIEW BASE — PRAISE IN GREY, CRITICISM IN RED, ONE SCALE



Two themes recur across unrelated panels, which is what makes them worth quoting: reporting, and the learning curve. TrustRadius codes reporting as “limited and unintuitive, a sentiment shared by 36% of reviewers”.

Well rated. The complaints are about progression

4.1/5

524 RATINGS

82%

WOULD RECOMMEND

90%

APPROVE OF THE CEO

71%

POSITIVE BUSINESS OUTLOOK

Glassdoor's own summary names the weak spots: **“limited upward mobility and discrepancies in compensation relative to market rates”**.

Glassdoor company page · captured signed-in, 2 September 2026

Not pay — progression

The proxy puts median employee pay at **\$164,000**.
This is not a low-paying company, so “limited upward mobility” is a complaint about the ceiling, not the floor

Still hiring hard

296 open roles across 15 hiring departments — including Bucharest, which matches a Romanian entity in the sub-processor list

AND WHAT THEY ARE HIRING FOR — THE FORWARD-LOOKING HALF

Sales **89** · Engineering **57** · Customer Experience **38** · Marketing 24 · Solutions Consulting 23 · Partnerships 16. **Go-to-market is 72.0% of the board against 19.6% for engineering and product — roughly 3.7 to one.** A requisition is an intention, not a person, and this is not a spend ratio — but it points the same way as slide 13's retention figures.

They name four. Buyers weigh eight.

NAMED IN THE 10-K — THEIR CHOICE

Adobe	Salesforce
Iterable	Klaviyo

GARTNER'S BUYER-DERIVED SHORTLIST — NOT THEIR CHOICE

Salesforce	Adobe	Iterable	Oracle
Optimove	Blueshift	MoEngage	CleverTap

THE GAP IS THE FINDING


The brighter names on each side are the ones the other list does not have. Three of their four appear on the buyer list, so this is not a vendor misreading its market — but **five vendors buyers compare them against appear nowhere in the 10-K**, and they are mostly mobile-engagement specialists rather than the suites Braze positions against.

PART II


The product

- 01 One campaign, seven stages, end to end
- 02 Where the limits are — in their words, not ours
- 03 Which channel has no middleman
- 04 What the AI actually is, on five lenses

THEIR TECHNICAL DOCUMENTATION UNLESS MARKED OTHERWISE



One campaign, end to end



Seven stages. **The two picked out are where the evidence turned out to be interesting** — how fresh the data is at stage one, and who physically sends the message at stage six. Everything in between is competent and largely as advertised, and I will say so as we go.

The way in is not the way out

60,000/min

INGEST — /USERS/TRACK, AT 75 OBJECTS EACH

250/min

READ PROFILES BACK BY
IDENTIFIER, AT 50 IDS EACH

250,000/hr

BULK EXPORT TO CLOUD STORAGE

STATE THIS FAIRLY — IT IS NOT A LOCK-IN STORY

The asymmetry is real

Writing reaches ~4.5m objects a minute. Reading profiles back by identifier reaches ~12,500 a minute for a customer who joined after August 2024


But they are different operations

Writing event objects versus reading whole profiles. The ratio shows where priority sits, not like-for-like throughput

And bulk export exists

Segment export to cloud storage sits under the default 250,000/hour limit. The sanctioned bulk route is generous

Their words: three of four are not real-time



Warehouse syncs run “from every 15 minutes to once per month”. Going faster is not self-serve: *“contact your customer success manager or use REST API ingestion.”* **If your customer data lives in a warehouse, fifteen minutes is the floor.**

A merge can fail and still return success

unlimited

ALIASES PER PROFILE

1

ALIAS PER LABEL —
UNIQUE ACROSS THE BASE

5

IDENTIFIER TYPES
ACCEPTED ON INGEST

1

IDENTIFIER TYPE FOR
WAREHOUSE SEGMENTS

A merge can decline and still report success

“If both profiles have invalid phone numbers, Braze does not merge them ... **The endpoint still returns 202 Accepted with a success message**, so the HTTP response does not indicate that the merge was skipped.”

Reporting splits after a merge

Dashboard summaries attribute a pre-merge send to the surviving profile. Currents — their paid data-export feed — along with Query Builder and Messaging History, attribute it to the orphaned one. Both are right by their own rules, and they disagree.

Two engines, two databases, one bought

Two systems decide who gets a message, and they sit on two different databases that do not talk to each other. One follows rules you write; the other runs models.

RULE-BASED — MONGODB

Segmentation

Custom events, attributes, purchases and most targeting

Segment Extensions

SQL, but served from Snowflake

Global Control Group

Holdouts — people deliberately left un-messaged, to measure lift

MODEL-BASED — SNOWFLAKE

Predictive Suite

Churn and event prediction — 7 focused doc pages

AI item recommendations

Snowflake-backed

Decisioning Studio

Was OfferFit until June 2025

The split has a customer-visible consequence Braze flags itself: “Removing data from one system does not automatically remove it from the other.” Deleting bad event data means doing it in MongoDB, separately from anything Snowflake-backed.

Canvas is the strongest thing in the platform

249

FOCUSED DOCUMENTATION
PAGES ON CANVAS

10

CANVAS API ENDPOINTS

385

G2 REVIEWS TAGGED 'EASE OF USE'

139

TAGGED 'LEARNING CURVE'

“Canvas makes it easy to build complex, branching lifecycle flows ... Being able to trigger contextual push notifications, in-app messages, and emails from live event streams — **without needing engineering for every small tweak** — is a huge win.”

Enterprise IT manager, G2 • 5/5 • August 2026

Liquid does the work, and draws the complaints

123

FOCUSED DOC PAGES
ON LIQUID TEMPLATING

43

ON CONNECTED CONTENT

38

ON CONTENT BLOCKS

78

ON CATALOGS

What it gives you

Liquid templating, Connected Content for live API calls at send time, reusable Content Blocks, and product catalogs for item-level personalisation.

What reviewers say about it

“Liquid personalization and Connected Content also make it straightforward to scale truly dynamic messaging” — and, from the same reviewer, “a challenging learning curve around Liquid syntax” for non-technical users.

Only two channels have a named middleman

CHANNELS WITH A NAMED DELIVERY SUB-PROCESSOR

Email — three of them

Amazon SES · Bird.com (SparkPost) · Twilio (SendGrid). Plus Mailgun for previewing

SMS / mobile messages — two

Infobip · Twilio

CHANNELS WITH NONE NAMED

Push

In-app

Content Cards

Banners

Webhooks

WhatsApp

LINE

KakaoTalk

Landing pages

Live notif.

Say what the disclosure names, not what exists. Absence here is not proof of no intermediary — APNs, FCM and the WhatsApp/LINE/Kakao business APIs may sit outside the definition of a sub-processor. What is certain is that **email has three interchangeable senders**, which is deliberate redundancy on their highest-volume channel.

The response comes back, billed differently

71

FOCUSED DOC PAGES ON WEBHOOKS

5 min

CURRENTS EXPORT CADENCE
— OR EVERY 15,000 EVENTS

add-on

CURRENTS IS “AN OPTIONAL BRAZE ADD-ON”

free

ENGAGEMENT EVENTS ARE NOT
BILLED AS DATA POINTS

Inbound and two-way handling is real: webhooks are a first-class documented channel, and engagement tracking — opens, clicks, push receipts — is explicitly excluded from data-point billing. **But getting the raw event stream back out is a paid add-on that exports on a five-minute cadence**, which is the mechanism behind the reporting complaint on slide 15.

WEBHOOKS IS THE MOST UNDER-SOLD THING IN THE PLATFORM

Fifth-deepest channel in the product

71 focused documentation pages puts Webhooks above WhatsApp, Content Cards, Banners and every channel except email, in-app, SMS and push.

And it has no marketing page at all

One of the five documented channels sold under no name. A prospect judging Braze from its product pages would not know this exists.

Which cuts both ways for us

It is a real integration surface we should not pretend is absent — and it is one Braze’s own sellers are least likely to lead with.

Thirteen documented. Ten marketed.

DOCUMENTED, WITH NO MARKETING PAGE

Banners

Transactional email

Landing pages

Live notifications

Webhooks

MARKETED, BUT MISSING FROM THEIR OWN DOCS CHANNEL INDEX

KakaoTalk

Four KakaoTalk documentation pages exist — setup, message creation, click tracking, reporting — and a product page at `/product/kakaotalk-messenger`. The channels index mentions it **zero** times.

Landing pages is the one to notice: eleven user-guide pages plus three partner pages, a substantial documented capability sold under no name at all.

Wide network, one page deep

150+

TECHNOLOGY PARTNERS — THEIR CLAIM

322

PARTNER DOCUMENTATION PAGES

~810

AVERAGE WORDS PER PARTNER PAGE

~2,708

AVERAGE WORDS PER DEVELOPER-GUIDE PAGE

The integration network is real and wide. But the documentation shape says something about its depth: **the partner section is the second-largest in the corpus by page count and close to the thinnest by words per page** — a lot of short entries, roughly one per partner. The developer guide, a third the size, carries more than three times the words per page.

What is bought

Delivery (Twilio, Infobip, SendGrid, SparkPost, SES), warehousing (Snowflake, Databricks, BigQuery), models (Anthropic, OpenAI, Google) — all from the sub-processor list

What is built

The ingestion layer, Canvas, Liquid, the identity model, Currents, and the integration surface itself

Fifteen clusters, seven subsystems — and one exception



one of the nine US clusters is the exception — slide 35

THE SEVEN SUBSYSTEMS, IDENTICAL IN EVERY CLUSTER

Dashboard

SDK Data Collection

Data Processing

REST APIs

Outbound Messaging

Currents

Cloud Data-Ingestion

Corroborated independently: the sub-processor disclosure lists AWS regions for the same six territories. **There is no US 09.**

Reporting is the most consistent complaint

36%

OF RECENT TRUSTRADIUS
REVIEWERS CALL REPORTING
LIMITED AND UNINTUITIVE

140

G2 REVIEWS TAGGED
'MISSING FEATURES'

27

FOCUSED DOC PAGES ON THE
GLOBAL CONTROL GROUP

2am ET

WHEN YOUR BILLING DASHBOARD
REFRESHES — DAILY, CACHED

What is genuinely there


Holdouts and a Global Control Group for incrementality, campaign and Canvas analytics endpoints, KPI endpoints, and a Snowflake data share that avoids copying data at all.

What reviewers keep hitting

“The out-of-the-box reporting still feels too basic unless you export raw data through paid add-ons like Currents” — and, separately, a wish for “more robust native reporting on long-term retention cohorts”.

Bought recently, running on other people’s models

LENS 1 — FOCUSED DOCUMENTATION PAGES



LENSES 2-5

0 of 135 API endpoints

28 namespaces. None of them AI, prediction, agent or decisioning

3 external model suppliers

Anthropic, OpenAI and Google — named in their own sub-processor disclosure

The engine was bought

Decisioning Studio was OfferFit until June 2025. \$303.2m

Reviewers name copywriting

Two G2 reviewers say the AI copy needs “a careful human hand”. None mentions decisioning


PART III

Strategy

- 01 Where seven years of revenue actually went
- 02 What is provisioned that has not been announced
- 03 What would survive a competitor doing the same thing

AUDITED, EXCEPT WHERE MARKED

Decelerating and getting more efficient at once



This killed the hypothesis it was built to test. The expectation was decelerating growth *propped up by* sales spend. **The evidence says the opposite: S&M has fallen more than twelve points as a share of revenue since FY2023, to its lowest in the series, while growth halved.** Operating cash flow turned positive in FY2024 and has grown every year since — to \$71.4m.

One instance is not on the same cloud

1 · Two company documents

The sub-processor list (1 June 2026) names **Amazon and Google** as hosting providers. The 10-K names “**Amazon Web Services and Rackspace**”. **Neither names Microsoft.**

2 · Their own documentation

Braze publishes the IPs you must allowlist, per instance. Every address listed for **US-08 is registered to Microsoft Corporation**. Every other instance — US-10, AU-01, ID-01, JP-01, KR-01 — is Amazon. Checked against ARIN.

3 · Certificate transparency

50 hosts sit on region codes p-aze-us, s-aze-us, d-aze-us — matching no AWS region identifier, where every other code does. They include sdk-us08 and subcenter-08.

HOW TO SAY THIS — AND HOW NOT TO

State the three observations and stop. A hosting arrangement may sit outside a sub-processor listing for reasons not visible from outside. Saying they failed to disclose would be a legal conclusion this evidence does not support. *The question is for them, not a conclusion for us.*

Three things a competitor cannot copy

Contracted revenue, not pipeline

\$1,033.0m of remaining performance obligation — 1.40× current revenue, already signed. You cannot displace what is not up for renewal. This is the strongest of the three.

Ten years of streaming plumbing

MongoDB, Snowflake, Kafka and Redis under fifteen regional clusters, with a decade of incident history to show it holds. Copyable in principle; expensive and slow in practice.

Marketer independence from engineering

The thing reviewers actually praise: changing a journey without filing a ticket. That is a workflow habit, and habits are stickier than features.

AND WHAT DOES NOT PROTECT THEM

The AI

Bought for \$303.2m, running on Anthropic, OpenAI and Google models, with no API surface of its own. Anyone can buy the same models. This is a feature race, not a moat.

The channel roster

Broad and well built — and matched by several of the five specialists buyers shortlist them against but they never name.

PART IV

Open questions

- 01 One hard question, answered properly
- 02 A decade of reliability, measured rather than claimed
- 03 The hypotheses that died in Braze's favour
- 04 What public sources could not answer, and what would close it
- 05 What to remember

WHERE THE EVIDENCE RUNS OUT, THIS SAYS SO

“Is it real-time?” — answered properly

ASK

Which path?

The answer differs by
a factor of hundreds

SDK / API

Near-real-time

Their words, with ‘async
processing’ attached

WAREHOUSE

15 minutes, floor

“Not real-time”,
three times over

FASTER?

Not self-serve

“Contact your customer
success manager”

EXPORT

5 minutes

And Currents is a paid add-on

Why this is the question that matters

Most enterprise buyers’ customer
data lives in a warehouse. For them
the honest answer is fifteen minutes,
and it is not a setting they can change.

Why it is not a gotcha

Braze publishes this themselves, in a
comparison table, and the SDK path
genuinely is near-real-time. The word covers
two architectures — it is not a false claim.

What to actually ask them

“Which ingestion path will *my* data take, and
what is the latency on that path?” The answer
is in their documentation before the meeting.

451 incidents, and the rate is falling



79 min

MEDIAN INCIDENT DURATION

29.6%

OF NON-MAINTENANCE INCIDENTS
WERE MAJOR OR CRITICAL

63 v 27

DASHBOARD INCIDENTS
VS OUTBOUND MESSAGING

97.3%

OF 845 UNSOLICITED PUBLIC
ISSUES CLOSED – MEDIAN 11 DAYS

9 of 9

SDK REPOS SHIPPED WITHIN
13 DAYS OF CAPTURE

The caveat belongs on the slide, not only in the notes: 2026 stands at 35 through August, roughly **double the 2025 monthly rate**. One year is not a trend, and none is claimed. And **never compare this to a vendor who publishes nothing**.

Three things we expected to find, and did not

H1 · “Growth is decelerating while sales spend holds”

Killed. Growth did decelerate — and sales and marketing fell **from 56.7% of revenue to 44.3%** over the same period, its lowest in the seven-year series. Braze is decelerating *and* getting more efficient. Operating cash flow turned positive in FY2024 and has grown every year since.

H4 · “Some supported platforms are effectively unmaintained”

Killed. All nine **SDK repositories** that publish releases had shipped within 13 days of capture. The one archived repo, Unreal, has had its documentation removed too — so the marketing and the maintenance record agree. One soft spot: `braze-roku-sdk`, 181 days idle with 38 Roku doc pages still live.

H9 · “Incident rate has risen with scale”

Killed. Incidents peaked at 60 in 2023 and fell to **27 in 2025**, the quietest full year on the status page — over a period when revenue grew 7.7×. With its caveat attached: 2026 runs at roughly double the 2025 monthly rate through August, and one year is not a trend.

WHY THIS SLIDE EXISTS

Ten hypotheses were written **before any source was read**, so they could be graded honestly afterwards. Four were killed, three of them favourably; one could not be tested at all. **A set of hypotheses that all confirmed would have meant they were written to confirm.** Everything critical in this deck should be read against the fact that these three were looked for and were not there.

Four gaps, and what would close them

Does satisfaction fall with customer size?

The one hypothesis that could not be tested at all. **All three review sites paywall exactly that breakdown**, and only 7 of 860 coded records carry a segment. *Closed by:* paid panel access, or an investor-day disclosure

What else is in the certificate estate?

The host list is **partial** — 833 hosts through a rate-limited fallback after crt.sh returned errors all day. Everything found stands; **nothing is claimed about what is absent**. *Closed by:* an API token

Can the customer roster be checked independently?

The only roster outside the 10-K is **178 self-published customer stories**, which is a marketing selection and not a sample. Independent detection was **not attempted**, and that is recorded as a gap rather than a zero. *Closed by:* tag crawls, CT records naming customer subdomains, or job ads naming Braze in the stack

Actual pricing, and why the export limit was cut

No vendor publishes a rate card, and the 10× cut to the profile-lookup limit on 22 August 2024 is documented without explanation. *Closed by:* a procurement award, a customer contract, a changelog entry

Four questions worth the next week

1 · Ask about US-08

The highest-value single answer in the backlog. It decides whether slide 35 is an infrastructure observation or a procurement question.

Route: the DPA schedule a customer receives, or ask them directly.

2 · Buy one panel's segment data

It closes the only hypothesis that went unresolved, and satisfaction-by-size is the question your sales team will ask first. **Route:** paid G2 or Gartner access.

3 · Detect the customer base independently

The only roster outside the audited count is 178 stories Braze chose to publish. Tag crawls, certificate records and job ads naming the stack would give a roster nobody curated. **Route:** a crawl, and it is not cheap.

4 · Re-run this in ninety days

Every number here is reproducible by script. The two to watch: whether the gross-margin decline continues, and whether the 2026 incident rate settles or climbs.

The full backlog — eight open questions, each with what would close it — is in docs/QUESTIONS.md, and every finding in this deck resolves to a row in docs/FACTS.md with a source path and a capture date.

One thing to remember

Their documentation is more honest than anyone's marketing — including their own.

Every hard limit in this deck came from a page Braze wrote for its own engineers: three of four ingestion paths labelled “not real-time”, a fifteen-minute warehouse floor, an export limit cut tenfold on a dated boundary, and a merge that returns success when it has quietly declined. None of it is hidden. All of it is unread.

1,352

DOCUMENTATION PAGES READ

451

INCIDENTS

845

PUBLIC ISSUES

17

SUB-PROCESSORS

7

AUDITED YEARS