

Braze

COMPETITOR ANALYSIS

A customer engagement platform — and a public company, which changes what can be known.

THEIR DOCUMENTATION

1,352 pages

SEC FILINGS

737 documents

AUDITED FINANCIALS

FY2019 → now

PUBLIC CODE

137 repos

STATUS PAGE

451 incidents

SUB-PROCESSORS

17 disclosed

41 slides · every figure traceable to a source path

Public sources only · captured 1–2 September 2026

How we approach it

PART I

The company

who owns it • what it costs • who buys it

PART II

The product

seven stages of one campaign

PART III

Strategy

where the money goes • what protects them

PART IV

Open questions

what public sources cannot answer

AND THREE RULES THAT APPLY TO ALL FOUR

Every claim is graded

Marked by how strong its source is. The bar at the foot of each slide shows where we are.

Marketing is labelled as marketing

Their own words appear
— but never as evidence.

Gaps become a backlog

Anything needing a non-public source goes on the open-questions list, and grows as you add to it.

Three grades, one bar at the foot of every slide

STRONG

Audited filings, and their own technical documentation

Signed off by auditors and filed under legal penalty — or behaviour the platform must actually support

MEDIUM

Independent records

Customer reviews, infrastructure traces, incident history, real transacted prices

WEAK

Marketing, or disputed

Their own claims, or sources that contradict each other

Five grades are kept in the evidence record. Three reach the slides, mapped in docs/EVIDENCE- GRADES .md. **A claim takes the grade of its weakest supporting source, never its best.**

PART I

The company

- 01 Who they are, and who controls them now
- 02 What they bought, and what it cost
- 03 What a customer pays — bounded, not guessed
- 04 Who buys it, and who buyers compare them against

AUDITED FILINGS UNLESS MARKED OTHERWISE

If you remember five things

Their own documentation is the best source against them

Three of four ingestion paths are labelled “not real-time” by Braze. The profile-export limit was cut 10× for customers joining after August 2024. Both are published, in tables.

The AI decisioning engine was bought, not built

OfferFit, acquired June 2025 for \$303.2m and renamed AI Decisioning Studio. The models come from Anthropic, OpenAI and Google — named in their own sub-processor disclosure.

One instance is not on the same cloud as the others

The IPs Braze tells you to allowlist for US-08 are registered to Microsoft. Every other instance’s are Amazon. Their sub-processor list names only Amazon and Google.

Decelerating and getting more efficient at once

Growth halved from 49.3% to 24.4% since FY2023 while sales and marketing fell from 56.7% of revenue to 44.3%. Operating cash flow has been positive three years running.

Buyers weigh them against twice as many vendors as they name

The 10-K names four competitors. Gartner’s buyer-derived shortlist has eight — and the five extra are specialists Braze never mentions.

Who they say they are

“Our **real-time** engagement between brands and their customers ... made **stream** processing architecture ... We have designed it to **listen like a human would**, process new information in context, and **react** **instantaneously**.”

Braze 10-K, Item 1 · filed 25 March 2026

The claim

A single platform, fed by streaming first-party data, reacting in the moment across every channel

Why it is quoted here

This is the company’s own narrative in a filed document — the strongest form of a marketing claim, not a technical one

What to hold on to

“Real-time” is doing a lot of work in that sentence. Slide 20 shows their own table grading it

Fifteen years, and a governance change nine months old



THE PART THAT IS EASY TO MISS

On **30 January 2026** the Class B stock converted to Class A. In the 10-K's words: *"our executive officers and early investors no longer hold super-voting rights. Consequently, our voting power is now more widely distributed among our public stockholders."*

Seven years of audited revenue



\$456.8m

NET IPO PROCEEDS, NOV 2021

7.7×

REVENUE GROWTH, FY2020 TO FY2026

\$124.3m

CASH AT FY2026 YEAR END

\$1,033.0m

CONTRACTED, NOT YET RECOGNISED

They bought their AI, and the filing says so

“the Company completed the acquisition of **OfferFit, Inc.** (‘OfferFit’) **which is now known as AI Decisioning Studio** for total consideration of **\$303.2 million** .”

Ernst & Young, critical audit matter · Braze 10-K, 25 March 2026

\$303.2m

TOTAL CONSIDERATION

77%

OF THE PRICE WAS GOODWILL

\$56.7m

DEVELOPED TECHNOLOGY,
AMORTISED TO COST OF REVENUE

2 Jun 2025

CLOSED

They bought a market, and the earn-out paid nothing

\$26.8m

TOTAL CONSIDERATION

\$26.0m

EARN-OUT AVAILABLE

\$0

EARN-OUT PAID

What it was

North Star Y, Pty Ltd — Braze’s **exclusive reseller in Australia and New Zealand**. Buying it took the market direct.

What the filing records

Braze “reduced the contingent consideration liability ... **to zero as it was determined that the sellers did not satisfy the earn-out qualifications.**”

How to read it

An earn-out that does not vest means the revenue targets written into the deal were not met. It does not say the acquisition failed — and the filings do not say that either.

— PRICE • BOUNDED, NOT GUESSED

Nobody publishes a price. You can still bound one

~\$283,000

REVENUE ÷ CUSTOMERS, FY2026 — A *BOUND*, NOT A PRICE

AND THE MECHANIC UNDERNEATH IT — FROM THEIR DOCUMENTATION

You are billed per data point

“a session start, session end, custom event, or purchase recorded, as well as **any attribute set** on an end user profile” — each one counts separately

Engagement is free

Push tokens, device info, email opens and push clicks are **not** counted. That is a genuinely customer-friendly boundary

Their own advice is to send less

“**Don’t waste data points. Only update changing data!**” — a platform sold on streaming everything, priced so you shouldn’t

How many customers depends on who is counting

2,609

THE 10-K'S DEFINED METRIC, 31 JAN 2026

333

CUSTOMERS AT \$500K+ ARR
— UP FROM 202 IN FY2024

178

SELF-PUBLISHED CUSTOMER STORIES

—

INDEPENDENT DETECTION: NOT ATTEMPTED

EXPANSION IS SLOWING, AND THE ENTERPRISE PREMIUM HAS NEARLY GONE

FY2024 · all customers



117%

FY2024 · \$500k+ ARR



120%

FY2026 · all customers



109%

FY2026 · \$500k+ ARR



110%

Wide, thin, and carrying a heavy base

54.9%

UNITED STATES, FY2026

45.1%

INTERNATIONAL

0

OTHER COUNTRIES ABOVE 10% OF REVENUE

The sentence that matters

“Other than the United States, **no other individual country accounted for 10% or more of total revenue** for any of the periods presented.”

Set against the footprint

15 regional clusters and **15 legal entities across 14 territories** —
from the status page and the sub-processor disclosure respectively.

So what


Forty-five per cent of revenue is international and none of it concentrates.
That is a wide, thin footprint carrying a heavy infrastructure and legal base.

Well liked, with two consistent complaints

4.5/5

G2 · 1,702 REVIEWS


WHAT G2'S OWN TAGS PRAISE



4.5/5

GARTNER · 263 RATINGS

AND WHAT THEY CRITICISE



8.8/10

TRUSTRADIUS · 348 REVIEWS

Two themes recur across unrelated panels, which is what makes them worth quoting: reporting, and the learning curve. TrustRadius codes reporting as “limited and unintuitive, a sentiment shared by 36% of reviewers”.

Strong culture; the complaints are about ceilings

4.1/5

524 RATINGS

82%

WOULD RECOMMEND

90%

APPROVE OF THE CEO

71%

POSITIVE BUSINESS OUTLOOK

Work-life balance is not the problem

It tracks the overall rating almost exactly over six months, around 4.0–4.1. No divergence — so no trend is claimed

What their own summary names

Management effectiveness and “clarity in direction”; “limited upward mobility and discrepancies in compensation relative to market rates”

Still hiring hard

Roughly 284–300 open roles across 15 functions and 26 locations — including Bucharest, which matches a Romanian entity in the sub-processor list

Buyers weigh them against twice as many vendors as they name

NAMED IN THE 10-K — THEIR CHOICE

Adobe	Salesforce
Iterable	Klaviyo

GARTNER'S BUYER-DERIVED SHORTLIST — NOT THEIR CHOICE

Salesforce	Adobe	Iterable	Oracle
Optimove	Blueshift	MoEngage	CleverTap

THE GAP IS THE FINDING


Three of their four names appear on the buyer list — so this is not a vendor misreading its market. But **five vendors buyers actually compare them against appear nowhere in the 10-K**, and they are mostly the mobile-engagement specialists rather than the suites Braze positions against.

PART II


The product

- 01 One campaign, seven stages, end to end
- 02 Where the limits are — in their words, not ours
- 03 Which channel has no middleman
- 04 What the AI actually is, on five lenses

THEIR TECHNICAL DOCUMENTATION UNLESS MARKED OTHERWISE



What actually happens when a campaign runs



Seven stages. **Two of them are where the evidence turned out to be interesting** — how fresh the data is at stage one, and who physically sends the message at stage six. Everything in between is competent and largely as advertised, and I will say so as we go.

The way in is not the way out

60,000/min

INGEST — /USERS/TRACK, AT 75 OBJECTS EACH

250/min

READ PROFILES BACK BY IDENTIFIER, AT 50 IDS EACH

250,000/hr

BULK EXPORT TO CLOUD STORAGE

STATE THIS FAIRLY — IT IS NOT A LOCK-IN STORY

The asymmetry is real

Writing reaches ~4.5m objects a minute. Reading profiles back by identifier reaches ~12,500 a minute for a customer who joined after August 2024

But they are different operations

Writing event objects versus reading whole profiles. The ratio shows where priority sits, not like-for-like throughput

And bulk export exists

Segment export to cloud storage sits under the default 250,000/hour limit. The sanctioned bulk route is generous

Their table. Their words. Three of four are not real-time

Standard CDI sync

“**Not real-time**; minimum 15-minute sync cadence”

CDI Segments

“**Not real-time**; refreshes on your Segment Extension schedule”

CDI Canvas triggers

“**Not real-time**; bounded by sync schedule (minimum 15 minutes)”

/users/track and the SDKs

“**Near-real-time** (async processing)”

Warehouse syncs run “from every 15 minutes to once per month”. Going faster is not self-serve: “*contact your customer success manager or use REST API ingestion.*” **If your customer data lives in a warehouse, fifteen minutes is the floor.**

Generous at the top, narrow at the bottom, quiet when it fails

unlimited

ALIASES PER PROFILE

1

ALIAS PER LABEL –
UNIQUE ACROSS THE BASE

5

IDENTIFIER TYPES ACCEPTED ON INGEST

1

IDENTIFIER TYPE FOR WAREHOUSE SEGMENTS

A merge can decline and still report success

“If both profiles have invalid phone numbers, Braze does not merge them ... **The endpoint still returns 202 Accepted with a success message**, so the HTTP response does not indicate that the merge was skipped.”

Reporting splits after a merge

Dashboard summaries attribute a pre-merge send to the surviving profile. Currents, Query Builder and Messaging History attribute it to the orphaned one. Both are right by their own rules, and they disagree.

Two engines, two databases, one of them bought

RULE-BASED — MONGODB

Segmentation

Custom events, attributes, purchases and most targeting

Segment Extensions

SQL, but served from Snowflake

Global Control Group

Holdouts, and profiles in it never merge

MODEL-BASED — SNOWFLAKE

Predictive Suite

Churn and event prediction — 7 focused doc pages

AI item recommendations

Snowflake-backed

Decisioning Studio

Was OfferFit until June 2025

The split has a customer-visible consequence Braze flags itself: “Removing data from one system does not automatically remove it from the other.” Deleting bad event data means doing it in MongoDB, separately from anything Snowflake-backed.

Canvas is the strongest thing in the platform

249

FOCUSED DOCUMENTATION PAGES ON CANVAS

10

CANVAS API ENDPOINTS

385

G2 REVIEWS TAGGED 'EASE OF USE'

139

TAGGED 'LEARNING CURVE'

“Canvas makes it easy to build complex, branching lifecycle flows ... Being able to trigger contextual push notifications, in-app messages, and emails from live event streams —

Enterprise IT manager, G2 • 5/5 • August 2026

**without needing engineering
for every small tweak**

— is a
huge win.”

Liquid is the substrate, and it is also the complaint

123

FOCUSED DOC PAGES ON LIQUID TEMPLATING

43

ON CONNECTED CONTENT

38

ON CONTENT BLOCKS

78

ON CATALOGS

What it gives you

Liquid templating, Connected Content for live API calls at send time, reusable Content Blocks, and product catalogs for item-level personalisation.

What reviewers say about it

“Liquid personalization and Connected Content also make it straightforward to scale truly dynamic messaging” — and, from the same reviewer, “a challenging learning curve around Liquid syntax” for non-technical users.

Only two channels have a named middleman

CHANNELS WITH A NAMED DELIVERY SUB-PROCESSOR

Email — three of them

Amazon SES · Bird.com (SparkPost) · Twilio (SendGrid). Plus Mailgun for previewing

SMS / mobile messages — two

Infobip · Twilio

CHANNELS WITH NONE NAMED

Push

In-app

Content Cards

Banners

Webhooks

WhatsApp

LINE

KakaoTalk

Landing pages

Live notif.

Say what the disclosure names, not what exists. Absence here is not proof of no intermediary — APNs, FCM and the WhatsApp/LINE/Kakao business APIs may sit outside the definition of a sub-processor. What is certain is that **email has three interchangeable senders**, which is deliberate redundancy on their highest-volume channel.

The response comes back, and it is billed differently

71

FOCUSED DOC PAGES ON WEBHOOKS

5 min

CURRENTS EXPORT CADENCE
— OR EVERY 15,000 EVENTS

add-on

CURRENTS IS “AN OPTIONAL BRAZE ADD-ON”

free

ENGAGEMENT EVENTS ARE NOT
BILLED AS DATA POINTS

Inbound and two-way handling is real: webhooks are a first-class documented channel, and engagement tracking — opens, clicks, push receipts — is explicitly excluded from data-point billing. **But getting the raw event stream back out is a paid add-on that exports on a five-minute cadence**, which is the mechanism behind the reporting complaint on slide 14.

Thirteen documented. Ten marketed. The drift runs both ways

DOCUMENTED, WITH NO MARKETING PAGE

Banners

Transactional email

Landing pages

Live notifications

Webhooks

MARKETED, BUT MISSING FROM THEIR OWN DOCS CHANNEL INDEX

KakaoTalk

Four KakaoTalk documentation pages exist — setup, message creation, click tracking, reporting — and a product page at `/product/kakaotalk-messenger`. The channels index mentions it **zero** times.

Landing pages is the one to notice: eleven user-guide pages plus three partner pages, a substantial documented capability sold under no name at all.

Broad, and documented shallowly

150+

TECHNOLOGY PARTNERS — THEIR CLAIM

322

PARTNER DOCUMENTATION PAGES

~810

AVERAGE WORDS PER PARTNER PAGE

~2,708

AVERAGE WORDS PER DEVELOPER-GUIDE PAGE

The integration network is real and wide. But the documentation shape says something about its depth: **the partner section is the second-largest in the corpus by page count and close to the thinnest by words per page** — a lot of short entries, roughly one per partner. The developer guide, a third the size, carries more than three times the words per page.

What is bought

Delivery (Twilio, Infobip, SendGrid, SparkPost, SES), warehousing (Snowflake, Databricks, BigQuery), models (Anthropic, OpenAI, Google) — all from the sub-processor list

What is built

The ingestion layer, Canvas, Liquid, the identity model, Currents, and the integration surface itself

Fifteen clusters, seven subsystems each — and one exception

THE SEVEN SUBSYSTEMS IN EVERY CLUSTER

Dashboard
SDK Data Collection
Data Processing
REST APIs
Outbound Messaging
Currents
Cloud Data-Ingestion

THE FIFTEEN CLUSTERS

US 01	US 02	US 03	US 04	US 05
US 06	US 07	US 08	US 10	EU 01
EU 02	AU 01	ID 01	JP 01	KR 01

Corroborated independently: the sub-processor disclosure lists AWS regions for the US, EU, Australia, Indonesia, Japan and South Korea — the same six geographies. **There is no US 09.**

The most consistent complaint in the corpus

36%

OF RECENT TRUSTRADIUS REVIEWERS CALL REPORTING LIMITED AND UNINTUITIVE

140

G2 REVIEWS TAGGED 'MISSING FEATURES'

27

FOCUSED DOC PAGES ON THE GLOBAL CONTROL GROUP

2am ET

WHEN YOUR BILLING DASHBOARD REFRESHES — DAILY, CACHED

What is genuinely there


Holdouts and a Global Control Group for incrementality, campaign and Canvas analytics endpoints, KPI endpoints, and a Snowflake data share that avoids copying data at all.

What reviewers keep hitting

“The out-of-the-box reporting still feels too basic unless you export raw data through paid add-ons like Currents” — and, separately, a wish for “more robust native reporting on long-term retention cohorts”.

Bought recently, and running on other people’s models

LENS 1 — FOCUSED DOCUMENTATION PAGES



LENSES 2-5

0 of 135 API endpoints

28 namespaces. None of them AI, prediction, agent or decisioning

3 external model suppliers

Anthropic, OpenAI and Google — named in their own sub-processor disclosure

The engine was bought

Decisioning Studio was OfferFit until June 2025. \$303.2m

Reviewers name copywriting

Two G2 reviewers say the AI copy needs “a careful human hand”. None mentions decisioning

PART III


Strategy

- 01 Where seven years of revenue actually went
- 02 What is provisioned that has not been announced
- 03 What would survive a competitor doing the same thing


AUDITED, EXCEPT WHERE MARKED

Decelerating and getting more efficient at the same time

SALES & MARKETING AS A SHARE OF REVENUE



REVENUE GROWTH OVER THE SAME YEARS



This killed the hypothesis it was built to test. The expectation was decelerating growth *propped up by* sales spend. **The evidence says the opposite: S&M has fallen more than twelve points as a share of revenue since FY2023, to its lowest in the series, while growth halved.** Operating cash flow turned positive in FY2024 and has grown every year since — to \$71.4m.

One instance is not on the same cloud as the others

1 · The compelled disclosure

Their sub-processor list, revision 1 June 2026, names **two** hosting providers: Amazon Web Services and Google Cloud. **Microsoft is not named anywhere in it.**

2 · Their own documentation

Braze publishes the IPs you must allowlist, per instance. Every address listed for **US-08 is registered to Microsoft Corporation**. Every other instance — US-10, AU-01, ID-01, JP-01, KR-01 — is Amazon. Checked against ARIN.

3 · Certificate transparency

50 hosts sit on region codes p-aze-us, s-aze-us, d-aze-us — matching no AWS region identifier, where every other code does. They include sdk-us08 and subcenter-08.

HOW TO SAY THIS — AND HOW NOT TO

State the three observations and stop. A hosting arrangement may sit outside a sub-processor listing for reasons not visible from outside. Saying they failed to disclose would be a legal conclusion this evidence does not support. *The question is for them, not a conclusion for us.*

Three things a competitor cannot simply copy

Contracted revenue, not pipeline

\$1,033.0m of remaining performance obligation

— 1.40× current revenue, already signed. You cannot displace what is not up for renewal. This is the strongest of the three.

Ten years of streaming plumbing

MongoDB, Snowflake, Kafka and Redis under fifteen regional clusters, with a decade of incident history to show it holds. Copyable in principle; expensive and slow in practice.

Marketer independence from engineering

The thing reviewers actually praise: changing a journey without filing a ticket. That is a workflow habit, and habits are stickier than features.

AND WHAT DOES NOT PROTECT THEM

The AI

Bought for \$303.2m, running on Anthropic, OpenAI and Google models, with no API surface of its own. Anyone can buy the same models. This is a feature race, not a moat.

The channel roster

Broad and well built — and matched by several of the five specialists buyers shortlist them against but they never name.

PART IV

Open questions

- 01 One hard question, answered properly
- 02 A decade of reliability, measured rather than claimed
- 03 What public sources could not answer, and what would close it
- 04 What to remember

WHERE THE EVIDENCE RUNS OUT, THIS SAYS SO

“Is it really real-time?” — answered properly

ASK

Which path?

The answer differs by
a factor of hundreds

SDK / API

Near-real-time

Their words, with ‘async
processing’ attached

WAREHOUSE

15 minutes, floor

“Not real-time”,
three times over

FASTER?

Not self-serve

“Contact your customer
success manager”

EXPORT

5 minutes

And Currents is a paid add-on

Why this is the question that matters

Most enterprise buyers’ customer
data lives in a warehouse. For them
the honest answer is fifteen minutes,
and it is not a setting they can change.

Why it is not a gotcha

Braze publishes this themselves, in a
comparison table, and the SDK path
genuinely is near-real-time. The word covers
two architectures — it is not a false claim.

What to actually ask them

“Which ingestion path will *my* data take, and
what is the latency on that path?” The answer
is in their documentation before the meeting.

451 incidents since 2016, and the rate is falling



79 min

MEDIAN INCIDENT DURATION

311 min

P90 — THE LONG TAIL IS LONG

29.6%

OF NON-MAINTENANCE INCIDENTS
WERE MAJOR OR CRITICAL

63 v 27

DASHBOARD INCIDENTS
VS OUTBOUND MESSAGING

Never compare this to a vendor who publishes nothing. A decade of visible incidents is a disclosure practice, not a defect count — and silence has not earned anyone a better record.

Four gaps, and what would close each one

Does satisfaction fall with customer size?

The one hypothesis that could not be tested at all. **All three review sites paywall exactly that breakdown**, and only 7 of 860 coded records carry a segment. *Closed by:* paid panel access, or an investor-day disclosure

What else is in the certificate estate?

The host list is **partial** — 833 hosts through a rate-limited fallback after crt.sh returned errors all day. Everything found stands; **nothing is claimed about what is absent**. *Closed by:* an API token

The real split of open roles by function

The careers board's filter would not drive reliably, so only the 15-function taxonomy and a front-of-list sample were captured. **No percentage is offered**. *Closed by:* the Greenhouse board API

Actual pricing, and why the export limit was cut

No vendor publishes a rate card, and the 10× cut to the profile-lookup limit on 22 August 2024 is documented without explanation. *Closed by:* a procurement award, a customer contract, a changelog entry

Four questions worth the next week

1 · Ask about US-08

The highest-value single answer in the backlog. It decides whether slide 34 is an infrastructure observation or a procurement question.

Route: the DPA schedule a customer receives, or ask them directly.

2 · Buy one panel's segment data

It closes the only hypothesis that went unresolved, and satisfaction-by-size is the question your sales team will ask first. **Route:** paid G2 or Gartner access.

3 · Pull the Greenhouse board API

Cheap, and turns '15 functions' into where headcount is actually going — the strategy stated in hiring rather than in a keynote.

4 · Re-run this in ninety days

Every number here is reproducible by script. The two to watch: whether the gross-margin decline continues, and whether the 2026 incident rate settles or climbs.

The full backlog — fourteen open questions, each with what would close it — is in `docs/QUESTIONS.md`, and every finding in this deck resolves to a row in `docs/FACTS.md` with a source path and a capture date.

Their documentation is more
honest than anyone's marketing
— including their own.

Every hard limit in this deck came from a page Braze wrote for its own engineers: three of four ingestion paths labelled “not real-time”, a fifteen-minute warehouse floor, an export limit cut tenfold on a dated boundary, and a merge that returns success when it has quietly declined. None of it is hidden. All of it is unread.

1,352

DOCUMENTATION PAGES READ

451

INCIDENTS

845

PUBLIC ISSUES

17

SUB-PROCESSORS

7

AUDITED YEARS