Braze, Inc. — evidence record

Public sources only · every claim graded and traceable

How to read this. The chapters are the *Record*: every fact is stated **once**, in full, with its caveats in the same sentence as the number. The Slide Map is the deck in order and holds no evidence — it points back into the chapters.

Grades run AUDITED INFRASTRUCTURE DOCUMENTED THIRD-PARTY CLAIMED, ordered by how hard it would be for the vendor to have made the statement untrue. **A claim takes the grade of its weakest source.**

1 · The company

> Braze is a fifteen-year-old New York software company that has been public since November 2021 and, since January 2026, is no longer controlled by its founders through super-voting stock – though the staggered board that is the other structural defence was kept. Two of its three cofounders still run it, while three of six executive officers turned over inside a year. The most useful single thing here is a ratio: revenue has grown three times faster than headcount since FY2022, and the compensation plan shows that this was deliberate.

1.1 · The entity, and why the fiscal calendar matters

FACT	VALUE	GRADE	SOURCE
SEC registrant	Braze, Inc.	AUDITED	`data/filings.csv`
CIK / ticker	0001676238 / BRZE (Nasdaq)	AUDITED	`data/filings.csv`
SIC classification	7372 Services-Prepackaged Software	AUDITED	`data/filings.csv`
Fiscal year end	31 January	AUDITED	`data/filings.csv`
Founded and incorporated	2011, in Delaware	AUDITED	`sources/filings/2026-03-25_10-K_000013.txt:249,598`
Headquarters	28 East 28th Street, Floor 12, New York, NY 10016	AUDITED	`sources/filings/2026-03-25_10-K_000013.txt:598`
Total SEC filings	737, 2017-07-20 → 2026-08-28	AUDITED	`data/filings.csv`

The 31 January year end is not administrative trivia, and it is the first thing that will trip up anyone re-deriving these numbers. "FY2026" means the year *ending* 31 January 2026 — so it is mostly calendar 2025. Worse, XBRL's `fy` field labels the **filing**, not the period: the FY2026 10-K carries figures tagged `fy2026` that describe FY2024 and FY2025. Keying on that field silently mislabels every year in the series. `tools/sec_facts.py` derives a correct `period` column from `start`/`end` instead, and every figure in chapter 2 keys on that. Anyone who pulls the same XBRL facts and gets different years has hit this, not a disagreement with us.

`source: data/filings.csv @ 2026-09-01`

1.2 · Going public, and the end of founder control

FACT	VALUE	GRADE	SOURCE
S-1 filed	22 October 2021	AUDITED	`data/filings.csv`
Final prospectus (424B4)	18 November 2021	AUDITED	`data/filings.csv`
Net IPO proceeds	~$456.8m after discounts, commissions and expenses	AUDITED	`sources/filings/2022-03-31_10-K_000005.txt:2061`
Dual-class structure	Ended 30 January 2026	AUDITED	`sources/filings/2026-03-25_10-K_000013.txt:1197`

The governance change is nine months old at the time of writing and is easy to miss because it appears in a risk factor rather than a press release. On 30 January 2026 the Class B common stock was retired and automatically converted into Class A. The 10-K states the consequence without hedging: "our executive officers and early investors no longer hold super-voting rights. Consequently, our voting power is now more widely distributed among our public stockholders."

It is worth being precise about what this does and does not tell you, because it is the kind of fact that invites over-reading. It does **not** say anything about intent, performance or strategy. What it changes is the *mechanism*: for four years after the IPO, a shareholder who disliked the direction of the company had limited ability to do anything about it, and from January 2026 that is no longer true. For a competitor the practical reading is about tempo rather than direction. A company answerable to a distributed shareholder base responds faster to a bad quarter than one insulated from it, and Braze has now had two consecutive years in which growth decelerated (chapter 2). Whether that produces pressure is not knowable from the filings; that it *can* now is.

`source: sources/filings/2026-03-25_10-K_000013.txt:1197 @ 2026-03-25`

1.2b · Who runs it, who sits on the board, and what neither has to do with the votes

The first draft of this analysis answered "who controls it now" entirely with the Class B conversion, which is half an answer. Control is who holds the votes **and** who holds the seats, and the seats did not change at all this year.

The people, named

PERSON	ROLE	SINCE	GRADE	SOURCE
Bill Magnuson	Chairman, CEO, President, cofounder	CTO Jul 2011; director Aug 2014; CEO Jan 2017; President Jun 2025	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:780,1123
Jonathan Hyman	Chief Technology Officer, cofounder	CIO Jul 2011; CTO Dec 2016	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:1187
Isabelle Winkles	Chief Financial Officer	Jan 2020 — resigned 29 May 2026	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:1186,1177
Astha Malik	Chief Business Officer	CMO Jun 2022; CBO Jun 2023	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:1188
Edward McDonnell	Chief Revenue Officer	Jul 2025, from Asana	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:1189
Susan Wiseman	General Counsel and Secretary	Jan 2016 — retiring by 30 Jun 2026	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:1190,1181

Braze was founded by three people and two of them still run it. Braze's own company page names them: *"Braze CEO Bill Magnuson and CTO Jon Hyman participated in a 2011 hackathon that led to a chance meeting with their third cofounder, Mark Ghermezian."* CLAIMED

`source: sources/external/braze-company-page_2026-09-02.md @ 2026-09-02`

That sentence is graded `claimed` and it has to be, which is itself worth a line. **No SEC filing in this corpus uses the word "cofounder" or names Mark Ghermezian at all.** What the filings do establish, under legal penalty, is the tenure: Magnuson has held executive office since July 2011 and Hyman since July 2011, which is the founding year. So the strong version of the claim — two of the three founding team are still running the company fifteen years on — rests on an audited fact for the tenure and a marketing page for the word "founder". Say it that way.

The reference project this method comes from made a finding of exactly this shape: six founders, all still in post after fourteen years, which says something real about stability. **Braze's version is a mixed result and the mixture is the interesting part.** Two founders in post is genuine continuity at the top. The third is absent from every filing —

and the largest disclosed holder of the stock is a private Canadian investment company, MCG7 Capital Inc., holding 6.0% partly through an entity named **"Appboy BH LLC"**. Appboy was Braze's original name. The proxy connects MCG7 to no named individual and expressly says "no single person has voting or dispositive authority over such securities", so **that is where this observation stops.** AUDITED

`source: sources/filings/2026-05-18_DEF-14A_021908.txt:6339 @ 2026-05-18`

One person holds four titles, and three of six seats turned over

Magnuson is Chairman, Chief Executive, President and a director. The President title arrived in June 2025 for a specific reason the proxy states in a footnote: *"Mr. Kleeger resigned as our President and Chief Commercial Officer on June 1, 2025 and, effective immediately thereafter, Mr. Magnuson was appointed our President."* The role was absorbed rather than refilled. AUDITED

`source: sources/filings/2026-05-18_DEF-14A_021908.txt:1175 @ 2026-05-18`

That is one of three departures from a six-person executive team inside about a year:

- **Myles Kleeger**, President and Chief Commercial Officer — resigned **1 June 2025**.
- **Isabelle Winkles**, Chief Financial Officer — announced **28 April 2026**, resigning

29 May 2026, after six years in the role.

- **Susan Wiseman**, General Counsel and Secretary — announced **7 April 2026**, retiring

on or before **30 June 2026**, after ten years.

AUDITED

`source: sources/filings/2026-05-18_DEF-14A_021908.txt:1175,1177,1181 @ 2026-05-18`

The proxy gives no reason for any of the three, and this record asserts none. The temptation here is obvious and must be named so it can be refused: a chief financial officer resigning in the same year the company disclosed a material weakness in financial-reporting controls (§2.1b) is a pattern a reader will assemble on their own. There is nothing in any filing connecting them, the weakness produced no misstatement and no restatement, and the departure was announced two months after the year end. **Stating the two facts and declining to join them is the correct treatment**, and a briefing that joins them has made a claim it cannot support in a room where someone will ask.

What *can* be said is narrower and still useful. Braze replaced its President, its chief financial officer and its general counsel inside about fifteen months, while its founding CEO and CTO stayed. Continuity is at the founder layer; turnover is in the professional layer around it.

Both successions have since completed. Braze's own leadership page in September 2026 names **Pankaj Malik** as Chief Financial Officer and Chief Accounting Officer and **Chris Lal** as General Counsel and Corporate Secretary — neither of whom appears in any filing in this corpus. That page lists ten leaders against the proxy's six executive officers, which is mostly a difference of scope rather than of headcount, and is recorded as conflict **C-10**. CLAIMED

`source: sources/external/braze-ir-leadership_2026-09-02.md @ 2026-09-02`

The board, and the defence that did not get retired

FACT	VALUE	GRADE	SOURCE
Seats	Seven	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:700-756
Independent	Six of seven; Magnuson is not, "due to his position as our President and Chief Executive Officer"	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:793
Chair	The CEO, with Phillip Fernandez as Lead Independent Director	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:796
Structure	Classified into three staggered classes	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:690-756
Investor seat	Neeraj Agrawal, General Partner of Battery Ventures, director since April 2016; Battery entities and his own holdings total 5,711,181 shares — 5.1%	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:6255,6341

Six independent directors out of seven is a good ratio and should be said plainly. The board is also experienced in a way that is specific rather than decorative: Phillip Fernandez cofounded and ran Marketo, which is the closest thing to a direct predecessor company in this category; David Obstler is the sitting CFO of Datadog; Yvonne Wassenaar ran Puppet and sits on three other public boards.

Two things about that board are worth a competitor's attention, and only one of them is critical.

The chair is the chief executive. The proxy does not gloss this — it argues for it, saying that separating the roles "has the potential to give rise to divided leadership, which could interfere with good decision-making". That is a real argument, made openly, and the mitigation — a Lead Independent Director — is the standard one. Quote their reasoning rather than characterising it.

The board is classified, into three staggered classes, so only about a third of it stands for election in any year. This is the finding that pairs with §1.2 and it cuts against the easy reading of the Class B retirement. Braze gave up super-voting stock in January 2026 and the disclosure language was generous about what that meant for shareholders. A staggered board is the other structural defence, it is arguably the more durable of the two, and **it did not change.** So the honest summary of Braze's governance in 2026 is that one takeover defence was retired and the other was kept — which is neither a scandal nor the unqualified opening-up the risk-factor language suggests on its own.

And the votes, which are now ordinary

HOLDER	SHARES	%	GRADE
Entities affiliated with MCG7 Capital Inc.	6,709,408	6.0%	AUDITED
Neeraj Agrawal / Battery Ventures	5,711,181	5.1%	AUDITED
Bill Magnuson	5,485,123	4.9%	AUDITED
Jonathan Hyman	2,029,437	just under 2%	AUDITED

`source: sources/filings/2026-05-18_DEF-14A_021908.txt:6255 @ 2026-05-18`

Against 111,783,722 Class A shares outstanding, and with the Class B stock gone, **no holder and no director has a blocking position.** MCG7 is the only 5%-or-greater holder Braze discloses. For anyone modelling how quickly Braze could be pushed to respond to a shareholder, that is the arithmetic — and it is new this year.

The related-party line worth reading twice

The proxy discloses that Braze bought **approximately $3.8m of Datadog services in FY2026**, and that director David Obstler is Datadog's Chief Financial Officer. AUDITED source: sources/filings/2026-05-18_DEF-14A_021908.txt:6420 @ 2026-05-18

That is an ordinary related-party disclosure, properly made, and **nothing improper is alleged.** It is here because of what it connects to. Chapter 5 records, from an entirely different document, that Braze "may provide End User metadata, such as user identifiers, to DataDog for support and application troubleshooting". So a monitoring vendor that receives end-user identifiers is also a company whose chief financial officer sits on Braze's board — two facts, disclosed independently in two documents neither of which references the other, and neither disclosure is deficient.

The reason to hold both is practical rather than insinuating: a buyer's data-protection review will ask about the sub-processor, and a buyer's procurement team may separately ask about board relationships. Knowing they are the same company means neither question surprises you.

One last small thing, recorded because absence of it would be the finding: **the proxy discloses its own late filings.** Three Section 16 reports covering thirteen transactions were filed late by Battery Partners XI and its affiliates, and one initial statement by Edward McDonnell. That is a housekeeping disclosure of the sort every proxy carries; it is noted here only so a reader who finds it later knows it was seen and judged immaterial. AUDITED

source: sources/filings/2026-05-18_DEF-14A_021908.txt:6367 @ 2026-05-18

1.3 · Headcount — and the ratio that actually matters

FISCAL YEAR	FULL-TIME EMPLOYEES	REVENUE PER EMPLOYEE
FY2022	1,164	$204k
FY2023	1,501	$237k
FY2024	1,548	$305k
FY2025	1,699	$349k
FY2026	1,988	$371k

AUDITED

source: the "we had a total of N full-time employees" line in each 10-K, FY2022-FY2026, with revenue from data/financials_annual.csv @ 2026-09-02

A single headcount figure is nearly useless. The series is not.

Over five fiscal years Braze added **71%** more people and **210%** more revenue. Revenue per full-time employee rose from $204,000 to $371,000 — an **82%** improvement. That is operating leverage, it is audited, and it is the single most favourable thing in this entire record. It is also the fact that makes chapter 2's story coherent rather than merely reassuring: a company whose sales-and-marketing share of revenue falls twelve points is not doing it by cutting marketing spend in isolation, it is doing it by getting more revenue out of each person it employs.

One year in that table deserves its own sentence. **In FY2024, headcount grew 3.1% while revenue grew by nearly a third.** That is not gradual efficiency; that is a hiring pause held through a year of strong growth, and it is where most of the 82% improvement was banked. Hiring resumed afterwards — FY2026 added 289 people, the largest absolute increase in the series — but the ratio never went back.

The counterweight, and it belongs here rather than in a footnote: the same period is when customer expansion slowed (chapter 7) and share-based compensation began exceeding the entire net loss (chapter 2). Efficiency improved *and* growth decelerated. Both are true, they are visible in the same years, and an account that gives you only one of them is selling something.

What the number does not cover

1,988 counts full-time employees only. Braze's own proxy, drawing a median employee on the same date, uses a different population — "all of our full-time **and part-time** employees… We did not include any independent contractors" — and never sizes it. So there are at least three populations in play (full-time; full-time plus part-time; and those plus contractors), Braze publishes one of them, and no independent headcount was captured in this run to triangulate against. This is conflict **C-07**, and its ruling is simple: say "1,988 full-time employees", never "headcount of 1,988".

`source: sources/filings/2026-05-18_DEF-14A_021908.txt:5769 @ 2026-05-18`

1.4 · What management is paid to optimise

This is the question `QUESTIONS.md` #33 asked and the first pass left open, because it never opened a proxy statement. The answer is unusually legible.

FACT	VALUE	GRADE	SOURCE
FY2026 bonus-plan weighting	Net CARR **60%**, non-GAAP operating income (loss) **40%**	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:1994
Profitability gate	Missing a specified non-GAAP operating income target as a share of gross profit results in "0% achievement of this performance metric"	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:1994
Growth cap	If retention fell below target, Net CARR achievement "would be limited to 110%"	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:1994
CEO total compensation, FY2026	$14,230,598	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:5771
Median employee total compensation	$164,000	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:5771
CEO pay ratio	87 to 1	AUDITED	sources/filings/2026-05-18_DEF-14A_021908.txt:5771

Read the three mechanics together and the plan describes a strategy more precisely than any keynote does.

Sixty per cent of the financial component pays for **net contracted recurring revenue** — growth, but *contracted* growth net of churn, which is a harder target than bookings. Forty per cent pays for **non-GAAP operating income**, and that portion has a hard floor: miss a specified operating-income-to-gross-profit threshold and the metric pays **nothing**, regardless of how close you came. And the growth half is capped at 110% if retention misses, which explicitly forbids buying CARR at the expense of the existing base.

That is a compensation plan built to produce exactly the pattern chapter 2 observes: decelerating growth alongside improving efficiency. It does not prove causation. It does mean that when a competitor asks "will Braze discount aggressively to win this deal?", the honest answer from the public record is that the people approving that discount are paid on a metric with a profitability gate attached, and were paid that way through the year in which sales-and-marketing spend fell to its lowest share of revenue in seven years.

The pay ratio is 87 to 1 on a median of $164,000. Both numbers are worth quoting together rather than separately: a high median softens the ratio, and it also sits awkwardly beside the Glassdoor complaint recorded in chapter 7 about "discrepancies in compensation relative to market rates". A $164,000 median is not a low-paying company. The grievance, if the reviews are read closely, is about *progression* rather than level — which is the same thing the "limited upward mobility" comment says.

1.5 · The legal footprint, and one entity that is not what it looks like

FACT	VALUE	GRADE	SOURCE
Group entities	15, across 14 territories	DOCUMENTED	`sources/clean/braze-subprocessors.md:41`
Territories named	Australia, Brazil, Canada, France, Germany, Ireland, Spain, United States, Ireland & Romania, Japan, South Korea, United Kingdom, UAE, Singapore, Indonesia	DOCUMENTED	`sources/clean/braze-subprocessors.md:41`

This list comes from the sub-processor disclosure rather than from marketing, which is what makes it worth having: it is compelled to be complete, and it names its own territories rather than gesturing at "global presence".

Two things in it repay attention.

Braze KK, in Japan, is not a wholly-owned subsidiary. Outside investors — Japan Cloud Computing and M30 LLC — bought $10.0m of its stock across 2020 and 2021; employees hold options over its shares; and a redeemable non-controlling interest sits outside permanent equity on the balance sheet. Braze holds a majority and consolidates it as a Variable Interest Entity. Anyone reasoning about Braze's Japanese operation from the entity list alone would assume ownership and control that the filings do not support. This is conflict **C-09**.

"Braze Ireland Procurement Limited" is located in "Ireland & Romania", and the careers board independently lists Bucharest as a hiring location. Two documents with no relationship to each other — a data-protection disclosure and a job board — describing the same Romanian engineering presence is the kind of corroboration this method exists to find, and it is the only evidence in the corpus that Braze operates in Romania at all. Neither document alone would have been enough to say so.

Set the footprint against the revenue. Fifteen legal entities and fifteen regional clusters (chapter 4) support a business where, outside the United States, **no single country reaches ten per cent of revenue** (chapter 2). That is a deliberately wide, thin international structure, and it is expensive to run. It is also the necessary shape if data-residency commitments are part of what you sell, which chapter 4 argues they are.

`source: sources/clean/braze-subprocessors.md:41 @ 2026-09-02`

What would change this chapter

A merger, take-private or change of control — all disclosed by 8-K within four business days. A materially different headcount in the next 10-K, or, more usefully, any disclosure that sizes the part-time population and closes C-07. A buy-out of the Braze KK minority, which would resolve C-09 and would appear in the equity note. A change to the bonus-plan weightings in the next proxy, which is the earliest public signal that the growth-versus-profitability balance has been re-struck. The entity list changes whenever the sub-processor document is revised; this record uses revision 1 June 2026 and any later revision supersedes it.

2 · The money

> Growth is decelerating and efficiency is improving at the same time, and the two are connected rather than coincidental. Revenue growth has halved since FY2023 while sales and marketing fell from 56.7% of revenue to 44.3%; operating cash flow has been positive and growing for three years; and the net loss persists mainly because share-based compensation now exceeds it. One caveat qualifies everything here, and it is stated first.

Every figure in this chapter is audited, filed under legal penalty, and re-derivable by running `python3 tools/sec_facts.py`, which writes `data/financials_annual.csv` and `data/financials_quarterly.csv` from the SEC XBRL company-facts API. Fiscal years end 31 January — see §1.1 for why that matters when re-deriving them.

2.1b · The caveat that travels with every figure in this chapter

Braze disclosed a **material weakness in its internal control over financial reporting** as at 31 January 2026. The CEO and CFO concluded that disclosure controls were "not effective at the reasonable assurance level". The weakness is specific: "ineffective information technology general controls, or ITGCs, in the areas of **user access and program change management** over the information technology systems that support our financial reporting processes" — and the company states that automated and IT-dependent manual controls relying on that environment were consequently ineffective too.

AUDITED

source: `sources/filings/2026-03-25_10-K_000013.txt:3258,3286 @ 2026-03-25`

Both halves belong in the same breath, because either one alone misleads.

The reassuring half is substantial, and it is the company's own and the auditor's: the material weakness "did not result in any identified misstatements to the financial statements, and there were no changes to previously issued financial results"; management concluded the statements "fairly present in all material respects"; and Ernst & Young issued an attestation report on internal control over financial reporting (`sources/filings/2026-03-25_10-K_000013.txt:3286`). Remediation is under way with "significant progress" claimed and no completion date given. OfferFit was excluded from the assessment, as SEC guidance permits in an acquisition's first year — it is 1.4% of total assets and 1.8% of revenue (`sources/filings/2026-03-25_10-K_000013.txt:3258`), so the exclusion is immaterial to the figures below.

The unreassuring half is that the control environment producing every number in this chapter was judged not effective by the people who signed it, and the specific area is user access and change management.

How much should that change your confidence? Less than the phrase "material weakness" suggests, and more than zero. A material weakness is a statement about the *probability* that an error could occur and go undetected, not a finding that one did. Here the independent checks all came back clean: no misstatement identified, nothing restated, an auditor attestation, and — separately and mechanically — this project's own restatement sweep across twenty-nine XBRL concepts found no superseded value — `data/financials_restated.csv`, produced by `tools/sec_facts.py` (conflict **C-05**). Three confirmations that the outputs are sound, against one disclosure that the process producing them is not yet controlled to standard.

Two disciplines apply. **Do not extrapolate this into a claim about customer data security.** It is scoped to financial-reporting systems and says nothing about the platform, and the temptation to make that leap — for a company whose product is data infrastructure — is exactly why the scope should be restated every time. And do not bury it: a competitor briefing that omits this and is later asked about it has damaged its own credibility for no gain.

2.1 · Seven years of revenue and margin

FISCAL YEAR	REVENUE ($M)	GROWTH	GROSS MARGIN
FY2020	96.4	—	63.0%
FY2021	150.2	55.9%	63.7%
FY2022	238.0	58.5%	67.0%
FY2023	355.4	49.3%	67.4%
FY2024	471.8	32.7%	68.7%
FY2025	593.4	25.8%	69.1%
FY2026	738.2	24.4%	67.1%

AUDITED

source: data/financials_annual.csv @ 2026-09-02

Two things happen in the final row and they should be read together. Growth continues to decelerate — it has roughly halved over three years — and gross margin falls for the first time in six years, by 2.0 points.

The deceleration is the less interesting of the two, because it is what happens to every subscription business as its base grows: adding $145m to a $593m base is a larger absolute achievement than adding $117m to a $355m base, and the percentage falls anyway. Note also that FY2026's growth is close to flat against FY2025 — the halving happened between FY2023 and FY2025, and has since largely stopped. An account that says "growth is collapsing" is describing a period that ended two years ago.

The margin fall is the one that repays attention, and Braze explains it themselves in terms specific enough to check: "primarily due to acquisition related operating costs, including personnel costs of acquired workforce and amortization expense of acquired technology, in addition to increased costs related to our tech stack."

The acquisition is OfferFit (chapter 3), whose developed technology is amortised **to cost of revenue** rather than to operating expense. That accounting choice is why an AI acquisition shows up in the gross-margin line rather than below it, and it is why the decline is best read not as pricing pressure but as the price of buying an AI product. What the sentence does not do is split the two causes. "Increased costs related to our tech stack" is unquantified, and for a company that has just added three foundation-model providers to its sub-processor list (chapter 6), the size of that component is the single most useful undisclosed number in the filing. It is open question 59.

source: sources/filings/2026-03-25_10-K_000013.txt:1565 @ 2026-03-25

2.2 · Losses, cash and the equity cost

FISCAL YEAR	NET LOSS ($M)	OPERATING CASH FLOW ($M)	SHARE-BASED COMP ($M)
FY2022	−76.7	−35.4	47.2
FY2023	−139.0	−22.3	72.2
FY2024	−129.2	+6.8	97.2
FY2025	−103.7	+36.7	115.1
FY2026	−131.3	+71.4	143.7

AUDITED

source: data/financials_annual.csv @ 2026-09-02

The favourable reading and the uncomfortable one are both in this table and both are true.

Favourably: operating cash flow crossed into positive territory in FY2024 and has grown every year since, roughly doubling each time. A company generating $71.4m of operating cash is not a company in difficulty, whatever the loss line says, and this is the number to reach for when someone characterises Braze as loss-making and therefore fragile. It is not fragile. It funds itself.

Uncomfortably: share-based compensation has exceeded the entire net loss for two consecutive years — 111% of it in FY2025 and 109% in FY2026. The arithmetic consequence is worth stating plainly, because it is the fact most often fudged in both directions: **on these figures, excluding share-based compensation, the business is around break-even.** That is not the same as saying it is profitable. Share-based compensation is a real cost; it is simply borne by shareholders through dilution rather than by the company through cash.

The dilution is measurable. Weighted-average diluted shares went from 94.6m in FY2023 to 107.9m in FY2026 (`data/financials_annual.csv`) — a rise of 14.1% in three years. So the honest summary is that Braze funds its own operations from cash while transferring roughly a seventh of the equity to employees over three years, and reports the result as a loss. Each of those three clauses gets misused on its own.

For a competitor the practical reading is that **the loss line is not a pressure point**. Anyone building a sales narrative around Braze running out of money is arguing against three years of rising operating cash flow and $124.3m of cash on the balance sheet.

2.3 · Where the money goes

FISCAL YEAR	S&M SHARE OF REVENUE	R&D SHARE OF REVENUE
FY2020	59.5%	21.1%
FY2022	53.4%	24.8%
FY2023	56.7%	27.4%
FY2024	52.4%	25.4%
FY2025	47.6%	22.6%
FY2026	44.3%	22.6%

AUDITED

`source: data/financials_annual.csv @ 2026-09-02`

This table killed a hypothesis. The analysis went in expecting growth to be decelerating *while sales spend held* — the familiar pattern of a company buying its last few points of growth. The evidence says the opposite: S&M has fallen more than twelve points as a share of revenue since FY2023, to its lowest in the seven-year series, while growth decelerated. Braze is decelerating and getting more efficient at the same time.

Three other parts of this record explain how, and together they make it a finding rather than a curiosity. Revenue per employee rose 82% over the same period (§1.3), with the bulk of it banked in a single year when hiring nearly stopped. The compensation plan pays 40% of its financial component on non-GAAP operating income, with a hard gate that pays nothing at all if a threshold is missed (§1.4). And operating cash flow turned positive in the same year the hiring pause happened. A cost discipline visible in four independent places is a policy, not an accident.

R&D tells a quieter story. It peaked at 27.4% of revenue in FY2023 and has settled at 22.6% for two years. In absolute terms it is still growing — $167.1m in FY2026, from `data/financials_annual.csv` — but as a share of revenue Braze now spends less on R&D than it did three years ago, during a period in which it also bought its AI product rather than building it (chapter 3). Those two facts are not proof of a strategy, but they point the same way, and a competitor should notice that the R&D line did not spike to fund an AI programme.

Sales and marketing still outspends research and development by 1.96 times — $327.0m against $167.1m in FY2026 — which is ordinary for enterprise software and worth stating plainly rather than treating as a finding.

2.4 · Contracted revenue and geography

FACT	VALUE	GRADE	SOURCE
Remaining performance obligation, FY2026	$1,033.0m	AUDITED	`data/financials_annual.csv`
RPO relative to FY2026 revenue	1.40×	AUDITED	derived from `data/financials_annual.csv`
United States revenue, FY2026	$405.1m — 54.9%	AUDITED	`sources/filings/2026-03-25_10-K_000013.txt:2421`
International revenue, FY2026	$333.1m — 45.1%	AUDITED	`sources/filings/2026-03-25_10-K_000013.txt:2421`

Remaining performance obligation crossed a billion dollars in FY2026, having risen every year from $234.2m in FY2021 (`data/financials_annual.csv`). At 1.40 times current-year revenue it is a genuine forward-visibility measure and, unlike a pipeline number, it is contracted and audited.

This is the most commercially useful number in the chapter and it is routinely under-used. RPO is revenue Braze has already signed and not yet recognised. For anyone selling against Braze it says, with an auditor's signature on it, that **roughly a year and a half of current revenue is not available to compete for** — not because customers would not move, but because they are inside contract terms. It reframes the competitive question from "can we win this account" to "when does it come up", and for most of the base the answer is not this year. Chapter 7's retention figures then tell you what happens when it does.

The geographic split is an audited disclosure rather than an inference, and it carries a sentence that matters more than the percentages: "Other than the United States, no other individual country accounted for 10% or more of total revenue for any of the periods presented." International share has been near-flat for three years — 43.4%, 45.0%, 45.1%, from the segment note at `sources/filings/2026-03-25_10-K_000013.txt:2421` — so the international business is growing at roughly the same rate as the whole, not faster.

Put that beside the infrastructure. Forty-five per cent of revenue is international and none of it concentrates, supported by fifteen regional clusters (chapter 4) and fifteen legal entities across fourteen territories (§1.5). That is a wide, thin footprint carrying a heavy fixed base — expensive, and the necessary shape if regional data residency is part of what is being sold. It also means no single non-US market is large enough that losing it would be visible in the consolidated numbers, which cuts both ways: resilient to a bad year in any one country, and lacking a second home market to fall back on.

What would change this chapter

A restatement — none has occurred, and both the mechanical sweep and the 10-K confirm it (**C-05**). Remediation of the material weakness, which would remove §2.1b's caveat and would be disclosed in the next 10-K or 10-Q. A second consecutive year of gross-margin decline, which would move FY2026 from an acquisition effect to a pattern. Any disclosure that splits "increased costs related to our tech stack" from acquisition amortisation (question 59). And a change in the S&M trajectory, which would say the efficiency programme had run its course.

3 · What they bought

> Braze has made two material acquisitions, and the pair reads as one strategy stated twice: buy the capability rather than build it, and buy the market rather than partner into it. One bought a reseller and its earn-out paid nothing; the other bought the product Braze now markets as its own artificial intelligence, for three hundred and three million dollars. Both prices are disclosed, which a private vendor almost never permits.

3.1 · OfferFit — bought June 2025, renamed AI Decisioning Studio

FACT	VALUE	GRADE	SOURCE
Date	2 June 2025	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3130
Total consideration	$303.2m	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:1788
Structure	$195.3m cash, $107.6m in Braze Class A common stock	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3130
Goodwill	$233.4m	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3134
Identified intangibles	$66.6m	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3134
Developed technology	$56.7m, amortised over six years to cost of revenue	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3137,3139
Customer relationships	$9.0m, six years, amortised to sales and marketing	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3137,3141
Trademarks	$0.9m	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3137

The single most useful sentence in the whole filing set, for anyone assessing Braze's AI position, sits in the auditor's critical-audit-matter note: **"OfferFit, Inc. ('OfferFit') which is now known as AI Decisioning Studio."** The product Braze markets as BrazeAI Decisioning Studio is a company it bought fifteen months ago, renamed. The 10-K describes what was acquired as "OfferFit's multi-agent decisioning engine", which is also the answer to what the agentic layer is: acquired, not built.

What the purchase-price allocation tells you

An allocation is a valuer's opinion, but it is a *disclosed* opinion with an auditor's attention on it, and reading it carefully answers questions the press release does not.

Goodwill is 77.0% of the price. That is the residual after identifiable assets, and a figure that high says most of what was bought is not on the balance sheet: the team, the customer momentum, and the position. It is entirely normal for an acquihire-plus-product deal and entirely abnormal for buying a mature revenue stream. Set it beside the share of consolidated revenue OfferFit contributed in its first eight months, which §2.1b records and which is small, and the shape is clear. Braze did not buy revenue. It bought a capability and the people who built it, at a price set by what building that capability would have cost in time.

Only $66.6m of $303.2m was identifiable. Of that, developed technology is $56.7m — 85% of the identified intangibles. So the thing being valued is overwhelmingly the software itself rather than the customer base ($9.0m) or the brand ($0.9m). The trademark valuation is the quiet confirmation of intent: at $0.9m over a $303.2m deal, nobody expected the OfferFit name to survive, and it did not.

The amortisation lands in cost of revenue. Braze amortises the developed technology straight-line over six years, and books it to cost of revenue rather than operating expense. At full run-rate that is roughly $9.5m a year against the gross margin line. This is the accounting mechanism behind chapter 2's first margin decline in six years, and it is why an AI acquisition is visible in gross margin rather than below it. It also means the effect persists: the margin drag does not wash out next year, it runs to 2031.

Ernst & Young flagged the valuation as a critical audit matter — "complex due to the significant estimation required by management", with the fair value estimated on an excess-earnings method whose key assumptions are forecast sales from the acquired business and an expected obsolescence rate. That is a routine disclosure about valuation difficulty, not a finding about the business, and it should be described that way. What it does usefully confirm is that the $56.7m rests on a forecast of *future* OfferFit revenue. If that forecast disappoints, the impairment would be disclosed, which is one of the few forward-looking signals a competitor gets for free.

`source: sources/filings/2026-03-25_10-K_000013.txt:1788-1792,3130-3141 @ 2026-03-25`

Why this matters more than a typical tuck-in

Three hundred and three million dollars is not a tuck-in for a company with $738.2m of revenue and $124.3m of cash — it is roughly two and a half times the year-end cash balance, which is why $107.6m of it was paid in stock. Braze was willing to dilute its own shareholders, in a year when three-year dilution was already at the level §2.2 records, to own this capability rather than partner for it.

Set that against the R&D line, which did not spike in the same period and has been flat as a share of revenue for two years (§2.3). The public record shows a company that chose to buy its decisioning engine rather than fund one internally, and paid a price that required equity to complete. Chapter 6 measures what arrived.

3.2 · North Star Y — bought June 2023, and the earn-out that paid nothing

FACT	VALUE	GRADE	SOURCE
Date and what it was	1 June 2023 — "Braze's exclusive reseller in Australia and New Zealand"	AUDITED	sources/filings/2025-03-31_10-K_000054.txt:2972
Total consideration	$26.8m — $20.6m cash, $6.1m stock, $1.8m contingent at fair value	AUDITED	sources/filings/2025-03-31_10-K_000054.txt:2974
Earn-out available	Up to $10.0m in year one and $16.0m in year two, on "qualified revenue performance metrics"	AUDITED	sources/filings/2025-03-31_10-K_000054.txt:2974
Earn-out paid	Nil	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:2569
Allocation	Goodwill $28.4m; intangibles $3.8m	AUDITED	sources/filings/2025-03-31_10-K_000054.txt:2976
Indemnification holdback	$2.8m, released in full during FY2025	AUDITED	sources/filings/2025-03-31_10-K_000054.txt:2978

Braze bought its own Australia and New Zealand reseller to take the market direct. The sellers could have earned up to a further $26.0m on "qualified revenue performance metrics" over the two years following completion.

They earned none of it. The FY2026 10-K records that Braze "reduced the contingent consideration liability related to the acquisition of North Star Y, Pty Ltd to zero as it was determined that the sellers did not satisfy the earn-out qualifications."

How to read an earn-out that pays zero

This is the finding most easily over-read in the entire record, so the reasoning matters more than the fact.

An earn-out is a price-setting device for a disagreement. The buyer thinks the business will do X; the seller thinks it will do more; rather than argue, they agree the seller gets paid the difference if the seller turns out to be right. **A zero payout therefore tells you one thing precisely: the revenue trajectory the sellers signed up to in June 2023 did not happen.** It tells you that from the buyer's own audited filing, in a document where getting it wrong has consequences.

It does **not** say the acquisition failed, and the filings do not say that either. Braze still owns the AU/NZ business, still runs an AU-01 cluster (chapter 4), and never took an impairment against the $28.4m of goodwill — which it would have had to disclose. A plausible and entirely benign reading is that the targets were set aggressively precisely because an earn-out is where optimism gets priced, and that the business is performing adequately but not at the seller's number.

What is fair to say is narrower and still useful. **Braze paid $26.8m for direct control of Australia and New Zealand, structured up to half the potential value as performance contingent, and the performance conditions were not met.** For anyone competing in that region, the public record contains no evidence of the acceleration the deal was priced for. That is worth knowing and it appears on no website.

Two further details close the loop. The $2.8m indemnification holdback was **released in full**, which means no indemnity claims were made against the sellers — so whatever happened, it was a performance shortfall rather than a dispute about what was sold. And the results of North Star were "not material" to Braze's consolidated statements in the years after acquisition, which is consistent with a regional reseller of that size and tells you the shortfall never threatened the group numbers.

`source: sources/filings/2026-03-25_10-K_000013.txt:2569 @ 2026-03-25`

3.3 · The two together

Read as a pair, the acquisitions describe a consistent posture rather than two unrelated events.

Both were **capability or control purchases, not revenue purchases** — goodwill was 77.0% of the OfferFit price and $28.4m of the $26.8m-plus-contingent North Star price, and in neither case did the acquired revenue matter to the consolidated line. Both were paid partly in **stock**, which for a company already diluting at the rate §2.2 records is a deliberate choice to preserve cash. And in both cases Braze bought something it had previously accessed through a partner: a reseller relationship in one, and — on the evidence of chapter 6 — an entire product category in the other.

The pattern a competitor should take from this is that **Braze will buy rather than wait**, that it will use equity to do it, and that its most recent and largest purchase was made in the category where it was furthest behind.

What would change this chapter

A third acquisition, which would appear as an 8-K and then in the following business-combination note. A goodwill impairment against either purchase, which would be disclosed and would change how the OfferFit price should be read — and, for North Star, would convert a missed earn-out into something considerably more serious. Measurement-period adjustments to the OfferFit allocation are still possible; the figures here are the FY2026 10-K's as-adjusted values. And a disclosure of AI Decisioning Studio revenue separate from the consolidated line, which would let anyone judge whether $303.2m was well spent.

4 · The platform

> Braze's documentation is markedly more honest than its marketing about what the platform will and will not do, and the limits it admits are commercial as much as technical: three of four ingestion paths are labelled "not real-time" by Braze itself, the endpoint that reads user profiles back out was cut tenfold for customers who joined after August 2024, and the billing unit is the individual attribute write.

4.1 · The corpus this chapter rests on

FACT	VALUE	GRADE	SOURCE
Documentation pages captured	1,352	DOCUMENTED	`data/docs_index.csv`
Words in the corpus	1,565,479	DOCUMENTED	`data/docs_sections.csv`
Documented REST endpoints	135, across 28 top-level namespaces	DOCUMENTED	`data/api_endpoints.csv`
Public repositories	137	DOCUMENTED	`data/repos.csv`

The corpus is shaped unevenly, and the shape is itself informative: 696 user-guide pages, 322 partner pages, 200 API pages, 106 developer-guide pages. The developer guide averages about 2,700 words a page; the partner section averages about 810. The integration network is broad and documented shallowly — a lot of short pages, one per partner.

`source: data/docs_sections.csv @ 2026-09-02`

4.2 · How fresh the data is — their table, not ours

Braze publishes a comparison of its four ingestion paths and grades each one's latency itself. This is the single most load-bearing table in the corpus.

PATH	BRAZE'S OWN LATENCY DESCRIPTION
Standard Cloud Data Ingestion sync	"Not real-time; minimum 15-minute sync cadence"
CDI Segments (Connected Sources)	"Not real-time; refreshes on your Segment Extension schedule"
CDI Canvas triggers	"Not real-time; bounded by sync schedule (minimum 15 minutes)"
`/users/track` and the SDKs	"Near-real-time (async processing)"

DOCUMENTED

`source: sources/docs/docs__user_guide__example_library__data__compare_data_ingestion_options.md:86-90 @ 2026-09-02`

Three of the four are labelled not real-time by the vendor. The fourth is qualified as near-real-time with asynchronous processing. Warehouse syncs "can run from every 15 minutes to once per month", and going faster than fifteen minutes is not self-serve: "contact your customer success manager or use REST API ingestion."

This is the substance of conflict **C-01**. The 10-K describes "high-volume, continuous streaming of user data to be processed in real time" and the developer guide says "the platform can handle any data in real time" — both true of the SDK and API path. The correct sentence is not that Braze overstates, but that the word covers two architectures: if your customer data lives in a warehouse, fifteen minutes is the floor unless you move to the API.

Other admitted ceilings on the same page: 75 objects per `/users/track` request combined across attributes, events and purchases; a 60-minute query runtime cap per connected source; approximately 3.75 million Canvas entries per hour per sync run. Currents, the streaming export, is "an optional Braze add-on" that exports "every five minutes, or every 15,000 events, whichever comes first."

4.3 · Data in versus data out

DIRECTION	ENDPOINT	LIMIT
In	`/users/track`	3,000 requests per 3 seconds for customers with data points in their pricing — 60,000 a minute — at 75 objects each
Out, by identifier	`/users/export/ids`	**250 requests per minute** if onboarded on or after 22 August 2024; **2,500** if before
Out, in bulk	`/users/export/segment`	The default 250,000 requests per hour

DOCUMENTED

`source: sources/docs/docs__api__api_limits.md:29-39 @ 2026-09-02`

Two things here, and the second is the interesting one.

The asymmetry is real but must be stated fairly. Writing can reach roughly 4.5 million objects a minute; reading profiles back by identifier is capped at about 12,500 profiles a minute for a customer who joined after August 2024, at 50 identifiers per request. Those are different operations — writing event objects against reading whole profiles — so the ratio shows where the design and commercial priority sits, not a like-for-like throughput comparison. And bulk export exists and is generously limited, so this is not a data-lock-in story and should not be told as one.

The dated boundary is the finding. On 22 August 2024 the profile-lookup limit was cut by a factor of ten for new customers, and existing customers kept the old limit. Two customers on the same product, doing the same thing, are subject to limits an order of magnitude apart depending on when they signed. Braze documents this openly in a table. Why it changed, and whether anyone was told, is open question 55.

4.4 · What it is built on

FACT	VALUE	GRADE	SOURCE
Named backing stores	Snowflake, Kafka, MongoDB, Redis	DOCUMENTED	sources/docs/docs__developer_guide__getting_started__architec
MongoDB-backed	Custom events, custom attributes, user profiles, purchases, most segmentation	DOCUMENTED	…architecture_overview.md:46-56
Snowflake-backed	SQL Segment Extensions, Prediction Suite, AI Personalized Item Recommendations, Estimated Real Open Rate	DOCUMENTED	…architecture_overview.md:58-68
Regional clusters	15: US 01–08, US 10, EU 01–02, AU 01, ID 01, JP 01, KR 01	DOCUMENTED	data/status_components.csv
Subsystems per cluster	7, identical in every one	DOCUMENTED	data/status_components.csv

The status page is an architecture disclosure the vendor made by accident. Every cluster exposes the same seven subsystems — Dashboard, SDK Data Collection, Data Processing, REST APIs, Outbound Messaging, Currents, Cloud Data-Ingestion — which is a functional decomposition of the product, published live and updated during outages. Note that Currents and CDI are first-class per-cluster subsystems rather than features.

The cluster geography is corroborated independently: the sub-processor disclosure lists AWS regions for the United States, European Union, Australia, Indonesia (with backup in Singapore), Japan and South Korea — the same six geographies the cluster names use. Channel delivery is modelled differently, as a single global group rather than per cluster. There is no US 09; the US clusters run 01 to 08 and then 10.

The customer-visible consequence of the architecture is not the clusters. It is the storage split, and Braze flags it themselves under an "important" callout: **"Removing data from one system does not automatically remove it from the other."** Erroneous custom-event data must be addressed in MongoDB, separately from anything Snowflake-backed. For a buyer with deletion obligations, that is a two-system problem stated in the vendor's own documentation.

`source: sources/docs/docs__developer_guide__getting_started__architecture_overview.md:70 @ 2026-09-02`

4.5 · One instance is not on the same cloud as the others

This is the clearest result the method produced, and it comes from three sources that have nothing to do with each other.

- The sub-processor disclosure of 1 June 2026 names exactly two hosting providers:

Amazon Web Services, and Google LLC for Google Cloud Platform. Microsoft is not named anywhere in the document. DOCUMENTED

- A **second, independent company document** says the same thing. The 10-K risk factors

state: "We rely upon third-party providers of cloud-based infrastructure, **including Amazon Web Services and Rackspace**, to host our products." Microsoft is not named there either. AUDITED

- Braze's own documentation publishes, per instance, the IP addresses a customer must

allowlist so Braze can reach their warehouse. Every address listed for **US-08** is registered to **Microsoft Corporation**. Every address listed for every other instance — US-10, AU-01, ID-01, JP-01, KR-01 and the generic US block — is registered to an Amazon entity. The check covered every instance for which Braze publishes a list. INFRASTRUCTURE

- Certificate transparency shows 50 hosts on region codes `p-aze-us`, `s-aze-us` and

`d-aze-us` — a code matching no AWS region identifier, where every other code in the estate does. Hosts on it include `sdk-us08`, `subcenter-08` and `itp-api-08`, tying the code to instance 08. INFRASTRUCTURE

`source: sources/external/rdap-instance-ip-ownership_2026-09-02.json @ 2026-09-02`

`source: sources/docs/docs__user_guide__data__unification__cloud_ingestion__connected_sources.md:239 @ 2026-09-02`

State the three observations and stop. A hosting arrangement may sit outside a sub-processor listing for reasons not visible from outside, the document may be behind the estate, or there may be an explanation nobody external can see. Asserting a disclosure failure would be a legal conclusion this evidence does not support. The question — why the disclosure and the infrastructure differ — is recorded as open question 52 and as conflict **C-03**.

Note also that the certificate-transparency host list is **partial**: crt.sh returned HTTP 502 throughout and the Cert Spotter fallback was rate-limited without a token. No claim is made here that anything is *absent* from CT, because the list was not exhaustive.

4.6 · Identity, and where it is quietly lossy

FACT	VALUE	GRADE	SOURCE
Aliases per profile	No limit	DOCUMENTED	sources/docs/docs__user_guide__data__unification__user_data
Aliases per label	One, and unique across the user base	DOCUMENTED	sources/docs/docs__api__endpoints__user_data__post_user_ide
Identifiers accepted on ingest	Five: external_id, user alias, braze_id, email, phone	DOCUMENTED	…compare_data_ingestion_options.md:113-118
Identifiers accepted for warehouse segmentation	One: `external_user_id` only	DOCUMENTED	…compare_data_ingestion_options.md:113-118

The identity model is generous at the top and narrow at the bottom: unlimited aliases on a profile, but a warehouse-driven segment can only be built on a single identifier type.

Two admitted behaviours deserve a buyer's attention because neither is visible from marketing and both fail quietly.

A merge can decline and still report success. If both profiles carry invalid phone numbers Braze does not merge them, and their documentation says so directly: "The endpoint still returns 202 Accepted with a success message, so the HTTP response does not indicate that the merge was skipped."

Reporting splits across surfaces after a merge. Dashboard campaign summaries attribute a pre-merge send to the surviving profile; Currents, Query Builder and Messaging History still attribute it to the orphaned profile's ID. Both are correct by their own rules and they do not agree, which is a reconciliation problem for anyone joining Braze data to a warehouse. Orphaned profiles are also not eligible to receive messages.

`source: sources/docs/docs__api__endpoints__user_data__post_users_merge.md:178 @ 2026-09-02`

4.7 · What a customer is billed for

The billable unit is defined in the documentation with contract-grade precision: a "data point" is "a billable unit of use of the Braze Services, measured by a session start, session end, custom event, or purchase recorded, as well as any attribute set on an end user profile", and each such item "shall each count as a single data point". A session start and a session end are two.

Engagement is explicitly *not* billed — push tokens, device information, and all campaign engagement tracking such as email opens and push clicks are excluded, as are user deletion, Connected Content, subscription-state changes and external-ID renames. That is a genuinely customer-friendly boundary and belongs in any fair account of the pricing.

The tension is in the advice Braze gives alongside it: **"Don't waste data points. Only update changing data!"**, and "we recommend setting up a program to prevent sending the same unchanging data." A platform positioned on continuous streaming of behavioural data prices in a way that makes its customers build programmes to send less of it. Both facts are theirs; the tension between them is the finding, and it is corroborated from an unrelated direction by a Gartner reviewer citing "steep pricing and strict data-point limits, where overages can become expensive very quickly" (chapter 7).

`source: sources/docs/docs__user_guide__data__infrastructure__data_points.md:18,20,36-38 @ 2026-09-02`

4.8 · A decade of reliability, and an engineering record

FACT	VALUE	GRADE	SOURCE
Incidents published	451, 2016-10-09 → 2026-08-05	DOCUMENTED	`data/incidents.csv`
Median duration	79 minutes	DOCUMENTED	`data/incidents.csv`
p90 duration	311 minutes	DOCUMENTED	`data/incidents.csv`
Major or critical	98 of 331 non-maintenance incidents — 29.6%	DOCUMENTED	`data/incidents.csv`
Public issues captured	845, all from non-members of the org	DOCUMENTED	`data/issues.csv`
Issues closed	822 — 97.3% — median 11 days	DOCUMENTED	`data/issues.csv`

Incidents peaked in 2023 at 60 and fell to 27 in 2025, the quietest full year on the status page, over a period when revenue grew 7.7 times. Incident rate per unit of scale has fallen sharply, which kills the hypothesis that it would rise. The caveat travels with it: 2026 stands at 35 through August, roughly double the 2025 monthly rate.

The component distribution has a shape worth naming. The Dashboard appears in 63 incidents, more than any other component and more than twice Outbound Messaging's 27. The control plane that marketers use fails more often than the sending path.

None of this may be compared against a vendor that publishes no status page. A decade of visible incidents is a disclosure practice, not a defect count, and a competitor with silence has not thereby earned a better record.

On the engineering side the record is straightforwardly good. All nine SDK repositories that publish releases shipped within thirteen days of capture. The one platform whose repository is archived — Unreal — has had its documentation removed entirely, so the marketing and the maintenance agree. The single soft spot is `braze-roku-`

`sdk`: 181 days since its last push while 38 Roku documentation pages remain live. And 845 unsolicited public issues closed at 97.3% with a median of 11 days is a real operational number that Braze does not control the denominator of.

The issue corpus describes the SDK surface, not the dashboard, and its authors are developers rather than buyers. It is not a satisfaction measure and is not used as one.

What would change this chapter

A revised sub-processor list naming Microsoft would resolve §4.5 in one direction; a statement that US-08 is something other than an Azure deployment would resolve it in another. A change to the documented rate limits, or the removal of the two-tier export limit, would date §4.3. Two consecutive quiet quarters would make the 2026 incident uptick noise; two more like April 2026 would make it a trend.

5 · Channels and delivery

> Thirteen channels are documented, which is broad rather than narrow — the hypothesis that Braze's channel list would prove thinner than its positioning was killed by the evidence. What the roster does show is a documentation site and a marketing site that have drifted apart in both directions at once, and a compelled disclosure showing that only two of the thirteen channels have a named delivery intermediary at all.

5.1 · The roster, counted from both ends

Braze's own documentation channel index lists twelve:

In-product — in-app messages, Content Cards, Banners. **Out-of-product** — Email, Transactional email, Landing pages, LINE, Live notifications, Push, SMS/MMS/RCS, Webhooks, WhatsApp.

DOCUMENTED

`source: sources/docs/docs__user_guide__channels.md:14-64 @ 2026-09-02`

A thirteenth exists and is not on that list. **KakaoTalk** has four documentation pages — setup, message creation, click tracking and reporting, which is a fully documented channel rather than a stub (`sources/docs/docs__user_guide__channels__kakaotalk__*.md`) — and a marketing product page at `/product/kakaotalk-messenger`, in `data/site_inventory.csv`. It is mentioned zero times on Braze's own channels index.

The drift runs the other way too. Five documented channels have no dedicated marketing product page anywhere on the site: **Banners, Transactional email, Landing pages, Live notifications and Webhooks**. The largest of them is not small — Landing pages carries eleven user-guide pages plus three partner pages (`sources/docs/docs__user_guide__messaging__landing_pages*.md`), a substantial documented capability sold under no name.

Why the drift is the finding, not the count

The honest answer to "how many channels does Braze have" is three different numbers depending on which document you trust, which is why this is conflict **C-02**. But the count is the least interesting thing here.

A documentation site and a marketing site are maintained by different people on different cadences, and the gap between them is a reliable measure of which parts of a product the company is currently *selling* versus which parts it merely *supports*. On that reading:

- **Landing pages, Banners, Transactional email, Live notifications and Webhooks are

supported but not sold.** Eleven documentation pages exist for landing pages because customers use them and need help; no marketing page exists because landing pages are not how Braze wins a deal against Salesforce. These are retention features, not acquisition features, and their documentation-to-marketing ratio says so.

- **KakaoTalk is sold but not indexed**, which is the rarer direction and more likely an

oversight than a strategy — a channel added to the marketing taxonomy and missed when the docs index was last revised.

Investment also differs sharply by channel in a way the roster hides. `/product/line` exists in four languages; `/product/kakaotalk-messenger` in two. Both are East Asian messaging channels; one has had twice the localisation attention. For anyone competing in Korea specifically, that ratio is a more honest signal of commitment than either channel's presence on a list.

`source: data/site_inventory.csv @ 2026-09-01`

The hypothesis this killed

The analysis went in expecting the channel list to be narrower than the positioning implied — the usual finding when a vendor claims omnichannel. It is not. Thirteen documented channels including WhatsApp, LINE, KakaoTalk, RCS, Banners, Landing pages, Live notifications and Webhooks is a genuinely broad set, broader than several of the specialists buyers shortlist Braze against (chapter 7). **Recorded as killed, in Braze's favour**, because a hypothesis quietly dropped is a bias.

The roster measured, not just listed

A channel list says nothing about how much of each channel exists. The documentation corpus can be counted the same way chapter 6 counts the AI, and the result orders the roster by how much Braze has had to explain.

CHANNEL	FOCUSED DOC PAGES	API ENDPOINTS
Email	347	16
In-app messages	115	0
SMS / MMS / RCS	89	2
Push (mobile)	73	1
Webhooks	71	2
WhatsApp	56	0
Content Cards	47	0
Banners	28	0
Landing pages	21	0
LINE	11	0
Web push	11	0
Live notifications	8	2
KakaoTalk	5	0
Transactional email	4	1

DOCUMENTED

`source: data/capabilities.csv @ 2026-09-02`

Three readings come out of this table that the roster alone cannot give you.

The distribution is extremely uneven. Email at 347 focused pages is more than the next three channels combined, and roughly seventy times KakaoTalk's five. "Thirteen channels" is true and it describes a set in which one channel is the product and twelve are completeness.

Depth tracks marketing, with two exceptions that go opposite ways. Broadly the best-documented channels are the marketed ones — but **Webhooks at 71 pages is the fifth deepest channel in the platform and has no marketing page at all**, while **KakaoTalk at 5 pages has one.** Webhooks is a genuinely substantial, genuinely unmarketed capability; KakaoTalk is marketed at roughly the documentation depth of a stub. If you are judging what Braze can actually do from its product pages, those are the two places you would be most wrong in each direction.

Channel breadth is not channel parity, and nobody should present it as such. A prospect in Korea evaluating KakaoTalk support is looking at five focused pages and no API endpoints. A prospect asking about LINE is looking at eleven. Those channels exist and are real, and they are not comparable to what an email or push buyer receives. The fair sentence is that Braze covers thirteen channels and covers three of them deeply.

5.2 · Who actually delivers each channel

The sub-processor disclosure is legally compelled to be complete, which is what gives it force. Revision 1 June 2026 names 17 third-party sub-processors.

CHANNEL	NAMED DELIVERY SUB-PROCESSOR
Email	Three — Amazon SES, Bird.com (SparkPost), Twilio (SendGrid)
SMS / mobile messages	Two — Infobip, Twilio
Everything else	None named

DOCUMENTED

`source: sources/clean/braze-subprocessors.md:21-37 @ 2026-09-02`

Email has three delivery providers named, plus a fourth supplier — Mailgun — for Email on Acid previewing (`sources/clean/braze-subprocessors.md:21-37`). Mobile messages have two. **Push, in-app messages, Content Cards, Banners, Webhooks, WhatsApp, LINE, KakaoTalk, Landing pages and Live notifications have no delivery sub-processor named at all.**

What that absence does and does not mean

The caveat travels in the same breath as the finding, because the finding is easy to misuse. Absence from a sub-processor list is **not** proof of no intermediary. Platform transports such as APNs and FCM, and the WhatsApp, LINE and Kakao business APIs, may not be classified as sub-processors processing personal data on Braze's behalf — a message handed to Apple for delivery to a device may sit outside the definition entirely. What can be said is what the compelled disclosure names, and it names middlemen for email and SMS and for nothing else.

Two conclusions do survive that caveat, and both are commercially useful.

Email deliverability is not a single point of failure. Three interchangeable providers on the highest-volume channel is deliberate redundancy, and it is expensive to maintain — three sets of IP reputation, three integrations, three contracts. A competitor hoping to attack Braze on email reliability is attacking the part of the estate that has been engineered hardest against exactly that. This is a favourable finding and it comes from a document Braze had no marketing reason to publish.

The in-product channels are genuinely Braze's own. In-app messages, Content Cards and Banners are rendered by Braze's SDK inside the customer's own application. There is no carrier, no inbox provider, and no third party in the path — which is why no sub-processor is named for them, and why those channels are the ones least exposed to a supplier relationship going wrong. It is also why they are the channels where Braze's latency story is strongest, since nothing leaves the device-to-Braze path.

5.3 · The rest of the supplier list, read for what it discloses

Beyond delivery, the same document discloses things no marketing page would.

- **Two hosting providers named** — Amazon Web Services, and Google LLC for Google Cloud

Platform. What that omits is chapter 4's subject and the most significant single finding in this analysis.

- **End-user profiles are stored by a third party.** Rackspace US, Inc. provides "Database

Administration as a Service (DBaaS), a managed database service provider that **hosts and stores End User profiles**." A buyer doing a security review should know that the profile store is operated by a managed-database vendor rather than by Braze directly — and the 10-K independently confirms it, naming Rackspace alongside AWS as infrastructure (`sources/filings/2026-03-25_10-K_000013.txt`) Braze relies on.

- **Monitoring receives user identifiers.** "Braze may provide End User metadata, such as

user identifiers, to DataDog for support and application troubleshooting." That is ordinary practice, disclosed properly, and it is the kind of line a data-protection officer will ask about.

- **Three foundation-model suppliers** — Anthropic, OpenAI and Google — dealt with in

chapter 6.

- Databricks, dbt Labs, Domino Data Lab, ClickHouse, Snowflake, Cloudflare and Fastly make

up the remainder: analytics, transformation, CDN and traffic management.

Braze separately claims "over 150 technology partners, which we call 'Alloys'". That is a marketing figure appearing inside technical documentation and is graded accordingly. CLAIMED
`source: sources/docs/docs__developer_guide__getting_started__architecture_overview.md:82 @ 2026-09-02`

The shape of the integration layer, measured

The partner network can be measured rather than accepted. The documentation corpus carries **322 partner pages** — the second-largest section by page count — averaging about 810 words each, against a developer guide that runs more than three times as many words per page across a third as many pages (chapter 4). Both counts come from `data/docs_sections.csv`, produced by `tools/index_docs.py`.

That is the documentation profile of a broad network described one short page at a time. It is not a criticism: a connector may need only a short page, and 322 of them is a real network that takes real effort to maintain. But it does mean the integration layer is **wide and shallow in documentation terms**, and anyone told the integration story is deep should measure it this way rather than counting logos. The depth is in the SDKs and the API, which is where the words actually are.

What would change this chapter

A revision of the sub-processor document — this record uses revision 1 June 2026 and any later revision supersedes it entirely, particularly if it adds or removes a delivery provider. A new channel appearing in the docs index, or KakaoTalk being added to it, which would close half of C-02. A marketing page appearing for any of the five unmarketed channels, which would tell you Braze had decided to start selling one of them. And any change to the three-provider email arrangement, which would be visible in the same disclosure and would matter more than it sounds.

6 · The AI

> Five independent lenses point the same way: the decisioning engine was bought rather than built, the models come from three named external suppliers, the AI surface carries no API of its own, and its documentation footprint is a fraction of the platform's established capabilities. None of that makes it thin — it makes it *new*, *supplied*, and *measurable*, which is a different and far more useful claim than the one a competitor would prefer to make.

The rule this chapter follows, carried from the reference project: **never assert thinness.** State the shipped machine learning and the agentic layer as separate, counted things, and let the ratio speak. A claim of this kind is only defensible when several unrelated lenses agree, so five were used and each is recorded separately below. If four had agreed and one had not, that would be in here too.

6.1 · Lens one — how much documentation exists

Focused pages, meaning pages where the capability's vocabulary appears at least five times in the body — a threshold that separates "mentions it" from "is about it". Counted against the documentation corpus described in chapter 4.

CAPABILITY	FOCUSED PAGES
Recommendations	73
BrazeAI (umbrella brand)	32
Generative AI	26
Decisioning Studio	22
Agents	17
Intelligent Timing / Channel	13
Predictive Suite	7
— for comparison —	
Email	347
Canvas (journeys)	249
Segmentation	242
Liquid templating	123

DOCUMENTED

`source: data/capabilities.csv @ 2026-09-02`

The pattern set behind these counts is a judgement and is published as such, in `docs/CAPABILITY-TAXONOMY.tsv` and in the `pattern` column of `data/capabilities.csv`, precisely so a reader who disagrees can change it and re-run. It was revised on 2026-09-02 to use Braze's own product names rather than generic category words; counts from before that revision are not comparable with these, and the change is logged in the corrections table (chapter 8).

How to read a documentation count without over-reading it

Documentation volume measures what a vendor has had to *explain*, which is a proxy for maturity and surface area, not for quality. Three caveats keep it honest.

A new feature is under-documented by construction. Predictive Suite at seven focused pages could mean an immature product or a simple one, and the count cannot distinguish those. Decisioning Studio at 22 pages is fifteen months old inside Braze; Canvas at 249 has had a decade. Comparing them directly and concluding Canvas is "35 times better" would be nonsense.

Some capabilities are inherently short to document. Intelligent Timing is a toggle with a model behind it. There is not much to say, and thirteen pages may be complete.

But a floor still exists. A product marketed as a pillar of the platform, with its own top-level navigation and its own trademark, that carries fewer focused pages than Connected Content (43) or Content Blocks (38) is telling you something about how much surface it currently exposes to customers. Both comparators sit in the same file and were counted by the same rule — `data/capabilities.csv`. That is the finding — not thinness, but *surface area*, which is measurable and which the next four lenses corroborate.

6.2 · Lens two – the published API surface

Of Braze's 135 documented REST endpoints, spread across 28 top-level namespaces, the number sitting in an AI, ML, prediction, recommendation, agent or decisioning namespace is **zero**. The check was exhaustive over every documented path.

The namespaces are: `app_group`, `apps`, `campaigns`, `canvas`, `catalogs`, `cdi`, `content_blocks`, `custom_attributes`, `data_objects`, `email`, `events`, `export`, `fcm`, `kpi`, `media_library`, `messages`, `preference_center`, `purchases`, `scim`, `segments`, `sends`, `sessions`, `sms`, `subscription`, `templates`, `transactional`, `users`, `v2`.

For contrast, Canvas has 10 endpoints, Catalogs 9, subscription management 9, identity 8, email 16.

DOCUMENTED

`source: data/api_endpoints.csv @ 2026-09-02`

This is the lens that matters most to a technical buyer, and its meaning is specific. A capability with no API surface is not necessarily immature — it may be deliberately dashboard-only, and plenty of good features are. But it has three practical consequences that a prospect will hit:

- **The AI features cannot be orchestrated from outside the product.** There is no way to

trigger a prediction, retrieve a recommendation, or configure an agent programmatically.

- **They cannot be tested in a CI pipeline** the way a campaign trigger can.
- **They cannot be composed with a customer's own models.** Everything must happen inside

Braze's interface, on Braze's terms.

For an organisation that already runs its own data science, that is the difference between a platform capability and a black box. It is also reversible — an API namespace could appear in any release — which is exactly why this chapter's closing section names it as the thing to watch.

6.3 · Lens three – who supplies the models

The sub-processor disclosure names three external providers, each with the same description: "provides artificial intelligence models and machine learning infrastructure to support data processing and advanced platform features."

- **Anthropic PBC**
- **OpenAI OpCo, LLC**
- **Google LLC**

DOCUMENTED

source: `sources/clean/braze-subprocessors.md:22,31,34 @ 2026-09-02`

Using foundation models from suppliers is ordinary and not a criticism — very few companies train their own, and choosing three rather than one is a sensible hedge against any single provider's pricing, availability or policy changes. What is unusual is that Braze is legally obliged to name them and therefore does, where a private vendor's marketing never would.

The consequence for a competitive conversation is narrow and firm: **any claim that Braze's AI is proprietary at the model layer is contradicted by Braze's own compelled disclosure.** What is proprietary is the integration, the data model the models see, and the decisioning logic acquired with OfferFit. That is a real asset and this chapter does not diminish it. It is simply not a model advantage, and it is available to anyone else willing to write the same integrations.

There is a second-order point worth noting: three model suppliers is also a **cost** line, and it appears in the only place the filings leave unquantified — "increased costs related to our tech stack" at `sources/filings/2026-03-25_10-K_000013.txt:1565`, the unsplit half of the FY2026 gross-margin decline (§2.1). Whether AI inference is a material and growing share of cost of revenue is the single most useful number Braze does not disclose.

6.4 · Lens four — the engine itself was bought

The decisioning product is OfferFit, acquired 2 June 2025 for $303.2m and renamed AI Decisioning Studio, as chapter 3 sets out from the auditor's own note at `sources/filings/2026-03-25_10-K_000013.txt:1788`. The agentic layer Braze markets is, in the 10-K's words, "OfferFit's multi-agent decisioning engine".

AUDITED

source: `sources/filings/2026-03-25_10-K_000013.txt:3130 @ 2026-03-25`

That is the fourth lens and the one that dates the other four: whatever the documentation footprint says, the engine is fifteen months old *inside Braze*. It is older than that as software — OfferFit existed as a company — but the integration, the data plumbing into Braze's profile store, and the documentation are all fifteen months old at most. A twenty-two-page documentation footprint for a product of that age is unremarkable. It would be remarkable at five years.

This is why the counts in §6.1 are presented as surface area rather than as a verdict. Read together with the acquisition date, they describe a product that is **early**, not one that is hollow — and early products get documented, so the honest expectation is that these numbers rise. That is a falsifiable prediction and re-running `capability_count.py` in ninety days tests it: the script is `tools/capability_count.py` and its output is `data/capabilities.csv`.

6.5 · Lens five — the words customers use

Not planned as a lens, but the review corpus turned out to say something specific enough to include. Across the captured panels, AI is named thirteen times.

What the positives name: "AI Advisor... a strong addition, providing robust capabilities and useful insights"; "SQL Query Builder with AI assistant"; "the AI personalization tool".

What the criticisms name: two G2 reviewers independently single out the **AI copywriting assistant** as needing "a careful human hand" before anything reaches a customer; a TrustRadius reviewer says "AI features feel half baked" (`sources/panels/trustradius.md`); and TrustRadius's own synthesis records that "the platform's cost, particularly for advanced AI features, is a recurring consideration".

THIRD-PARTY

`source: sources/panels/g2.md:137,189 @ 2026-09-02`

The distribution is the interesting part. **Every AI comment in the captured sample is about generative assistance — copywriting, query building, personalisation. Not one mentions decisioning, agents or prediction.** That is exactly what you would expect fifteen months after an acquisition whose product is still being integrated, and it is consistent with lens four rather than an independent confirmation of it.

The sample is small — fourteen review bodies, coded in `data/review_coding.csv` — and is used here for vocabulary only, never for a percentage. The denominator discipline is set out in chapter 7 and in conflict **C-06**.

6.6 · What the five lenses together support

They support this and nothing stronger: **Braze's AI is real, shipping, recently and largely acquired, running on models from three named external suppliers, without an API surface of its own, and documented at a fraction of the depth of the platform's established capabilities.** Every clause there is sourced separately above, from a different kind of source, and no clause depends on another.

They do not support "their AI is thin". That sentence is unfalsifiable, the evidence does not need it, and a competitor who says it will be contradicted the first time a prospect opens the product and sees a working decisioning engine. The sourced version survives that demonstration; the unfalsifiable one does not.

The genuinely useful question a prospect can ask, which follows directly from lens two, is: *"Can I call your AI from my own systems?"* Today the documented answer is no.

What would change this chapter

An AI namespace appearing in the REST API — the single most consequential change, and the one that would invalidate lens two outright. A jump in the Decisioning Studio or Agents page counts on a re-run of `capability_count.py`; the measurement is repeatable, which is the point of publishing the pattern set. A sub-processor revision adding or removing a model provider, which would say something about cost or about a shift to in-house models. Review vocabulary beginning to name decisioning rather than copywriting, which would show the acquired engine reaching customers. And any disclosure separating AI inference cost from the rest of cost of revenue.

7 · The market

> Braze sells to 2,609 customers where a "customer" is an ultimate parent entity, no single one of which is 10% of revenue. Expansion within that base is slowing and the premium large accounts used to enjoy has nearly vanished. And the vendors buyers actually weigh Braze against are twice as many as the four it names — with the five extra all being specialists it does not mention.

7.1 · Three customer counts that must never be merged

ROSTER	COUNT	GRADE	SOURCE
The 10-K's defined metric	2,609 as at 31 January 2026	AUDITED	`sources/filings/2026-03-25_10-K_000013.txt:1379`
Self-published customer stories	178 `customers/` URLs	CLAIMED	`data/site_inventory.csv`
Independent detection	Not attempted in this run	—	open question 57

The first is a defined, audited metric. The second is a marketing selection. The third was not done, and is recorded as a gap rather than as a zero — nobody should read "178" as the customer base or as a sample of it.

What the 10-K metric actually counts, and why it changes the price

Braze defines the metric precisely, and the definition does more work than the number: **"We define a customer as the separate and distinct, ultimate parent-level entity that has an active subscription with us to use our products."**

AUDITED

`source: sources/filings/2026-03-25_10-K_000013.txt:1381 @ 2026-03-25`

Ultimate parent-level entity. A holding company running ten retail brands on Braze counts once. A conglomerate with separate contracts for three divisions counts once. This is a deliberately conservative way to count and it flatters nothing — but it has a direct consequence for the only pricing figure this analysis can produce.

Revenue divided by customers is **about $283,000.** That is a bound, not a price, and the definition tells you what kind of bound: it is **per corporate group**, not per brand, per seat or per contract. A single-brand startup and a group running ten brands under one master agreement are both one unit in that denominator. The real distribution behind $283,000 is therefore far wider than the average suggests, and skewed by construction — which is why the 333 customers at $500k+ of annual recurring revenue is the more useful segmentation, and why "bounded at" is the only honest verb. This is conflict **C-08**.

Two further disciplines. **No transacted price appears anywhere in the corpus** — no rate card, no procurement award, no quoted contract value in any review, filing or documentation page. The search was exhaustive and found nothing, so the bound stands alone and unsupported by any real observed price. And the 178 customer stories are brands, not parents, so they are not a subset of 2,609 in any arithmetic sense.

The one genuinely reassuring number in this chapter

No customer accounted for 10% or more of total revenue in FY2026 or FY2025.

AUDITED

`source: sources/filings/2026-03-25_10-K_000013.txt:2173 @ 2026-03-25`

This is a required disclosure and its absence of drama is the point. A business with 2,609 parent-level customers and no concentration above 10% has no single account whose loss would be visible in the consolidated numbers. For anyone modelling competitive risk against Braze, there is no keystone logo to unseat — the base has to be taken account by account, which is slow, and which the contracted-revenue position in §2.4 makes slower still.

7.2 · Expansion is slowing, and the enterprise premium has almost gone

MEASURE	FY2024	FY2025	FY2026
Dollar-based net retention, all customers	117%	111%	109%
Dollar-based net retention, customers ≥$500k ARR	120%	114%	110%
Customers with ARR ≥$500k	202	247	333
Total customers	2,044	2,296	2,609

AUDITED

`source: sources/filings/2026-03-25_10-K_000013.txt:1400,1379 @ 2026-03-25`

Three movements here, and the second and third are easy to miss.

Net retention is falling, and Braze explains why itself (`sources/filings/2026-03-25_10-K_000013.txt:1400`): "primarily due to customer turnover and renewals at lower subscription levels... customers renew their contracts at levels more closely aligned with their current needs, rather than opting for larger commitments based on anticipated future demand." That is a demand-environment explanation, it is the company's own, and it is plausible — every subscription vendor said a version of it over this period. It is worth noting what it is *not*: it is not an explanation involving competitive loss, and Braze does not offer one anywhere in `sources/filings/2026-03-25_10-K_000013.txt`.

The enterprise premium has nearly closed. Large customers used to expand three points faster than the average — 120% against 117% in FY2024. In FY2026 they expand one point faster, 110% against 109%. Whatever advantage the large-account cohort had in expansion has almost gone. This is the subtler and more consequential movement, because expansion within large accounts is where a platform business earns its multiple, and it is the mechanism by which land-and-expand is supposed to work.

And yet the number of large accounts is growing fast — 202 to 333 in two years, a rise of about 65% while the total customer count grew about 28%. So Braze is winning substantially more big accounts and growing each of them more slowly.

Those three facts together describe a specific transition, and naming it is more useful than any one of them alone: **Braze is shifting from an expansion business to an acquisition business.** Growth increasingly comes from new large logos rather than from existing ones getting larger. That is a harder, more expensive way to grow — it is consistent with a sales-and-marketing line that, despite falling every year, is still the largest single cost in the business (§2.3) — and it is the pattern that makes the RPO figure in §2.4 the right thing to watch. Retention above 100% still means the base grows without a single new customer; but at 109% it grows slowly, and the burden shifts to the new-logo engine.

Note also what these figures cannot tell you. Dollar-based net retention is not satisfaction, and it is not a competitive win-rate. It nets expansion against contraction and churn across a whole cohort, so a company could lose accounts to a competitor and still report 109% if the survivors expanded enough. It is the closest audited proxy available for the health of the base, and it is a proxy.

7.3 · Who they name, and who buyers actually shortlist

The 10-K names four competitors: **Adobe, Salesforce, Iterable, Klaviyo**, alongside the claim that "none of our competitors currently offer comparable comprehensive customer engagement solutions" — `sources/filings/2026-03-25_10-K_000013.txt`.

Gartner Peer Insights publishes what buyers *also considered*, derived from the buyers rather than the vendor. That list has eight: **Salesforce, Adobe, Iterable, Oracle, Optimove, Blueshift, MoEngage, CleverTap**.

	NAMED BY BRAZE	ON THE BUYER SHORTLIST
Adobe, Salesforce, Iterable	yes	yes
Klaviyo	yes	no
Oracle, Optimove, Blueshift, MoEngage, CleverTap	no	yes

THIRD-PARTY

`source: sources/panels/gartner.md:113-121 @ 2026-09-02`

The overlap is real — three of the four names are shared, so this is not a vendor misreading its own market, and the hypothesis that Braze would be shortlisted against entirely different vendors was only **partially** borne out. That partial result is the honest one and it is recorded as such.

The asymmetry is what matters, and it runs one way. Buyers weigh Braze against five vendors it never names, and the character of those five is the finding: **Oracle is a suite vendor, but Optimove, Blueshift, MoEngage and CleverTap are mobile-engagement and CDP specialists.** Braze's own competitive frame is "we are more comprehensive than the suites"; the buyers' frame includes "is this better than a focused specialist". Those require different answers, and a competitive briefing built only from Braze's comparison pages would prepare for the first conversation and walk into the second.

Where Gartner's reviewers rate Braze *above* the two largest alternatives, they name the same things twice: service and support, and ease of integration and deployment — with evaluation and contracting added against Salesforce. That is a consistent, independent, buyer-sourced statement that Braze is easier to buy from and easier to deploy than Adobe or Salesforce, and it should be quoted as readily as anything critical in this record.

7.4 · What customers say

Site-level aggregates, which are the figures to quote. The captured review bodies are a sample of fourteen records used for vocabulary and theme, never for a percentage.

PANEL	RATING	BASE
G2	4.5 / 5	1,702 reviews
Gartner Peer Insights	4.5 / 5	263 ratings
TrustRadius	8.8 / 10	348 reviews

THIRD-PARTY

`source: sources/panels/g2.md:46 @ 2026-09-02`

Braze is well liked, and that is the first thing to say. Two panels at 4.5 out of 5 and a third at 8.8 out of 10, across more than two thousand ratings, is a strong result that no amount of subsequent criticism should be allowed to obscure.

The bases themselves disagree, which is conflict **C-06**: Glassdoor's search page says 563 reviews where its company page says 524 ratings; Gartner shows 263 Ratings and 267 Verified Reviews; TrustRadius shows 348 in the header and 162 under the default filter. The gaps are small on Gartner and very large on TrustRadius, where 186 of 348 are excluded by the site's own default view. The ruling is to quote each site's headline rating with its own headline base and never to mix them — and this is the second reason, after sample size, that no percentage in this analysis is computed from captured review bodies.

G2's own coded tags (`sources/panels/g2.md:46`), counted over its full review base rather than over the sample captured here, put the praise at Ease of Use (385), Intuitive (188), Customer Support (151), Helpful (148) and Features (138); and the criticism at Missing Features (140), Learning Curve (139), Limitations (102), Limited Features (93) and Steep Learning Curve (86). The most common praise tag is nearly three times the most common criticism, which is the ratio to carry away.

Two themes recur across unrelated panels and are therefore worth more than the rest: **reporting is the most consistent complaint**, and **the learning curve is the second.** TrustRadius codes reporting as "limited and unintuitive, a sentiment shared by 36% of reviewers" (`sources/panels/trustradius.md`).

The reporting complaint corroborates independently with the platform evidence in chapter 4, and the corroboration is what makes it solid rather than anecdotal: out-of-the-box reporting is thin enough that Currents — a paid add-on exporting on a five-minute cadence — is the route to raw data. A complaint from buyers and a mechanism in the documentation, arrived at from unrelated directions, describing the same thing.

Running `tools/code_reviews.py` over the whole corpus produces theme percentages against a denominator of 860 records in `data/review_themes.csv`, but 841 of those are GitHub issues from `data/issues.csv` rather than buyer reviews. That denominator is stated wherever the script's output is used, and its themes describe the SDK surface, not satisfaction.

7.5 · What employees say

Glassdoor, captured signed-in: **4.1 out of 5 from 524 ratings, 82% would recommend, 90% approve of the CEO, 71% positive business outlook.**

THIRD-PARTY

source: sources/panels/glassdoor.md:68-92 @ 2026-09-02

Work-life balance tracks the overall rating almost exactly over the six months to September 2026 — both hovering around 4.0 to 4.1 — so there is no divergence worth calling a trend, and none is claimed. The analysis went looking for a work-life-balance decline, which is the usual story at this stage of a company's life, and did not find one.

Glassdoor's own summary names the weak spots as management effectiveness and "clarity in direction", and "limited upward mobility and discrepancies in compensation relative to market rates". That last phrase is worth setting against §1.4: the proxy discloses a median employee total compensation of $164,000 (`sources/filings/2026-05-18_DEF-14A_021908.txt:5769`). This is not a low-paying company, so the grievance is better read as being about *progression* than about level — which is precisely what "limited upward mobility" says, and what the "limit to the number of folks who can score a 4/5" comment in the captured reviews describes.

What they are hiring for, which is the forward-looking half

The careers board carries **296 open roles across fifteen hiring departments**, of twenty-one listed.

DEPARTMENT	OPEN ROLES
Sales	89
Engineering	57
Customer Experience	38
Marketing	24
Solutions Consulting	23
Partnerships	16
Growth	15
Information Technology	13
People	7
Business Development	5
Finance	4
GTM Strategy, GTM Operations, Legal, Product	1–2 each

DOCUMENTED

`source: data/careers_departments.csv @ 2026-09-02`

The shape is unambiguous. **Go-to-market functions account for 213 of the 296 roles — 72.0% — against 58 for engineering and product, 19.6%.** Braze is hiring go-to-market people over product-builders by roughly **3.7 to one**.

Read that against §7.2 and it stops being a personnel statistic. The retention figures say expansion within the existing base is slowing and that growth is shifting to new logos; a hiring board that is nearly three-quarters go-to-market is what executing that shift looks like from the outside, and it is a leading indicator where retention is a lagging one. The single largest department is Sales at 89 roles, more than engineering and product combined.

Two disciplines on how far that can be pushed. **A requisition is an intention, not a person** — a board is a plan, and plans get cut; nothing here contradicts or updates the audited employee counts in §1.3. And **this is not a spend ratio**: a sales requisition costs materially less than a senior engineering one, so 3.7:1 in headcount is consistent with the audited 1.96:1 in money (§2.3) rather than in conflict with it. The two point the same way at different magnitudes, which is what corroboration looks like when the measures are genuinely different.

This closes open question 56. The first pass recorded the split as uncapturable, because the board's Department filter would not drive under automation; the same board is published as unauthenticated JSON with the grouping already done. The correction is in chapter 8's log, and the transferable lesson is there too: **when a page will not yield, look for the API behind it before recording a gap.**

Product at a single open role is the one number here worth not over-reading. One requisition in a department is as likely to mean the team is fully staffed as anything else, and a fifteen-person product organisation with no vacancy would show exactly this.

What would change this chapter

A fourth consecutive year of falling net retention would turn a demand-environment explanation into a competitive one, and would be the single most important number in the next 10-K. A reversal in the large-account premium would say the expansion engine had restarted. A new name in the 10-K competition paragraph, or a change in the Gartner shortlist, would move §7.3. And access to any of the three paywalled by-company-size rating breakdowns would finally settle whether satisfaction falls with customer size — the one hypothesis this analysis could not test at all.

8 · What we could not answer

> Eight questions this analysis could not close from public sources, each with the source that would close it — plus the two places where a number moved and the old value is deliberately still visible, and the ledger of ten hypotheses graded honestly, three of which died in Braze's favour.

A gap written down is evidence. A gap not written down is a mistake. The full backlog lives in `docs/QUESTIONS.md`; what follows is the residue that matters to a reader of this record, in the order a competitor would care about.

8.0a · The four documents the deck rests on

Slide 5 exists because the operator watched the deck and could not follow it. Seventeen specialist terms were in use and none was defined; two of them — "10-K" and "sub-processor" — carried whole findings, and one of those landed on the title slide. The slide is not a glossary. It is the argument for why these four sources beat marketing, with the definitions riding along inside it.

DOCUMENT	WHAT IT IS	WHY IT IS HARD TO FALSIFY
10-K	The audited annual report to the SEC	Signed by the CEO and CFO, audited by Ernst & Young, wrong at legal risk. Seven years captured
DEF 14A (proxy)	Filed before the annual shareholder meeting	Must name every executive officer and director, their pay, and who owns the company
Sub-processor disclosure	A public list of outside suppliers that touch customer data	Data-protection law requires it to be *complete*, so it names middlemen marketing never would
Status page	A live record of every incident since 2016	Written during the outage, by someone trying to fix it

Nothing on that slide is a new finding, and it cites nothing this record does not already carry: the filings are in §1 and §2, the sub-processor disclosure runs through chapters 5 and 6, and the status page is chapter 4's operational record. **What it adds is the reason to believe the rest**, delivered before the first claim rather than after it.

The general lesson is worth recording next to the corrections log, because it is the same class of error. **A term the analyst has read four hundred times is invisible to them and opaque to everyone else.** Nothing in the evidence was wrong; the deck was simply unreadable to the person it was built for, and no amount of verification catches that. Only showing it to somebody does.

8.0 · The hypothesis ledger, and why three killed results are on a slide

Ten hypotheses were written into `docs/STRATEGY.md` **before any source was captured**, each with a stated condition that would kill it. That ordering is the whole point: a hypothesis written after the reading is a conclusion wearing a question's clothes, and a hypothesis quietly dropped when it fails is a bias that leaves no trace.

They ended: **four evidenced, four killed, one partially evidenced, one untestable.**

The distribution matters more than any single verdict. If ten out of ten had confirmed, the honest reading would not be that the analysis was unusually perceptive — it would be that the hypotheses were written to confirm, and every critical finding in this record would deserve to be discounted accordingly. Four failures out of ten is what a set of genuine priors looks like.

Three of the four killed results went in Braze's favour, and in the first draft of this project all three reached the record and none reached the deck. That asymmetry was not a decision anyone made; it is what happens by default, because a favourable finding has nowhere obvious to go in a competitive deck and an unfavourable one does. It is now slide 40, on the grounds that an analysis reading as uniformly critical invites its audience to discount all of it — including the parts that are hardest to argue with.

#	HYPOTHESIS, AS WRITTEN BEFOREHAND	HOW IT DIED	DETAIL
H1	Growth is decelerating while sales spend holds	Sales and marketing fell faster than growth did, to its lowest share in the seven-year series (`data/financials_annual.csv`)	ch2 §2.3
H4	Some supported platforms are effectively unmaintained	Every SDK repository publishing releases had shipped within 13 days of capture, and the archived one has had its documentation removed too (`data/sdk_releases.csv`, `data/repos.csv`)	ch4
H9	Incident rate has risen with scale	Incidents peaked in 2023 and fell to the quietest full year on the status page in 2025 (`data/incidents.csv`)	ch4

DOCUMENTED

`source: docs/STRATEGY.md @ 2026-09-02`

The fourth kill, H6, is in Braze's favour too but is left off the slide because chapter 5 already states it plainly where it belongs: the channel roster is broad, not narrow, and the finding inverted into something more interesting than the hypothesis.

One discipline to note about how these are quoted. Each killed hypothesis keeps its caveat attached in the same breath — 2026's incident count runs at roughly double the 2025 monthly rate, and `braze-roku-sdk` sits idle with its documentation still live. A favourable finding stripped of its caveat is exactly as dishonest as an unfavourable one stripped of its context, and it is more tempting because nobody objects.

8.1 · The eight open questions

#	QUESTION	WHAT WOULD CLOSE IT
52	Why does the sub-processor disclosure name only Amazon and Google as hosting providers when the allowlist IPs Braze publishes for instance US-08 are all registered to Microsoft?	A revised sub-processor list, the DPA schedule a customer receives, or a direct answer from Braze
53	Is the `aze` region code Azure, and is US-08 a distinct class of instance?	A Braze statement, a customer on US-08, or a wider certificate-transparency capture
54	Does satisfaction fall as customer size rises?	Paid access to the G2, Gartner or TrustRadius rating-by-company-size breakdowns — all three paywalled
55	Why was the profile-lookup rate limit cut tenfold for customers onboarding after 22 August 2024, and were existing customers told?	A changelog entry, a support answer, or customers either side of the date
57	Can the customer roster be verified independently of the self-published stories?	Tag crawls, certificate transparency naming customer subdomains, or job ads naming the stack
58	What else is in the certificate-transparency estate?	A `CERTSPOTTER_TOKEN`; crt.sh returned HTTP 502 throughout this run
59	Does the gross-margin decline continue, and how much of it is AI cost rather than acquisition amortisation?	The FY2027 10-K, or quarterly cost-of-revenue detail
46	Actual list pricing and discounting policy	A public-sector procurement award, or a customer contract

There were nine until late in the work. **Question 56 — the exact split of open roles by function — closed**, from the source this table had already named: the Greenhouse board API. The answer is in chapter 7 §7.5 and the correction is logged in §8.2. It is worth one sentence here because it is the only gap in this project that closed after being written down, and it closed because the row said what would close it. **A gap recorded with its remedy is an instruction to a later reader; a gap recorded as a shrug is not.**

Two of the remaining eight deserve emphasis because they bound claims made elsewhere in this record.

Question 58 bounds chapter 4.5. The host list is partial — 833 hosts in `data/subdomains.csv`, captured through a rate-limited fallback after the primary source returned errors all day; the failures are logged in `logs/fetch-failures.md`. Everything said about what *was* provisioned stands. Nothing is said about what is absent, because the check was not exhaustive and an absence from an incomplete list is not a finding.

Question 54 is the one hypothesis that could not be tested at all. The analysis went in expecting enterprise satisfaction to be lower than small-business satisfaction, as it usually is. All three review sites paywall exactly that breakdown, and only seven of the 860 coded records in `data/review_coding.csv` carry a customer-size segment — far too few for anything. The nearest audited proxy points mildly the other way and measures something different: net retention among customers with $500k or more of annual recurring revenue is slightly *above* the all-customer figure (chapter 7). That is expansion, not satisfaction, and it is not offered as a substitute.

8.2 · Where a number moved – the corrections log

Required by `RECORD-SPEC.md` and mirroring `docs/FACTS.md`, with the old values still visible so a stale copy of this document can be recognised on sight. A record that shows its own errors is the reason to trust the rest of it.

#	FACT	WAS	IS NOW	WHY IT CHANGED	DATE
1	Status-page grouping	"132 in 12 groups"	132 rows in **17 named groups**, of which **15 are regional clusters**	The setup-time reconnaissance counted by eye. Counting `data/status_components.csv` gives 17 named groups and 15 clusters	2026-09-02
2	Capability page counts	Journey orchestration 319 · Identity resolution 136 · Segmentation 339	Canvas 249 · Identity 110 · Segmentation 242	**The pattern set moved, not the product.** The taxonomy in `docs/CAPABILITY-TAXONOMY.tsv` was revised from generic category words to Braze's own product names. Counts from the two runs are not comparable, and both are reproducible from the `pattern` column of `data/capabilities.csv`	2026-09-02
3	KakaoTalk's status	Read at first pass as documented but not marketed	**Marketed**, but missing from Braze's own docs channel index	A truncated search appeared to show no marketing page. Checking the full `/product/` enumeration found one	2026-09-02
4	The money chapter's completeness	Written with no reference to internal control over financial reporting	Carries the **material weakness** disclosure and both its halves (§2.1b)	The first pass read the 10-K for what it was looking for and missed Item 9A entirely. Found in self-review	2026-09-02
5	`DECK-SPEC.md`'s cluster count	"12 named regional clusters"	Corrected in the specification itself to 15	Correction 1 was recorded here but the specification that repeated the error was left uncorrected for a further pass. A spec known to be wrong will mislead the next reader	2026-09-02
6	The hiring split by function	"~284–300 open roles", split by department recorded as **uncapturable**	**296 open roles across 15 hiring departments**, with the exact split (ch7 §7.5)	The board's web filter would not drive under automation, so the gap was written down — correctly, at the time. The same board is published as unauthenticated JSON with the grouping already done. **When a page will not yield, look for the API behind it before recording a gap**	2026-09-02

Correction 2 is the one that matters most to a sceptical reader, and it is the reason the capability taxonomy is published rather than merely described. **A number that moves because the measurement moved is not a finding**, and presenting it as one would be the easiest way to mislead an audience in this entire project.

Correction 4 is the one that matters most to a careful one. The first pass of this analysis wrote an entire financial chapter without noticing that the company had disclosed its financial-reporting controls were not effective. Nothing else in this record is more important as a warning about how the omission happened: the filing was searched for answers to questions already being asked, rather than read for what it volunteered.

8.2b · Disagreements found in a dedicated sweep

Five conflicts existed after the first pass, all of which surfaced incidentally while writing other chapters. A dedicated sweep on 2026-09-02 — the discipline the reference project used to find twenty-one — added four more, and they cluster in exactly the place incidental discovery misses: **the definitions under the most-quoted numbers.**

CONFLICT	THE DISAGREEMENT	RULING IN ONE LINE
C-06	Each review panel's own surfaces disagree about its review count — Glassdoor 563 vs 524, Gartner 263 vs 267, TrustRadius 348 vs 162	Quote each site's headline rating with its own headline base, never mixed
C-07	The headcount in §1.3 counts full-time employees; the proxy's pay-ratio population on the same date is full-time **and** part-time, and is never sized	Say "full-time employees", never "headcount"
C-08	A "customer" is an **ultimate parent entity**, so the §7.1 metric counts corporate groups while the 178 stories count brands	The $283,000 bound is per corporate group, and the two rosters are not comparable
C-09	The 15 "Braze Group" entities include Braze KK, a majority-held VIE with outside investors — not a wholly-owned subsidiary	Name the exception when citing the entity list

Nine conflicts is still well short of the twenty-one the reference project recorded, and the honest reason is not that Braze generates fewer disagreements. It is that four of the nine here were found in a single deliberate hour, and no comparable sweep has been run over the documentation corpus, the incident history or the site inventory. That is a known gap in method rather than a finding about the company, and it is recorded as such.

The middle row is the important one, and it is the reason the taxonomy is published rather than described. A number that moves because the measurement moved is not a finding, and presenting it as one would be the easiest way to mislead an audience in this entire project.

8.3 · What was substituted, and what that cost

BLOCKED	SUBSTITUTED WITH	WHAT IS LOST
crt.sh (HTTP 502 all day, twice)	Cert Spotter, rate-limited without a token	A partial host list. Stated as partial everywhere it is used
Rating by company size on all three review panels	Nothing equivalent exists publicly	Hypothesis 7 is unresolved rather than answered
Per-department headcount on the careers board	The department taxonomy and a front-of-list sample	No function split. No percentage is offered
Independent customer detection	Not attempted	The 178 published stories are not treated as a roster

Every review panel that returned HTTP 403 to a script was captured through a browser session instead, and Glassdoor required the operator to sign in before its content became readable at all — recorded here because the capture route is part of the evidence.

What would change this chapter

Any of the nine questions closing. The most valuable single answer would be question 52, because it decides whether chapter 4.5 is an observation about infrastructure or something a buyer's procurement team needs to raise.

Slide map

The deck in order, generated from the built deck so coverage cannot be overstated. No evidence lives here — if a number appears in this table it belongs in a chapter.

#	SLIDE	RECORD
01	Competitor Analysis	§4
02	How we approach it	§8
03	How we grade	§8
04	Part I: The company	§1
05	The four documents	§8
06	Five things	§1
07	Who they say they are	§6
08	Origins	§1
09	How they got this big	§2
10	Acquisition: OfferFit	§3
11	Acquisition: North Star Y	§3
12	What it costs	§7
13	Who uses it	§7
14	Where they operate	§7
15	What customers say	§7
16	What employees say	§7
17	Who they compete with	§7
18	Part II: The product	§4
19	How one campaign works	§4
20	How data moves	§4
21	Stage 1: Data	§4
22	Stage 2: Identity	§4
23	Stage 3: Decisioning	§4
24	Stage 4: Building	§4
25	Stage 5: Content	§4
26	Stage 6: Delivery	§5
27	Stage 7: Interaction	§4
28	Channels	§5
29	Integrations	§5
30	Infrastructure	§4
31	Analytics	§4
32	The AI, honestly	§6
33	Part III: Strategy	§8
34	Where the money goes	§2
35	What comes next	§4
36	What protects them	§6
37	Part IV: Open questions	§8

#	SLIDE	RECORD
38	Deep dive: how real-time is it	§4
39	Reliability, measured	§4
40	Where the evidence went their way	§8
41	What we could not answer	§8
42	What to research next	§8
43	Close	§4