

Braze, Inc. — evidence record

Public sources only · every claim graded and traceable

How to read this. The chapters are the *Record*: every fact is stated **once**, in full, with its caveats in the same sentence as the number. The [Slide Map](#) is the deck in order and holds no evidence — it points back into the chapters.

Grades run **AUDITED** **INFRASTRUCTURE** **DOCUMENTED** **THIRD-PARTY** **CLAIMED**, ordered by how hard it would be for the vendor to have made the statement untrue. **A claim takes the grade of its weakest source.**

1 • The company

Braze is a fifteen-year-old New York software company that has been public since November 2021, and as of January 2026 it is no longer controlled by its founders through super-voting stock. Almost everything in this record is knowable because of that listing.

1.1 • The entity

FACT	VALUE	GRADE	SOURCE
SEC registrant	Braze, Inc.	AUDITED	data/filings.csv
CIK / ticker	0001676238 / BRZE (Nasdaq)	AUDITED	data/filings.csv
SIC classification	7372 Services-Prepackaged Software	AUDITED	data/filings.csv
Fiscal year end	31 January	AUDITED	data/filings.csv
Founded and incorporated	2011, in Delaware	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:249,598
Headquarters	28 East 28th Street, Floor 12, New York, NY 10016	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:598

Fiscal years end on 31 January, so "FY2026" means the year to 31 January 2026. This matters more than it looks: XBRL's *fy* field labels the *filing*, not the period, and reading it directly mislabels every year in the series. `tools/sec_facts.py` derives a correct period column and every figure in chapter 2 keys on that.

source: data/filings.csv @ 2026-09-01

1.2 • Going public, and the end of founder control

FACT	VALUE	GRADE	SOURCE
S-1 filed	22 October 2021	AUDITED	data/filings.csv
Final prospectus (424B4)	18 November 2021	AUDITED	data/filings.csv
Dual-class structure	Ended 30 January 2026	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:1197

The governance change is recent enough to be easy to miss and material enough to matter. On 30 January 2026 the Class B common stock was retired and automatically converted into Class A, and the 10-K states the consequence plainly: "our executive officers and early investors no longer hold super-voting rights. Consequently, our voting power is now more widely distributed among our public stockholders."

For a competitor, the practical reading is that a company whose direction was until very recently insulated from its own shareholders is now not. Nothing in the evidence says what Braze will do with that; it is a change in who can apply pressure, and it is nine months old.

source: sources/filings/2026-03-25_10-K_000013.txt:1197 @ 2026-03-25

1.3 · People and footprint

FACT	VALUE	GRADE	SOURCE
Full-time employees	1,988 as at 31 January 2026	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:590
Legal entities in the group	15, across 14 territories	DOCUMENTED	sources/clean/braze-subprocessors.md:41
Total SEC filings	737, 2017-07-20 → 2026-08-28	AUDITED	data/filings.csv

The employee figure is the 10-K's own, as-of a single date, and counts differently from a LinkedIn headcount; the two are never merged here, and no LinkedIn figure was captured in this run.

The entity list comes from the sub-processor disclosure rather than from marketing, and it names its own territories: Australia, Brazil, Canada, France, Germany, Ireland, Spain, United States, Ireland & Romania, Japan, South Korea, United Kingdom, United Arab Emirates, Singapore and Indonesia. One line in it is worth noticing — "Braze Ireland Procurement Limited" is located in "Ireland & Romania", and the careers board lists Bucharest as a hiring location. Two unrelated documents describing the same Romanian presence is the kind of corroboration this method is built to find.

source: sources/clean/braze-subprocessors.md:41 @ 2026-09-02

What would change this chapter

A merger, a take-private, or a change of control — all of which would appear as an 8-K within four business days. A materially different headcount would appear in the next 10-K. The entity list changes when the sub-processor document is revised; this record uses revision 1 June 2026 and any later revision supersedes it.

2 • The money

Growth is decelerating and efficiency is improving at the same time: revenue growth has halved since FY2023 while sales and marketing has fallen from 56.7% of revenue to 44.3%, operating cash flow has been positive for three years, and the net loss persists mainly because share-based compensation now exceeds it.

Every figure in this chapter is audited, filed under legal penalty, and re-derivable by running `python3 tools/sec_facts.py`. Fiscal years end 31 January.

2.1 • Seven years of revenue and margin

FISCAL YEAR	REVENUE (\$M)	GROWTH	GROSS MARGIN
FY2020	96.4	—	63.0%
FY2021	150.2	55.9%	63.7%
FY2022	238.0	58.5%	67.0%
FY2023	355.4	49.3%	67.4%
FY2024	471.8	32.7%	68.7%
FY2025	593.4	25.8%	69.1%
FY2026	738.2	24.4%	67.1%

AUDITED

source: `data/financials_annual.csv` @ 2026-09-02

Two things happen in the final row and they should be read together. Growth continues to decelerate — it has roughly halved over three years — and gross margin falls for the first time in six years, by 2.0 points.

The company explains the margin fall itself, and the explanation is specific: "primarily due to acquisition related operating costs, including personnel costs of acquired workforce and amortization expense of acquired technology, in addition to increased costs related to our tech stack." The acquisition in question is OfferFit (chapter 3), whose developed technology is amortised **to cost of revenue**. So the first margin decline in six years is, on the company's own account, largely the price of buying its AI decisioning engine. How much of the remainder is the AI cost base rather than the acquisition is not split out, and is open question 59.

source: `sources/filings/2026-03-25_10-K_000013.txt:1565` @ 2026-03-25

2.1b • The caveat that travels with every figure in this chapter

Braze disclosed a **material weakness in its internal control over financial reporting** as at 31 January 2026. The CEO and CFO concluded that disclosure controls were "not effective at the reasonable assurance level". The weakness is specific: "ineffective information technology general controls, or ITGCs, in the areas of **user access and program change management** over the information technology systems that support our financial reporting processes" — and the company states that automated and IT-dependent manual controls relying on that environment were consequently ineffective too.

AUDITED

source: sources/filings/2026-03-25_10-K_000013.txt:3258,3286 @ 2026-03-25

Both halves belong in the same breath, because either one alone misleads.

The reassuring half is substantial and is the company's own and the auditor's: the material weakness "did not result in any identified misstatements to the financial statements, and there were no changes to previously issued financial results"; management concluded the statements "fairly present in all material respects"; and Ernst & Young issued an attestation report on internal control over financial reporting. Remediation is under way with "significant progress" claimed and no completion date given. OfferFit was excluded from the assessment, as SEC guidance permits in an acquisition's first year — it is 1.4% of total assets and 1.8% of revenue.

The unreassuring half is that the control environment producing every number in this chapter was judged not effective by the people who signed it, and the specific area is user access and change management.

Two disciplines apply. **Do not extrapolate this into a claim about customer data security** — it is scoped to financial-reporting systems and says nothing about the platform. And note that it is consistent with, not contradicted by, conflict **C-05**: the restatement check found nothing, and the 10-K independently confirms no previously issued results changed.

2.2 • Losses, cash and the equity cost

FISCAL YEAR	NET LOSS (\$M)	OPERATING CASH FLOW (\$M)	SHARE-BASED COMP (\$M)
FY2022	-76.7	-35.4	47.2
FY2023	-139.0	-22.3	72.2
FY2024	-129.2	+6.8	97.2
FY2025	-103.7	+36.7	115.1
FY2026	-131.3	+71.4	143.7

AUDITED

source: data/financials_annual.csv @ 2026-09-02

The favourable reading and the uncomfortable one are both in this table and both are true.

Favourably: operating cash flow crossed into positive territory in FY2024 and has grown every year since, roughly doubling each time. A company generating \$71.4m of operating cash is not a company in difficulty, whatever the loss line says.

Uncomfortably: share-based compensation has exceeded the entire net loss for two consecutive years — 111% of it in FY2025 and 109% in FY2026. Stated as an observation rather than an accusation: the loss is substantially a non-cash equity cost, and the people bearing it are shareholders through dilution. Weighted-average diluted shares went from 94.6m in FY2023 to 107.9m in FY2026, a rise of 14.1% in three years.

2.3 · Where the money goes

FISCAL YEAR	S&M SHARE OF REVENUE	R&D SHARE OF REVENUE
FY2020	59.5%	21.1%
FY2022	53.4%	24.8%
FY2023	56.7%	27.4%
FY2024	52.4%	25.4%
FY2025	47.6%	22.6%
FY2026	44.3%	22.6%

AUDITED

source: data/financials_annual.csv @ 2026-09-02

This table killed a hypothesis. The analysis went in expecting growth to be decelerating *while sales spend held* — the familiar pattern of a company buying its last few points of growth. The evidence says the opposite: S&M has fallen more than twelve points as a share of revenue since FY2023, its lowest in the seven-year series, while growth decelerated. Braze is decelerating and getting more efficient at the same time.

Sales and marketing still outspends research and development by 1.96 times — \$327.0m against \$167.1m in FY2026 — which is ordinary for enterprise software and worth stating plainly rather than treating as a finding.

2.4 · Contracted revenue and geography

FACT	VALUE	GRADE	SOURCE
Remaining performance obligation, FY2026	\$1,033.0m	AUDITED	data/financials_annual.csv
RPO relative to FY2026 revenue	1.40×	AUDITED	derived from data/financials_annual.csv
United States revenue, FY2026	\$405.1m	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:2421
International revenue, FY2026	\$333.1m — 45.1% of the total	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:2421

Remaining performance obligation crossed a billion dollars in FY2026, having risen every year from \$234.2m in FY2021. At 1.40 times current-year revenue it is a genuine forward visibility measure and, unlike a pipeline number, it is contracted and audited.

The geographic split is an audited disclosure rather than an inference, and it carries a sentence that matters more than the percentages: "Other than the United States, no other individual country accounted for 10% or more of total revenue for any of the periods presented." Forty-five per cent of revenue is international and none of it concentrates. Set that against fifteen regional clusters (chapter 4) and fifteen legal entities (chapter 1), and the shape is a wide, thin international footprint carrying a heavy infrastructure and legal base.

source: sources/filings/2026-03-25_10-K_000013.txt:2421 @ 2026-03-25

What would change this chapter

A restatement — none has occurred; across the twenty-nine key XBRL concepts tracked, data/financials_restated.csv is empty, and that negative result is recorded as conflict C-05 rather than left unstated. The next 10-K would change the trend lines. A second consecutive year of margin decline would move the FY2026 fall from an acquisition effect to a pattern.

3 · What they bought

Braze has made two material acquisitions. One bought a market; the other bought the product it now markets as its own artificial intelligence. Both prices are disclosed, which is a thing a private vendor almost never permits.

3.1 · OfferFit – bought June 2025, renamed AI Decisioning Studio

FACT	VALUE	GRADE	SOURCE
Date	2 June 2025	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3130
Total consideration	\$303.2m	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:1788
Structure	\$195.3m cash, \$107.6m in Braze Class A common stock	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3130
Goodwill	\$233.4m	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3134
Identified intangibles	\$66.6m	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3134
Developed technology	\$56.7m, amortised over six years to cost of revenue	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:3137, 3139

The single most useful sentence in the whole filing set, for anyone assessing Braze's AI position, is in the auditor's critical-audit-matter note: **"OfferFit, Inc. ('OfferFit') which is now known as AI Decisioning Studio."** The product Braze markets as BrazeAI Decisioning Studio is a company it bought fifteen months ago, renamed.

The 10-K describes what was acquired as "OfferFit's multi-agent decisioning engine", which is also the answer to what the agentic layer is: acquired, not built.

Three further details are worth stating precisely rather than in summary.

- Goodwill is 77.0% of the purchase price. That is high, and it is what a price looks

like when most of what is being bought is a team and a position rather than assets.

- The developed technology is amortised **to cost of revenue**, not to operating

expense, which is why this acquisition shows up in the gross-margin line in chapter 2 rather than below it.

- Ernst & Young flagged the \$56.7m developed-technology valuation as a critical audit

matter — "complex due to the significant estimation required by management". That is a routine disclosure about valuation difficulty, not a finding about the business, and it should be described that way.

source: sources/filings/2026-03-25_10-K_000013.txt:1788-1792,3130-3139 @ 2026-03-25

3.2 · North Star Y – bought June 2023, and the earn-out that paid nothing

FACT	VALUE	GRADE	SOURCE
Date	1 June 2023	AUDITED	sources/filings/2025-03-31_10-K_000054.txt:2972
What it was	Braze's exclusive reseller in Australia and New Zealand	AUDITED	sources/filings/2025-03-31_10-K_000054.txt:2972
Total consideration	\$26.8m — \$20.6m cash, \$6.1m stock, \$1.8m contingent	AUDITED	sources/filings/2025-03-31_10-K_000054.txt:2974
Earn-out available	Up to \$10.0m in year one and \$16.0m in year two	AUDITED	sources/filings/2025-03-31_10-K_000054.txt:2974
Earn-out paid	Nil	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:2569

Braze bought its own Australia and New Zealand reseller to take the market direct. The sellers could have earned up to a further \$26.0m on "qualified revenue performance metrics" over the two years following completion.

They earned none of it. The FY2026 10-K records that Braze "reduced the contingent consideration liability related to the acquisition of North Star Y, Pty Ltd to zero as it was determined that the sellers did not satisfy the earn-out qualifications."

This is stated as an observation and nothing more. An earn-out that does not vest means the revenue targets set at completion were not met; it does not by itself mean the acquisition failed, and the filings do not say it did. What it does mean is that the AU/NZ business did not perform to the plan written into the purchase agreement in June 2023, and that is a fact about a region Braze had considered important enough to buy.

source: [sources/filings/2026-03-25_10-K_000013.txt:2569](#) @ 2026-03-25

What would change this chapter

A third acquisition, which would appear as an 8-K and then in the following business-combination note. A goodwill impairment against either purchase would be disclosed and would change how the OfferFit price should be read. Measurement-period adjustments to the OfferFit allocation are still possible; the figures here are the FY2026 10-K's as-adjusted values.

4 · The platform

Braze's documentation is markedly more honest than its marketing about what the platform will and will not do, and the limits it admits are commercial as much as technical: three of four ingestion paths are labelled "not real-time" by Braze itself, the endpoint that reads user profiles back out was cut tenfold for customers who joined after August 2024, and the billing unit is the individual attribute write.

4.1 · The corpus this chapter rests on

FACT	VALUE	GRADE	SOURCE
Documentation pages captured	1,352	DOCUMENTED	data/docs_index.csv
Words in the corpus	1,565,479	DOCUMENTED	data/docs_sections.csv
Documented REST endpoints	135, across 28 top-level namespaces	DOCUMENTED	data/api_endpoints.csv
Public repositories	137	DOCUMENTED	data/repos.csv

The corpus is shaped unevenly, and the shape is itself informative: 696 user-guide pages, 322 partner pages, 200 API pages, 106 developer-guide pages. The developer guide averages about 2,700 words a page; the partner section averages about 810. The integration network is broad and documented shallowly — a lot of short pages, one per partner.

source: data/docs_sections.csv @ 2026-09-02

4.2 · How fresh the data is — their table, not ours

Braze publishes a comparison of its four ingestion paths and grades each one's latency itself. This is the single most load-bearing table in the corpus.

PATH	BRAZE'S OWN LATENCY DESCRIPTION
Standard Cloud Data Ingestion sync	"Not real-time; minimum 15-minute sync cadence"
CDI Segments (Connected Sources)	"Not real-time; refreshes on your Segment Extension schedule"
CDI Canvas triggers	"Not real-time; bounded by sync schedule (minimum 15 minutes)"
/users/track and the SDKs	"Near-real-time (async processing)"

DOCUMENTED

source: sources/docs/docs__user_guide__example_library__data__compare_data_ingestion_options.md:86-90 @ 2026-09-02

Three of the four are labelled not real-time by the vendor. The fourth is qualified as near-real-time with asynchronous processing. Warehouse syncs "can run from every 15 minutes to once per month", and going faster than fifteen minutes is not self-serve: "contact your customer success manager or use REST API ingestion."

This is the substance of conflict **C-01**. The 10-K describes "high-volume, continuous streaming of user data to be processed in real time" and the developer guide says "the platform can handle any data in real time" — both true of

the SDK and API path. The correct sentence is not that Braze overstates, but that the word covers two architectures: if your customer data lives in a warehouse, fifteen minutes is the floor unless you move to the API.

Other admitted ceilings on the same page: 75 objects per `/users/track` request combined across attributes, events and purchases; a 60-minute query runtime cap per connected source; approximately 3.75 million Canvas entries per hour per sync run. Currents, the streaming export, is "an optional Braze add-on" that exports "every five minutes, or every 15,000 events, whichever comes first."

4.3 · Data in versus data out

DIRECTION	ENDPOINT	LIMIT
In	<code>/users/track</code>	3,000 requests per 3 seconds for customers with data points in their pricing — 60,000 a minute — at 75 objects each
Out, by identifier	<code>/users/export/ids</code>	250 requests per minute if onboarded on or after 22 August 2024; 2,500 if before
Out, in bulk	<code>/users/export/segment</code>	The default 250,000 requests per hour

DOCUMENTED

source: `sources/docs/docs__api__api_limits.md:29-39` @ 2026-09-02

Two things here, and the second is the interesting one.

The asymmetry is real but must be stated fairly. Writing can reach roughly 4.5 million objects a minute; reading profiles back by identifier is capped at about 12,500 profiles a minute for a customer who joined after August 2024, at 50 identifiers per request. Those are different operations — writing event objects against reading whole profiles — so the ratio shows where the design and commercial priority sits, not a like-for-like throughput comparison. And bulk export exists and is generously limited, so this is not a data-lock-in story and should not be told as one.

The dated boundary is the finding. On 22 August 2024 the profile-lookup limit was cut by a factor of ten for new customers, and existing customers kept the old limit. Two customers on the same product, doing the same thing, are subject to limits an order of magnitude apart depending on when they signed. Braze documents this openly in a table. Why it changed, and whether anyone was told, is open question 55.

4.4 · What it is built on

FACT	VALUE	GRADE	SOURCE
Named backing stores	Snowflake, Kafka, MongoDB, Redis	DOCUMENTED	sources/docs/docs__developer_guide__getting_started__architec
MongoDB-backed	Custom events, custom attributes, user profiles, purchases, most segmentation	DOCUMENTED	...architecture_overview.md:46-56
Snowflake-backed	SQL Segment Extensions, Prediction Suite, AI Personalized Item Recommendations, Estimated Real Open Rate	DOCUMENTED	...architecture_overview.md:58-68
Regional clusters	15: US 01-08, US 10, EU 01-02, AU 01, ID 01, JP 01, KR 01	DOCUMENTED	data/status_components.csv
Subsystems per cluster	7, identical in every one	DOCUMENTED	data/status_components.csv

The status page is an architecture disclosure the vendor made by accident. Every cluster exposes the same seven subsystems — Dashboard, SDK Data Collection, Data Processing, REST APIs, Outbound Messaging, Currents, Cloud Data-Ingestion — which is a functional decomposition of the product, published live and updated during outages. Note that Currents and CDI are first-class per-cluster subsystems rather than features.

The cluster geography is corroborated independently: the sub-processor disclosure lists AWS regions for the United States, European Union, Australia, Indonesia (with backup in Singapore), Japan and South Korea — the same six geographies the cluster names use. Channel delivery is modelled differently, as a single global group rather than per cluster. There is no US 09; the US clusters run 01 to 08 and then 10.

The customer-visible consequence of the architecture is not the clusters. It is the storage split, and Braze flags it themselves under an "important" callout: **"Removing data from one system does not automatically remove it from the other."** Erroneous custom-event data must be addressed in MongoDB, separately from anything Snowflake-backed. For a buyer with deletion obligations, that is a two-system problem stated in the vendor's own documentation.

source: sources/docs/docs__developer_guide__getting_started__architecture_overview.md:70 @ 2026-09-02

4.5 · One instance is not on the same cloud as the others

This is the clearest result the method produced, and it comes from three sources that have nothing to do with each other.

- The sub-processor disclosure of 1 June 2026 names exactly two hosting providers:

Amazon Web Services, and Google LLC for Google Cloud Platform. Microsoft is not named anywhere in the document. DOCUMENTED

- A **second, independent company document** says the same thing. The 10-K risk factors

state: "We rely upon third-party providers of cloud-based infrastructure, **including Amazon Web Services and Rackspace**, to host our products." Microsoft is not named there either. **AUDITED**

- Braze's own documentation publishes, per instance, the IP addresses a customer must

allowlist so Braze can reach their warehouse. Every address listed for **US-08** is registered to **Microsoft Corporation**. Every address listed for every other instance — US-10, AU-01, ID-01, JP-01, KR-01 and the generic US block — is registered to an Amazon entity. The check covered every instance for which Braze publishes a list.

INFRASTRUCTURE

- Certificate transparency shows 50 hosts on region codes p-aze-us, s-aze-us and

d-aze-us — a code matching no AWS region identifier, where every other code in the estate does. Hosts on it include sdk-us08, subcenter-08 and itp-api-08, tying the code to instance 08. **INFRASTRUCTURE**

source: sources/external/rdap-instance-ip-ownership_2026-09-02.json @ 2026-09-02

source: sources/docs/docs__user_guide__data__unification__cloud_ingestion__connected_sources.md:239 @ 2026-09-02

State the three observations and stop. A hosting arrangement may sit outside a sub-processor listing for reasons not visible from outside, the document may be behind the estate, or there may be an explanation nobody external can see. Asserting a disclosure failure would be a legal conclusion this evidence does not support. The question — why the disclosure and the infrastructure differ — is recorded as open question 52 and as conflict **C-03**.

Note also that the certificate-transparency host list is **partial**: crt.sh returned HTTP 502 throughout and the Cert Spotter fallback was rate-limited without a token. No claim is made here that anything is *absent* from CT, because the list was not exhaustive.

4.6 · Identity, and where it is quietly lossy

FACT	VALUE	GRADE	SOURCE
Aliases per profile	No limit	DOCUMENTED	sources/docs/docs__user_guide__data__unification__user_data
Aliases per label	One, and unique across the user base	DOCUMENTED	sources/docs/docs__api__endpoints__user_data__post_user_id
Identifiers accepted on ingest	Five: external_id, user alias, braze_id, email, phone	DOCUMENTED	...compare_data_ingestion_options.md:113-118
Identifiers accepted for warehouse segmentation	One: external_user_id only	DOCUMENTED	...compare_data_ingestion_options.md:113-118

The identity model is generous at the top and narrow at the bottom: unlimited aliases on a profile, but a warehouse-driven segment can only be built on a single identifier type.

Two admitted behaviours deserve a buyer's attention because neither is visible from marketing and both fail quietly.

A merge can decline and still report success. If both profiles carry invalid phone numbers Braze does not merge them, and their documentation says so directly: "The endpoint still returns 202 Accepted with a success message, so the HTTP response does not indicate that the merge was skipped."

Reporting splits across surfaces after a merge. Dashboard campaign summaries attribute a pre-merge send to the surviving profile; Currents, Query Builder and Messaging History still attribute it to the orphaned profile's ID. Both are correct by their own rules and they do not agree, which is a reconciliation problem for anyone joining Braze data to a warehouse. Orphaned profiles are also not eligible to receive messages.

source: sources/docs/docs__api__endpoints__user_data__post_users_merge.md:178 @ 2026-09-02

4.7 · What a customer is billed for

The billable unit is defined in the documentation with contract-grade precision: a "data point" is "a billable unit of use of the Braze Services, measured by a session start, session end, custom event, or purchase recorded, as well as any attribute set on an end user profile", and each such item "shall each count as a single data point". A session start and a session end are two.

Engagement is explicitly *not* billed — push tokens, device information, and all campaign engagement tracking such as email opens and push clicks are excluded, as are user deletion, Connected Content, subscription-state changes and external-ID renames. That is a genuinely customer-friendly boundary and belongs in any fair account of the pricing.

The tension is in the advice Braze gives alongside it: "**Don't waste data points. Only update changing data!**", and "we recommend setting up a program to prevent sending the same unchanging data." A platform positioned on continuous streaming of behavioural data prices in a way that makes its customers build programmes to send less of it. Both facts are theirs; the tension between them is the finding, and it is corroborated from an unrelated direction by a Gartner reviewer citing "steep pricing and strict data-point limits, where overages can become expensive very quickly" (chapter 7).

source: sources/docs/docs__user_guide__data__infrastructure__data_points.md:18,20,36-38 @ 2026-09-02

4.8 · A decade of reliability, and an engineering record

FACT	VALUE	GRADE	SOURCE
Incidents published	451, 2016-10-09 → 2026-08-05	DOCUMENTED	data/incidents.csv
Median duration	79 minutes	DOCUMENTED	data/incidents.csv
p90 duration	311 minutes	DOCUMENTED	data/incidents.csv
Major or critical	98 of 331 non-maintenance incidents — 29.6%	DOCUMENTED	data/incidents.csv
Public issues captured	845, all from non-members of the org	DOCUMENTED	data/issues.csv
Issues closed	822 — 97.3% — median 11 days	DOCUMENTED	data/issues.csv

Incidents peaked in 2023 at 60 and fell to 27 in 2025, the quietest full year on the status page, over a period when revenue grew 7.7 times. Incident rate per unit of scale has fallen sharply, which kills the hypothesis that it would rise. The caveat travels with it: 2026 stands at 35 through August, roughly double the 2025 monthly rate.

The component distribution has a shape worth naming. The Dashboard appears in 63 incidents, more than any other component and more than twice Outbound Messaging's 27. The control plane that marketers use fails more often than the sending path.

None of this may be compared against a vendor that publishes no status page. A decade of visible incidents is a disclosure practice, not a defect count, and a competitor with silence has not thereby earned a better record.

On the engineering side the record is straightforwardly good. All nine SDK repositories that publish releases shipped within thirteen days of capture. The one platform whose repository is archived — Unreal — has had its documentation removed entirely, so the marketing and the maintenance agree. The single soft spot is braze-roku-

sdk: 181 days since its last push while 38 Roku documentation pages remain live. And 845 unsolicited public issues closed at 97.3% with a median of 11 days is a real operational number that Braze does not control the denominator of.

The issue corpus describes the SDK surface, not the dashboard, and its authors are developers rather than buyers. It is not a satisfaction measure and is not used as one.

What would change this chapter

A revised sub-processor list naming Microsoft would resolve §4.5 in one direction; a statement that US-08 is something other than an Azure deployment would resolve it in another. A change to the documented rate limits, or the removal of the two-tier export limit, would date §4.3. Two consecutive quiet quarters would make the 2026 incident uptick noise; two more like April 2026 would make it a trend.

5 · Channels and delivery

Thirteen channels are documented, which is broad rather than narrow — but the documentation index and the marketing site have drifted apart in both directions at once, and the sub-processor disclosure shows that only two of the thirteen have a named delivery intermediary at all.

5.1 · The roster, counted from both ends

Braze's own documentation channel index lists twelve:

In-product — in-app messages, Content Cards, Banners. **Out-of-product** — Email, Transactional email, Landing pages, LINE, Live notifications, Push, SMS/MMS/RCS, Webhooks, WhatsApp.

DOCUMENTED

source: sources/docs/docs__user_guide__channels.md:14-64 @ 2026-09-02

A thirteenth exists and is not on that list. **KakaoTalk** has four documentation pages — setup, message creation, click tracking and reporting, a fully documented channel — and a marketing product page at /product/kakaotalk-messenger. It is mentioned zero times on Braze's own channels index.

The drift runs the other way too. Five documented channels have no dedicated marketing product page anywhere on the site: **Banners, Transactional email, Landing pages, Live notifications and Webhooks**. The largest of them is not small — Landing pages carries eleven user-guide pages plus three partner pages, a substantial documented capability that is sold under no name.

So the honest answer to "how many channels" is three different numbers depending on which document you trust, which is why this is conflict **C-02**. The useful sentence is not the count: *five documented channels have no marketing page, and one marketed channel is missing from their own index*.

Investment also differs sharply by channel in a way the roster hides: /product/line exists in four languages, /product/kakaotalk-messenger in two.

source: data/site_inventory.csv @ 2026-09-01

5.2 · Who actually delivers each channel

The sub-processor disclosure is legally compelled to be complete, which is what gives it force. Revision 1 June 2026 names 17 third-party sub-processors.

CHANNEL	NAMED DELIVERY SUB-PROCESSOR
Email	Three — Amazon SES, Bird.com (SparkPost), Twilio (SendGrid)
SMS / mobile messages	Two — Infobip, Twilio
Everything else	None named

DOCUMENTED

source: sources/clean/braze-subprocessors.md:21-37 @ 2026-09-02

Email has three delivery providers named, plus a fourth supplier, Mailgun, for Email on Acid previewing. Mobile messages have two. **Push, in-app messages, Content Cards, Banners, Webhooks, WhatsApp, LINE, KakaoTalk,**

Landing pages and Live notifications have no delivery sub-processor named at all.

The caveat travels in the same breath as the finding: absence from a sub-processor list is not proof of no intermediary. Platform transports such as APNs and FCM, and the WhatsApp, LINE and Kakao business APIs, may not be classified as sub-processors processing personal data on Braze's behalf. What can be said is what the compelled disclosure names — and it names middlemen for email and SMS, and for nothing else.

Read for what it says about resilience rather than about channels, the email arrangement is the notable one: three interchangeable delivery providers for the highest-volume channel is a deliberate redundancy, and a reasonable competitor should assume email deliverability is not a single point of failure here.

5.3 · The rest of the supplier list, read for what it discloses

Beyond delivery, the same document discloses things no marketing page would:

- **Two hosting providers named** — Amazon Web Services, and Google LLC for Google Cloud

Platform. What that omits is dealt with in chapter 4.

- **End-user profiles are stored by a third party.** Rackspace US, Inc. provides

"Database Administration as a Service (DBaaS), a managed database service provider that **hosts and stores End User profiles.**"

- **Monitoring receives user identifiers.** "Braze may provide End User metadata, such as user identifiers, to DataDog for support and application troubleshooting."

- **Three foundation-model suppliers** — Anthropic, OpenAI and Google — dealt with in chapter 6.

- Databricks, dbt Labs, Domino Data Lab, ClickHouse, Snowflake, Cloudflare and Fastly make up the remainder: analytics, transformation, CDN and traffic management.

Braze separately claims "over 150 technology partners, which we call 'Alloys'". That is a marketing figure appearing inside technical documentation and is graded accordingly. **CLAIMED**

source: sources/docs/docs__developer_guide__getting_started__architecture_overview.md:82 @ 2026-09-02

What would change this chapter

A revision of the sub-processor document — this record uses 1 June 2026 and any later revision supersedes it entirely. A new channel appearing in the docs index, or KakaoTalk being added to it. A marketing page appearing for any of the five unmarketed channels, which would close the drift rather than resolve it.

6 • The AI

Four independent lenses point the same way: the decisioning engine was bought rather than built, the models come from three named external suppliers, the AI surface carries no API of its own, and its documentation footprint is a fraction of the platform's established capabilities. None of that makes it thin – it makes it *new*, *supplied*, and measurable, which is a different and more useful claim.

The rule this chapter follows, carried from the reference project: **never assert thinness**. State the shipped machine learning and the agentic layer as separate, counted things, and let the ratio speak. A claim like this is only defensible when several unrelated lenses agree, so four were used.

6.1 • Lens one – how much documentation exists

Focused pages, meaning pages where the capability's vocabulary appears at least five times in the body – a threshold that separates "mentions it" from "is about it". Counted against the documentation corpus described in chapter 4.

CAPABILITY	FOCUSED PAGES
Recommendations	73
BrazeAI (umbrella brand)	32
Generative AI	26
Decisioning Studio	22
Agents	17
Intelligent Timing / Channel	13
Predictive Suite	7
— for comparison —	
Email	347
Canvas (journeys)	249
Segmentation	242
Liquid templating	123

DOCUMENTED
source: data/capabilities.csv @ 2026-09-02

The pattern set behind these counts is a judgement and is published as such, in docs/CAPABILITY-TAXONOMY.tsv and in the pattern column of the CSV, precisely so that a reader who disagrees can change it and rerun. It was revised on 2026-09-02 to use Braze's own product names rather than generic category words; counts from before that revision are not comparable with these, and the change is logged in the corrections table.

The honest caveat: a new feature is under-documented by construction. Predictive Suite at seven focused pages could mean an immature product or a simple one. The counts do not distinguish those, which is why three further lenses follow.

6.2 · Lens two – the published API surface

Of Braze's 135 documented REST endpoints, spread across 28 top-level namespaces, the number sitting in an AI, ML, prediction, recommendation, agent or decisioning namespace is **zero**. The check was exhaustive over every documented path.

The namespaces are: app_group, apps, campaigns, canvas, catalogs, cdi, content_blocks, custom_attributes, data_objects, email, events, export, fcm, kpi, media_library, messages, preference_center, purchases, scim, segments, sends, sessions, sms, subscription, templates, transactional, users, v2.

For contrast, Canvas has 10 endpoints, Catalogs 9, subscription management 9, identity 8, email 16.

DOCUMENTED

source: data/api_endpoints.csv @ 2026-09-02

A capability with no API surface is not necessarily immature — it may be intended for use only inside the dashboard. But for a competitor, or for a buyer planning to drive Braze programmatically, it means the AI features cannot currently be orchestrated from outside the product.

6.3 · Lens three – who supplies the models

The sub-processor disclosure names three external providers, each with the same description: "provides artificial intelligence models and machine learning infrastructure to support data processing and advanced platform features."

- **Anthropic PBC**
- **OpenAI OpCo, LLC**
- **Google LLC**

DOCUMENTED

source: sources/clean/braze-subprocessors.md:22,31,34 @ 2026-09-02

Using foundation models from suppliers is ordinary and not a criticism — very few companies train their own. What is unusual is that Braze is legally obliged to name them and therefore does, which most vendors' marketing never will. Any competitor claiming Braze's AI is proprietary in the model layer is contradicted by Braze's own compelled disclosure.

6.4 · Lens four – the engine itself was bought

The decisioning product is OfferFit, acquired 2 June 2025 for \$303.2m and renamed AI Decisioning Studio, as chapter 3 sets out from the auditor's own note. The agentic layer Braze markets is, in the 10-K's words, "OfferFit's multi-agent decisioning engine".

That is the fourth lens and the one that dates the rest: whatever the documentation footprint says, the engine is fifteen months old inside Braze.

6.5 · Lens five – the words customers use

Not planned as a lens, but the review corpus turned out to say something specific enough to include. Across the captured panels, AI is named thirteen times.

What the positives name: "AI Advisor... a strong addition, providing robust capabilities and useful insights"; "SQL Query Builder with AI assistant"; "the AI personalization tool".

What the criticisms name: two G2 reviewers independently single out the **AI copywriting assistant** as needing "a careful human hand" before anything reaches a customer; a TrustRadius reviewer says "AI features feel half baked"; and TrustRadius's own synthesis records that "the platform's cost, particularly for advanced AI features, is a recurring consideration."

THIRD-PARTY

source: `sources/panels/g2.md:137,189` @ 2026-09-02

Reviewers are describing generative assistance, not decisioning. Nobody in the captured sample mentions Decisioning Studio or Agents at all — consistent with a product acquired fifteen months ago, and a reminder that this sample is small and used for vocabulary rather than for any percentage.

6.6 · What the five lenses together support

They support this and nothing stronger: **Braze's AI is real, shipping, recently and largely acquired, running on models from three named external suppliers, without an API surface of its own, and documented at a fraction of the depth of the platform's established capabilities.** Every clause there is sourced separately above.

They do not support "their AI is thin". That sentence is unfalsifiable, and the evidence does not need it.

What would change this chapter

An AI namespace appearing in the REST API. A jump in the Decisioning Studio or Agents page counts on a re-run of `capability_count.py` — the measurement is repeatable, which is the point. A sub-processor revision adding or removing a model provider. And review vocabulary beginning to name decisioning rather than copywriting, which would show the acquired engine reaching customers.

7 · The market

Braze sells to 2,609 customers, expansion within them is slowing, and the vendors buyers actually weigh it against are twice as many as the four it names – with the five extra all being specialists it does not mention.

7.1 · Three customer counts that must never be merged

ROSTER	COUNT	GRADE	SOURCE
The 10-K's defined metric	2,609 as at 31 January 2026	AUDITED	sources/filings/2026-03-25_10-K_000013.txt:1379
Self-published customer stories	178 customers/ URLs	CLAIMED	data/site_inventory.csv
Independent detection	Not attempted in this run	—	open question 57

The first is a defined, audited metric. The second is a marketing selection. The third was not done, and is recorded as a gap rather than as a zero — nobody should read "178" as the customer base or as a sample of it.

7.2 · Expansion is slowing, and the enterprise premium has almost gone

MEASURE	FY2024	FY2025	FY2026
Dollar-based net retention, all customers	117%	111%	109%
Dollar-based net retention, customers ≥\$500k ARR	120%	114%	110%
Customers with ARR ≥\$500k	202	247	333

AUDITED

source: sources/filings/2026-03-25_10-K_000013.txt:1400 @ 2026-03-25

Two movements, and the second is easy to miss.

Net retention is falling, and Braze explains why itself: "primarily due to customer turnover and renewals at lower subscription levels... customers renew their contracts at levels more closely aligned with their current needs, rather than opting for larger commitments based on anticipated future demand." That is a demand-environment explanation, and it is the company's own.

The subtler movement is the gap. Large customers used to expand three points faster than the average — 120% against 117% in FY2024. In FY2026 they expand one point faster, 110% against 109%. Whatever advantage the enterprise cohort had in expansion has nearly closed. Meanwhile the *number* of large customers has grown strongly, from 202 to 333. Braze is winning more big accounts and growing each of them more slowly.

Average revenue per customer can be bounded but not known: \$738.2m of revenue across 2,609 customers is **about \$283,000 each**. That is a bound, not a price — it mixes every contract size and includes professional services — and it should always be said as "bounded at", never as "costs".

7.3 · Who they name, and who buyers actually shortlist

The 10-K names four competitors: **Adobe, Salesforce, Iterable, Klaviyo**, alongside the claim that "none of our competitors currently offer comparable comprehensive customer engagement solutions."

Gartner Peer Insights publishes what buyers *also considered*, derived from the buyers rather than the vendor. That list has eight: **Salesforce, Adobe, Iterable, Oracle, Optimove, Blueshift, MoEngage, CleverTap**.

	NAMED BY BRAZE	ON THE BUYER SHORTLIST
Adobe, Salesforce, Iterable	yes	yes
Klaviyo	yes	no
Oracle, Optimove, Blueshift, MoEngage, CleverTap	no	yes

THIRD-PARTY

source: sources/panels/gartner.md:113-121 @ 2026-09-02

The overlap is real — three of the four names are shared, so this is not a vendor misreading its own market. The asymmetry is what matters: buyers weigh Braze against five vendors it never names, and those five are mostly the mobile-engagement specialists (MoEngage, CleverTap, Blueshift, Optimove) rather than the suite vendors Braze positions against. A competitor briefing built only from Braze's own comparison pages would miss five of the eight vendors in the room.

Where Gartner's reviewers rate Braze *above* the two largest alternatives, they name the same things twice: service and support, and ease of integration and deployment — against Salesforce, with evaluation and contracting added.

7.4 · What customers say

Site-level aggregates, which are the figures to quote. The captured review bodies are a sample of fourteen records used for vocabulary and theme, never for a percentage.

PANEL	RATING	BASE
G2	4.5 / 5	1,702 reviews
Gartner Peer Insights	4.5 / 5	263 ratings
TrustRadius	8.8 / 10	348 reviews

THIRD-PARTY

source: sources/panels/g2.md:46 @ 2026-09-02

G2's own coded tags, which are counted over its full review base rather than over the sample captured here, put the praise at Ease of Use (385), Intuitive (188), Customer Support (151), Helpful (148) and Features (138); and the criticism at Missing Features (140), Learning Curve (139), Limitations (102), Limited Features (93) and Steep Learning Curve (86). TrustRadius's synthesis adds a coded figure of its own: reporting is "often described as limited and unintuitive, a sentiment shared by 36% of reviewers", of the 36 reviews it published in the preceding eighteen months.

Two themes recur across unrelated panels and are therefore worth more than the rest: **reporting is the most consistent complaint**, and **learning curve is the second**. The reporting complaint is independently consistent with the platform evidence in chapter 4 — out-of-the-box reporting being thin enough that Currents, a paid add-on, is the route to raw data.

Running `tools/code_reviews.py` over the whole corpus produces theme percentages against a denominator of 860 records, but 841 of those are GitHub issues rather than buyer reviews. That denominator is stated wherever the

script's output is used, and its themes describe the SDK surface, not satisfaction.

7.5 · What employees say

Glassdoor, captured signed-in: **4.1 out of 5 from 524 ratings, 82% would recommend, 90% approve of the CEO, 71% positive business outlook.**

Work-life balance tracks the overall rating almost exactly over the six months to September 2026 — both hovering around 4.0 to 4.1 — so there is no divergence worth calling a trend, and none is claimed. Glassdoor's own summary names the weak spots as management effectiveness and "clarity in direction", and "limited upward mobility and discrepancies in compensation relative to market rates".

THIRD-PARTY

source: `sources/panels/glassdoor.md:68-92` @ 2026-09-02

The careers board shows roughly 284 to 300 open roles across fifteen functions and twenty-six locations. An exact split by function could not be captured — the board's filter would not drive reliably — so no percentage is offered, and the gap is open question 56 rather than an estimate.

What would change this chapter

A fourth consecutive year of falling net retention would turn a demand-environment explanation into a competitive one. A new name in the 10-K competition paragraph, or a change in the Gartner shortlist, would move §7.3. Access to any of the three paywalled by-company-size rating breakdowns would finally settle whether satisfaction falls with customer size — the one hypothesis this analysis could not test.

8 • What we could not answer

Nine questions this analysis could not close from public sources, each with the source that would close it — plus the two places where a number moved and the old value is deliberately still visible.

A gap written down is evidence. A gap not written down is a mistake. The full backlog lives in docs/QUESTIONS.md; what follows is the residue that matters to a reader of this record, in the order a competitor would care about.

8.1 • The nine open questions

#	QUESTION	WHAT WOULD CLOSE IT
52	Why does the sub-processor disclosure name only Amazon and Google as hosting providers when the allowlist IPs Braze publishes for instance US-08 are all registered to Microsoft?	A revised sub-processor list, the DPA schedule a customer receives, or a direct answer from Braze
53	Is the aze region code Azure, and is US-08 a distinct class of instance?	A Braze statement, a customer on US-08, or a wider certificate-transparency capture
54	Does satisfaction fall as customer size rises?	Paid access to the G2, Gartner or TrustRadius rating-by-company-size breakdowns — all three paywalled
55	Why was the profile-lookup rate limit cut tenfold for customers onboarding after 22 August 2024, and were existing customers told?	A changelog entry, a support answer, or customers either side of the date
56	What is the exact split of open roles by function?	The Greenhouse board API, or a working pass over the board's department filter
57	Can the customer roster be verified independently of the self-published stories?	Tag crawls, certificate transparency naming customer subdomains, or job ads naming the stack
58	What else is in the certificate-transparency estate?	A CERTSPOTTER_TOKEN; crt.sh returned HTTP 502 throughout this run
59	Does the gross-margin decline continue, and how much of it is AI cost rather than acquisition amortisation?	The FY2027 10-K, or quarterly cost-of-revenue detail
46	Actual list pricing and discounting policy	A public-sector procurement award, or a customer contract

Two of these deserve emphasis because they bound claims made elsewhere in this record.

Question 58 bounds chapter 4.5. The host list is partial — 833 hosts, captured through a rate-limited fallback after the primary source returned errors all day. Everything said about what *was* provisioned stands. Nothing is said about what is absent, because the check was not exhaustive and an absence from an incomplete list is not a finding.

Question 54 is the one hypothesis that could not be tested at all. The analysis went in expecting enterprise satisfaction to be lower than small-business satisfaction, as it usually is. All three review sites paywall exactly that breakdown, and only seven of the 860 coded records carry a customer-size segment — far too few for anything. The nearest audited proxy points mildly the other way and measures something different: net retention among customers with \$500k or more of annual recurring revenue is slightly *above* the all-customer figure (chapter 7). That is expansion, not satisfaction, and it is not offered as a substitute.

8.2 · Where a number moved

The corrections log, mirroring docs/FACTS.md, with the old values still visible so a stale copy of this document can be recognised on sight.

FACT	WAS	IS NOW	WHY IT CHANGED
Status-page grouping	"132 in 12 groups"	132 rows in 17 named groups, of which 15 are regional clusters	The setup-time reconnaissance counted by eye. Counting the CSV gives 17 named groups and 15 clusters. DECK-SPEC.md repeated the same error and is wrong for the same reason
Capability page counts	Journey orchestration 319, Identity resolution 136, Segmentation 339	Canvas 249, Identity 110, Segmentation 242	The pattern set moved, not the product. The taxonomy was revised from generic category words to Braze's own product names. Counts from the two runs are not comparable, and both are reproducible from the pattern column of data/capabilities.csv
KakaoTalk's status	Read at first pass as documented but not marketed	Marketed, but missing from Braze's own docs channel index	A truncated search appeared to show no marketing page. Checking the full /product/ enumeration found one

The middle row is the important one, and it is the reason the taxonomy is published rather than described. A number that moves because the measurement moved is not a finding, and presenting it as one would be the easiest way to mislead an audience in this entire project.

8.3 · What was substituted, and what that cost

BLOCKED	SUBSTITUTED WITH	WHAT IS LOST
crt.sh (HTTP 502 all day, twice)	Cert Spotter, rate-limited without a token	A partial host list. Stated as partial everywhere it is used
Rating by company size on all three review panels	Nothing equivalent exists publicly	Hypothesis 7 is unresolved rather than answered
Per-department headcount on the careers board	The department taxonomy and a front-of-list sample	No function split. No percentage is offered
Independent customer detection	Not attempted	The 178 published stories are not treated as a roster

Every review panel that returned HTTP 403 to a script was captured through a browser session instead, and Glassdoor required the operator to sign in before its content became readable at all — recorded here because the capture route is part of the evidence.

What would change this chapter

Any of the nine questions closing. The most valuable single answer would be question 52, because it decides whether chapter 4.5 is an observation about infrastructure or something a buyer's procurement team needs to raise.

Slide map

The deck in order, generated from the built deck so coverage cannot be overstated. No evidence lives here — if a number appears in this table it belongs in a chapter.

#	SLIDE	RECORD
01	Competitor Analysis	\$4
02	How we approach it	\$8
03	How we grade	\$8
04	Part I: The company	\$1
05	Five things	\$1
06	Who they say they are	\$6
07	Origins	\$1
08	How they got this big	\$2
09	Acquisition: OfferFit	\$3
10	Acquisition: North Star Y	\$3
11	What it costs	\$7
12	Who uses it	\$7
13	Where they operate	\$7
14	What customers say	\$7
15	What employees say	\$7
16	Who they compete with	\$7
17	Part II: The product	\$4
18	How one campaign works	\$4
19	How data moves	\$4
20	Stage 1: Data	\$4
21	Stage 2: Identity	\$4
22	Stage 3: Decisioning	\$4
23	Stage 4: Building	\$4
24	Stage 5: Content	\$4
25	Stage 6: Delivery	\$5
26	Stage 7: Interaction	\$4
27	Channels	\$5
28	Integrations	\$5
29	Infrastructure	\$4
30	Analytics	\$4
31	The AI, honestly	\$6
32	Part III: Strategy	\$8
33	Where the money goes	\$2
34	What comes next	\$4
35	What protects them	\$6
36	Part IV: Open questions	\$8
37	Deep dive: how real-time is it	\$4

#	SLIDE	RECORD
38	Reliability, measured	\$4
39	What we could not answer	\$8
40	What to research next	\$8
41	Close	\$4